

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 7

STRUCTURED QUERY LANGUAGE (SQL)

STRUCTURED QUERY LANGUAGE (SQL): OUTLINE

SQL Basic Structure and Operations

Additional Basic Operations

Aggregate Functions


Nested Subqueries and Set Operations

Database Definition

Database Modification


Using SQL in Applications

SQL OVERVIEW

- SQL is the most common relational query language.
 **SQL is used in all commercial relational DBMSs.**
- Commercial relational DBMSs have **different features** of SQL, but the **basic structure is the same**.
 - Data Manipulation Language
 - Data Definition Language
 - Integrity Constraint Specification
 - View Definition
 - Embedded/Dynamic SQL
 - Transaction Management
 - Security Management
 -
 -

BASIC STRUCTURE OF SQL QUERIES

- SQL is based on **set** and **relational algebra** operations with certain **modifications** and **enhancements**.
- An SQL query has the basic form:



- The result of an SQL query is a relation (**but it may contain duplicates**).

👉 **SQL queries can be nested (composed).**

EXAMPLE BANK RELATIONAL SCHEMA

Branch(branchName, branchCity, assets)

Client(clientName, clientStreet, clientCity)

Loan(loanNo, amount, *branchName*)

Account(accountNo, balance, *branchName*)

Borrower(*clientName*, *loanNo*)

Depositor(*clientName*, *accountNo*)

Attribute names in
italics are foreign
key attributes.

PROJECTION: SELECT CLAUSE

- The **select** clause corresponds to the relational algebra **projection** (π) operation.

Query: Find the names of all branches in the Loan relation.

```
select branchName
from Loan;
```

=

```
 $\pi_{\text{branchName}}(\text{Loan})$ 
```

- An asterisk (*) in the **select** clause denotes “**all attributes**”.

```
select *
from Loan;
```

=

```
select loanNo, amount, branchName
from Loan;
```

👉 **Attributes specified in the **select** clause must be defined in the relations in the **from** clause.**

👉 **SQL does not remove duplicates in the result by default.**



PROJECTION: DUPLICATE REMOVAL

- The keyword **distinct** forces the **removal of duplicates**.

Query: Find the unique names of all branches in the Loan relation.

```
select distinct branchName  
from Loan;
```

-----> **force** the DBMS to
remove duplicates

- The keyword **all** specifies that **duplicates should not be removed**.

```
select all branchName  
from Loan;
```

-----> **force** the DBMS not to
remove duplicates
(same as omitting **all**)



PROJECTION: ARITHMETIC OPERATIONS

- The **select** clause can contain arithmetic expressions involving the operators $+$, $-$, \div and \times that can operate on constants or attributes of tuples.

Query: Multiply the amount of each loan by 100.

```
select loanNo, amount*100, branchName  
from Loan;
```

This query returns a relation which is the same as the **Loan** relation, except that the attribute **amount** is multiplied by 100.



SELECTION: WHERE CLAUSE

- The **where** clause corresponds to the relational algebra **selection predicate** (σ) and specifies conditions that tuples in the relations in the **from** clause must satisfy.

Query: Find all loan numbers for loans made at the Tsimshatsui branch whose loan amount is greater than \$1200.

```
select loanNo
from Loan
where branchName='Tsimshatsui'
and amount>1200;
```

String values **must**
be enclosed in
single quotes.
Numeric values **do**
not require quotes.

≡

$$\pi_{\text{loanNo}} (\sigma_{\text{branchName}='Tsimshatsui' \wedge \text{amount}>1200}(\text{Loan}))$$

👉 **Attributes specified in a **where** clause must be defined in the relations in the **from** clause.**

SELECTION: WHERE CLAUSE (cont'd)

- SQL provides the **between** operator for convenience.

Query: Find the loan number of loans whose amount is between \$100,000 and \$200,000 (i.e., $\geq \$100,000$ and $\leq \$200,000$).

```
select loanNo
from Loan
where amount between 100000 and 200000;
```

- Can also use **not between** (i.e., $< \$100,000$ and $> \$200,000$).

```
select loanNo
from Loan
where amount not between 100000 and 200000;
```

- SQL allows **Boolean operators** **and**, **or** and **not** to be used in a **where** clause as well as **arithmetic expressions**.

NATURAL JOIN: WHERE CLAUSE

- A natural join can be specified by adding the appropriate join condition in the where clause.

Query: Find the name and loan number of all clients.

```
select clientName, borrower.loanNo
from Borrower, Loan
where Borrower.loanNo=Loan.loanNo;
```

Attribute names **must** be qualified if ambiguous.

- SQL provides a shorthand way to specify a natural join.

```
select clientName, loanNo
from Borrower natural join Loan;
```

Attribute names **cannot** be qualified in a natural join.

Why?

What must be true for these two queries to be equivalent?



CARTESIAN PRODUCT: FROM CLAUSE

- The **from** clause corresponds to the relational algebra **Cartesian-product operation** (\times).

Query: Find the Cartesian product of borrower and loan.

```
select *  
from Borrower, Loan;
```

- This can also be specified as

```
select *  
from Borrower cross join Loan;
```

✎ **A from clause with more than one relation is rarely used without a where clause.**

SET OPERATIONS: UNION, INTERSECT, EXCEPT

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations \cup , \cap and $-$.

Oracle Note

The keyword **minus** is used rather than **except**.

- Each of the set operations **automatically removes duplicates**.
- The operations **union all**, **intersect all** and **except all** keep all duplicates.
- Suppose a tuple occurs **m** times in **r** and **n** times in **s**, then it occurs:
 - m + n** times in **r union all s**
 - min(m, n)** times in **r intersect all s**
 - max(0, m-n)** times in **r except all s**

Oracle Note

Only **union all** is supported.

SET OPERATIONS: EXAMPLES

Query: Find all clients who have a loan, an account, or both.

```
(select clientName from Depositor)  
union  
(select clientName from Borrower);
```

Query: Find all clients who have both a loan and an account.

```
(select clientName from Depositor)  
intersect  
(select clientName from Borrower);
```

Query: Find all clients who have an account, but no loan.

```
(select clientName from Depositor)  
minus  
(select clientName from Borrower);
```



RENAME ATTRIBUTES: AS CLAUSE

- Attributes can be renamed using the **as** clause:

old-name as new-name

Query: Find the name and loan number of all clients having a loan at the Central branch; replace the attribute name loanNo with the name loanId.

```
select distinct clientName, Borrower.loanNo as loanId
from Borrower, Loan
where Borrower.loanNo=Loan.loanNo
and branchName='Central';
```

Oracle Note

The keyword **as** is optional in the **select** clause.

- The SQL standard also allows relations in the **from** clause to be renamed using the **as** clause.

Oracle Note

The keyword **as** is not allowed in the **from** clause.



RENAME RELATIONS

- Renaming relations is convenient for replacing long relation names used multiple times in a query with shorter ones.

Query: Find the client names and their loan numbers for all clients having a loan at *some* branch; replace the column name loanNo with the name loanId.

```
select distinct clientName, B.loanNo loanId
from Borrower B, Loan L
where B.loanNo=L.loanNo;
```

Oracle Note

Relations in the from clause are renamed using an identifier without the keyword as.

- An identifier for a relation (such as B and L above) is referred to as a *correlation name* in SQL.
 - Also known as *table alias*, *correlation variable* or *tuple variable*.




RENAME RELATIONS (CONT'D)

- Renaming a relation is required when we want to **compare tuples in the same relation** (self-join).

Query: Find the names of all branches that have greater assets than some (i.e., at least one) branch located in Hong Kong.

```
select distinct B1.branchName
from Branch B1, Branch B2
where B1.assets > B2.assets and B2.branchCity = 'Hong Kong';
```



STRING PATTERN MATCHING: LIKE OPERATOR

- The **like** operator is used for matching characters in strings.
- Character attributes can be compared to a pattern using:
 - % (percent) matches any substring.
 - _ (underscore) matches any single character.

Query: Find the name of all clients whose streets include the substring 'Main' (e.g., Mainroad, Mainly Avenue, Mainmount Street, ...).

```
select clientName
from Client
where clientStreet like '%Main%';
```

👉 **Pattern matching is *usually* case-sensitive.**



STRING PATTERN MATCHING: LIKE OPERATOR

(CONT'D)

- To include the special pattern matching characters in a string, SQL allows the specification of an escape character.
 - Suppose we use backslash (\) as the escape character.
 - `like '20\%%' escape '\'` matches all strings beginning with “20%”
 - `like 'pair_%' escape '\'` matches all strings beginning with “pair_”
- To include a single quote in a string, **use two single quotes**.
 - `like 'Toms"s%'` matches all strings beginning with “Tom’s”

STRING PATTERN MATCHING: REGEXP_LIKE OPERATOR

- The `regexp_like` operator is used for specifying patterns similar to that used in Unix regular expressions.

Query: Find the names of those clients whose names begin with Steven or Stephen (i.e., the name begins with 'Ste' followed by either 'v' or 'ph' followed by 'en' followed by any other characters).

```
select clientName
from Client
where regexp_like (clientName, '^Ste(v|ph)en');
```

Query: Find the names of those clients with a double vowel (i.e., double a, e, i, o or u) in their name, regardless of case.

```
select clientName
from Client
where regexp_like (clientName, '([aeiou])\1', 'i');
```



STRING PATTERN MATCHING: REGEXP_LIKE OPERATOR

Usage: `regexp_like(source_string, pattern, [match_parameter])`

where:

- *source_string* is a search value (usually an attribute name);
- *pattern* is a regular expression;
- *match_parameter* specifies a matching behaviour as follows
 - 'i' specifies case-insensitive matching.
 - 'c' specifies case-sensitive matching.
 - 'n' allows the period (.), which is normally the match-any-character wildcard character, to match the newline character.
 - 'm' treats the source string as multiple lines.

If *match_parameter* is omitted then:

- The default case sensitivity is used (usually case-sensitive).
- A period (.) does not match the newline character.
- The source string is treated as a single line.

ORDERING RESULT TUPLES: ORDER BY CLAUSE

Query: Find, in alphabetic order, the names of all clients having a loan at the Central branch.

```
select distinct clientName
from Borrower, Loan
where Borrower.loanNo=Loan.loanNo
      and branchName='Central'
[order by clientName]
```

Ordering options

asc - ascending
(default)

desc - descending

- Can sort on multiple attributes.
e.g., `order by` clientName `desc`, amount `asc`

✎ Since sorting a large number of tuples may be costly, it is desirable to use the `order by` clause only when necessary.