

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 8 EXERCISES

STRUCTURED QUERY LANGUAGE (SQL)

EXAMPLE RELATIONAL SCHEMA AND DATABASE

Sailor(sailorId, sName, rating, age)

Boat(boatId, bName, color)

Reserves(sailorId, boatId, rDate)

Attribute names in
italics are foreign
key attributes.

Sailor

<u>sailorId</u>	sName	rating	age
22	Dustin	7	45
29	Brutus	1	33
31	Lubber	8	55
32	Andy	8	25
58	Rusty	10	35
64	Horatio	7	35
71	Zorba	10	16
74	Horatio	9	35
85	Art	3	25
95	Bob	3	63
99	Chris	10	30

11 tuples

Reserves

<u>sailorId</u>	<u>boatId</u>	<u>rDate</u>
22	101	10/10/17
22	102	10/10/17
22	103	08/10/17
22	104	07/10/17
31	102	10/11/17
31	103	06/11/17
31	104	12/11/17
64	101	05/09/17
64	102	08/09/17
74	103	08/09/17
99	104	08/08/17

11 tuples

Boat

<u>boatId</u>	bName	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red
105	Serenity	Cyan

5 tuples

EXERCISE 1

Find the name and the number of reservations for each red boat.

☞ (Interlake, 3), (Marine, 3)

Is this a correct solution?

```
select bName, count(*) reservationCount
from Boat natural join Reserves
where color='red'
group by Boat.boatId;
```

Illegal!!!
Why?

☞ All non-aggregate attributes in the select clause must appear in the group by clause (i.e., bName must appear in the group by clause).

EXERCISE 1 (CONTD)

Find the name and the number of reservations for each red boat.

👉 (Interlake, 3), (Marine, 3)

Do you see any problems with this solution?

```
select bName, count(*) reservationCount
from Boat natural join Reserves
where color='red'
group by bName
```

sailorId	boatId	rDate	bName	color
22	102	10/10/17	Interlake	red
22	104	07/10/17	Marine	red
31	102	10/11/17	Interlake	red
31	104	12/11/17	Marine	red
64	102	08/09/17	Interlake	red
99	104	08/08/17	Marine	red

Reservations for red boats.

a group

a group

bName	reservationCount
Interlake	3
Marine	3

Name and count of the number of reservations for each red boat.

EXERCISE 1 (CONTD)

Suppose we change the query to this.

What is the result?

Find the name and the number of reservations for each boat.

```
select bName, count(*) reservationCount
from Boat natural join Reserves
group by bName
```

sailorId	boatId	rDate	bName	color
22	101	10/10/17	Interlake	blue
64	101	05/09/17	Interlake	blue
22	102	10/10/17	Interlake	red
31	102	10/11/17	Interlake	red
64	102	08/09/17	Interlake	red
22	103	08/10/17	Clipper	green
31	103	06/11/17	Clipper	green
74	103	08/09/17	Clipper	green
22	104	07/10/17	Marine	red
31	104	12/11/17	Marine	red
99	104	08/08/17	Marine	red

a group

a group

a group

bName	reservationCount
Interlake	5
Clipper	3
Marine	3

Since **bName** is not unique, grouping on it can get an incorrect result!



EXERCISE 1 (CONTD)

Find the name and the number of reservations for each boat.

Correct solution.

```
select bName, count(*) reservationCount
from Boat natural join Reserves
group by Boat.boatId, bName;
```

Recall:
attributes in the group by clause do not have to appear in the select clause.

sailorId	boatId	rDate	bName	color
22	101	10/10/17	Interlake	blue
64	101	05/09/17	Interlake	blue
22	102	10/10/17	Interlake	red
31	102	10/11/17	Interlake	red
64	102	08/09/17	Interlake	red
22	103	08/10/17	Clipper	green
31	103	06/11/17	Clipper	green
74	103	08/09/17	Clipper	green
22	104	07/10/17	Marine	red
31	104	12/11/17	Marine	red
99	104	08/08/17	Marine	red

a group

a group

a group

a group

bName	reservationCount
Interlake	2
Interlake	3
Clipper	3
Marine	3



EXERCISE 2

Find the sailor id and number of reservations made for each sailor.

☞ (22, 4), (29, 0), (31, 3), (32, 0), (58, 0), (64, 2),
(71, 0), (74, 1), (85, 0), (95, 0), (99, 1)

```
select sailorId, count(sailorId) reservationCount
from Reserves
group by sailorId;
```

How to include
all sailors?

sailorId	reservationCount
22	4
31	3
64	2
74	1
99	1

How about joining **Sailor** and **Reserves**?

```
select sailorId, count(sailorId) reservationCount
from Sailor natural join Reserves
group by sailorId;
```

What's the
problem?

sailorId	reservationCount
22	4
31	3
64	2
74	1
99	1

EXERCISE 2 (CONTD)

Find the sailor id and number of reservations made for each sailor.

☞ (22, 4), (29, 0), (31, 3), (32, 0), (58, 0), (64, 2),
(71, 0), (74, 1), (85, 0), (95, 0), (99, 1)

sailorId	sName	rating	age	boatId	rDate
22	Dustin	7	45	101	10/10/17
22	Dustin	7	45	102	10/10/17
22	Dustin	7	45	103	08/10/17
22	Dustin	7	45	104	07/10/17
31	Lubber	8	55	102	10/11/17
31	Lubber	8	55	103	06/11/17
31	Lubber	8	55	104	12/11/17
64	Horatio	7	35	101	05/09/17
64	Horatio	7	35	102	08/09/17
74	Horatio	9	35	103	08/09/17
99	Chris	10	30	104	08/08/17
29	Brutus	1	33	-	-
32	Andy	8	25	-	-
58	Rusty	10	35	-	-
71	Zorba	10	16	-	-
85	Art	3	25	-	-
95	Bob	3	63	-	-

```
select sailorId, count(sailorId) reservationCount
from Sailor natural join Reserves
group by sailorId;
```

sailorId	reservationCount
22	4
31	3
64	2
74	1
99	1

☞ Some Sailor tuples have no match in the Reserves relation.
How to deal with this problem?

EXERCISE 2 (CONT'D)

Find the sailor id and number of reservations made for each sailor.

☞ (22, 4), (29, 0), (31, 3), (32, 0), (58, 0), (64, 2),
(71, 0), (74, 1), (85, 0), (95, 0), (99, 1)

```
select sailorId, count(boatId) reservationCount
from Sailor natural left outer join Reserves
group by sailorId;
```

Note: **left outer join** keeps the common attributes, while
natural left outer join removes the common attributes.

Is this a
correct
solution?
NO! Why?

```
select sailorId, count(sailorId) reservationCount
from Sailor natural left outer join Reserves
group by sailorId;
```

Counting is done on the sailor ids and all
of them appear at least once in the result.



EXERCISE 3

Find the records (tuples) of the sailors with the highest rating.

☞ (58, Rusty, 10, 35), (71, Zorba, 10, 16), (99, Chris, 10, 30)

Is this a
correct
solution?
NO! Why?

```
select *  
from Sailor  
where rating=max(rating);
```

There is no **max(rating)** value to
compare in the **where clause**.
☞ The **max rating value** has to be
obtained by a **select statement!**

Is this a
correct
solution?
NO! Why?

```
select *, max(rating)  
from Sailor;
```

A query that returns multiple
tuples cannot contain an
aggregate function.

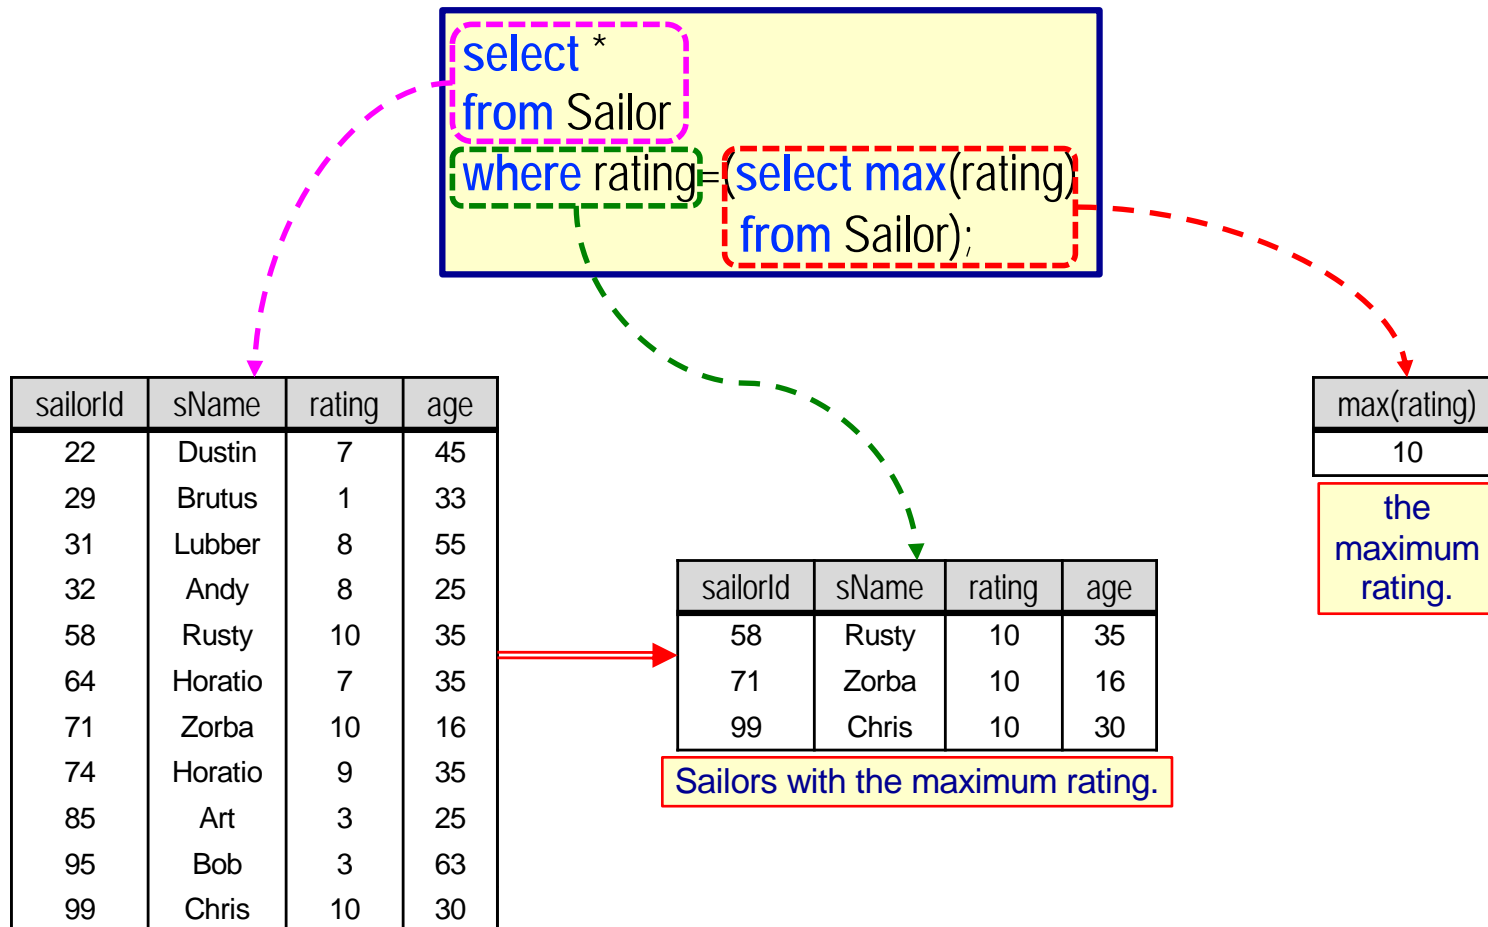
☞ There are multiple tuples in the
result, but only one max value!



EXERCISE 3 (CONTD)

Find the records (tuples) of the sailors with the highest rating.

✈️ (58, Rusty, 10, 35), (71, Zorba, 10, 16), (99, Chris, 10, 30)

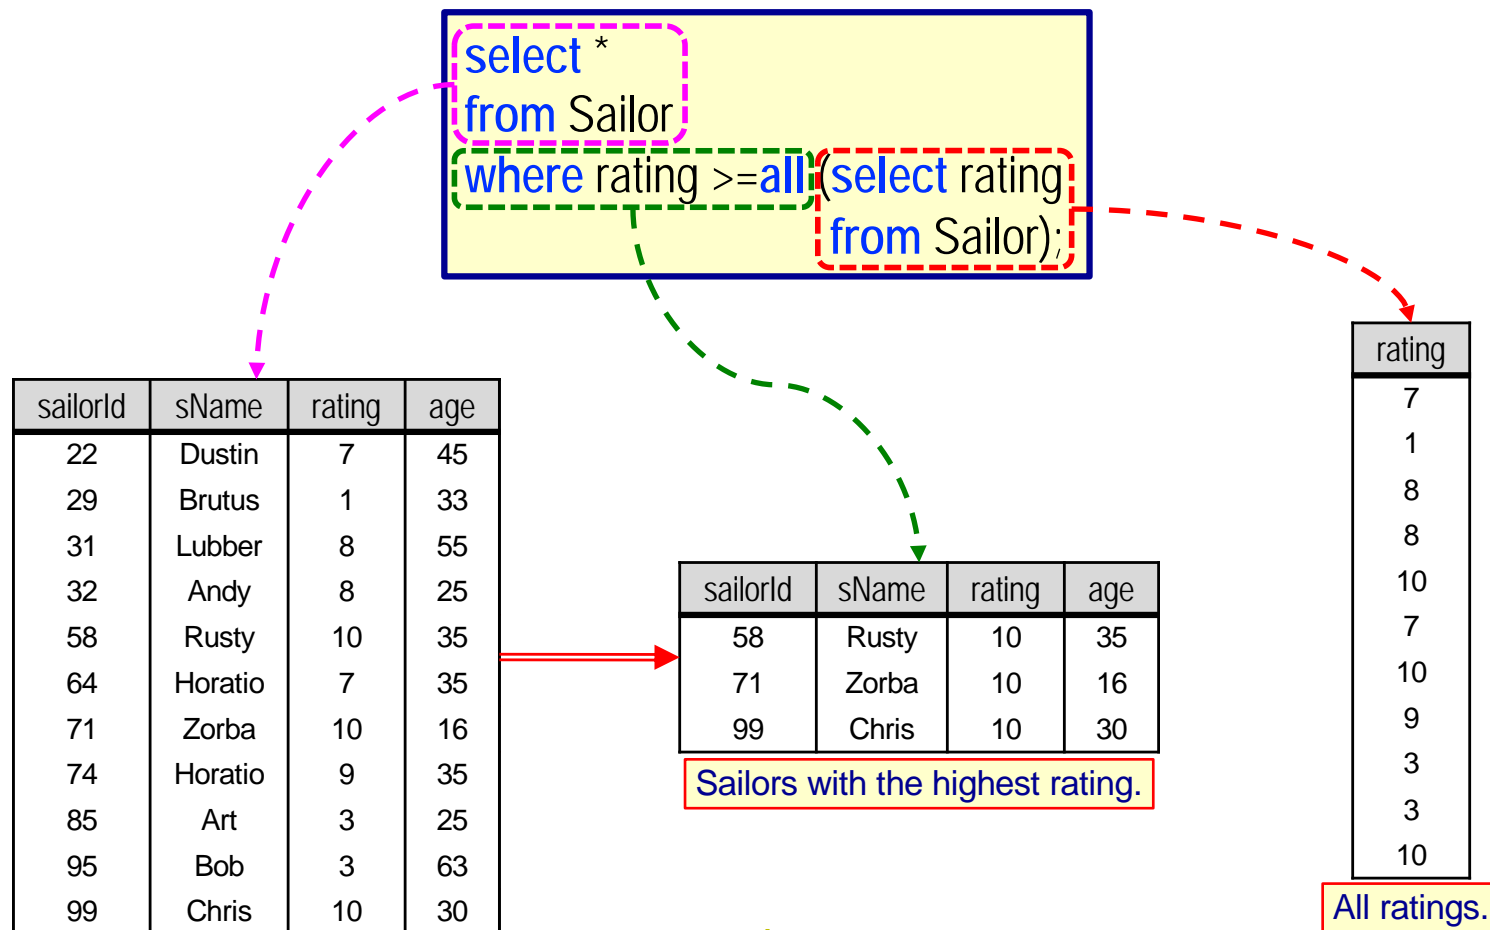


EXERCISE 3 (CONTD)

Use set
membership

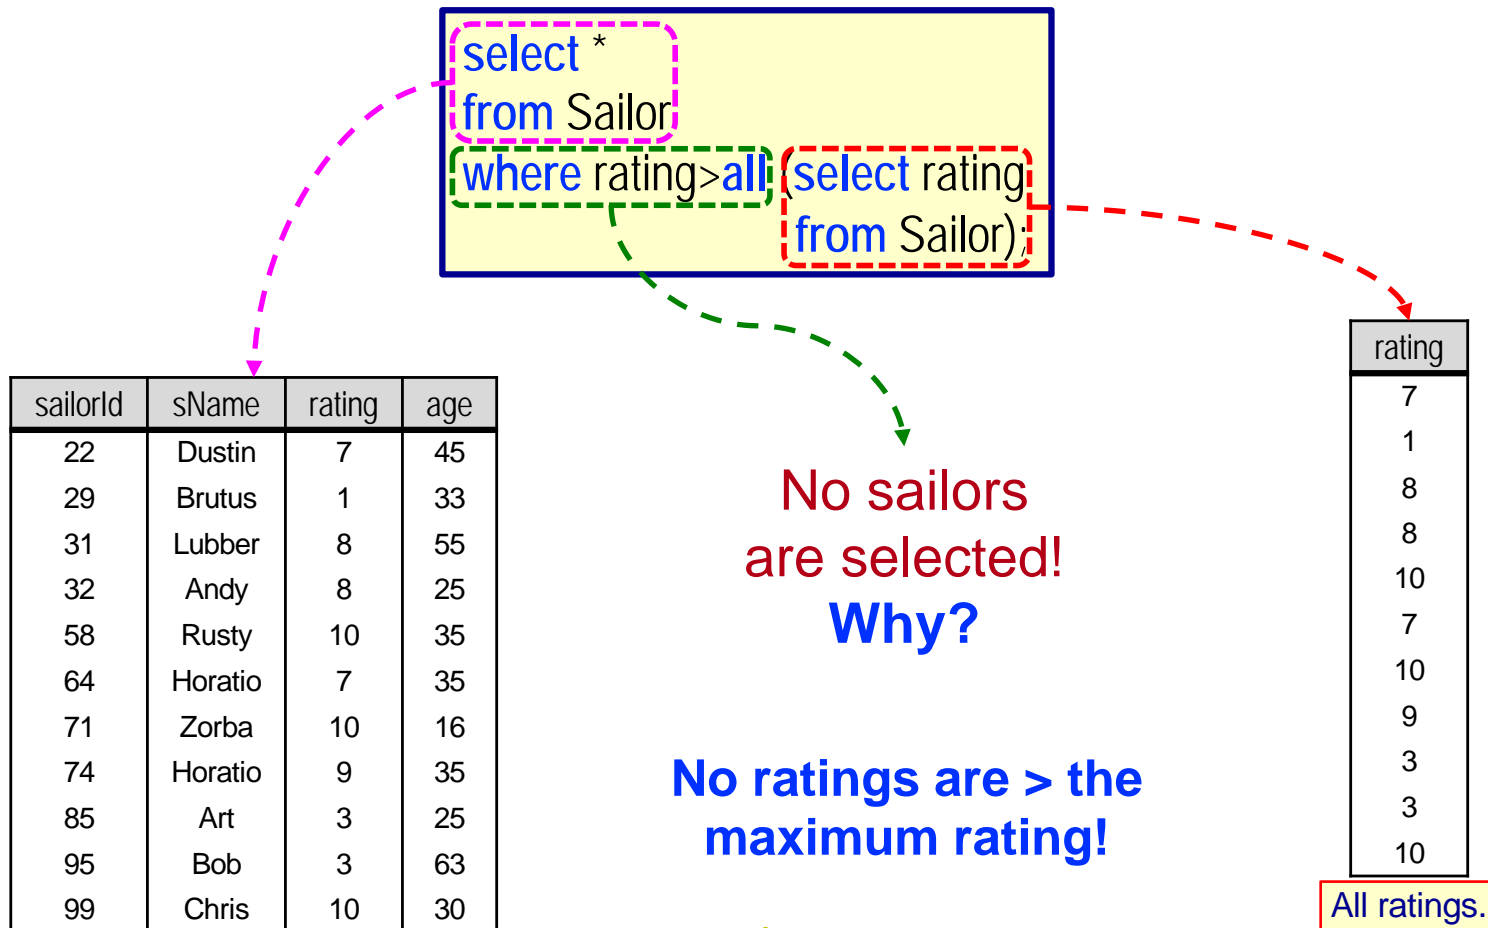
Find the records (tuples) of the sailors with the highest rating.

✈️ (58, Rusty, 10, 35), (71, Zorba, 10, 16), (99, Chris, 10, 30)



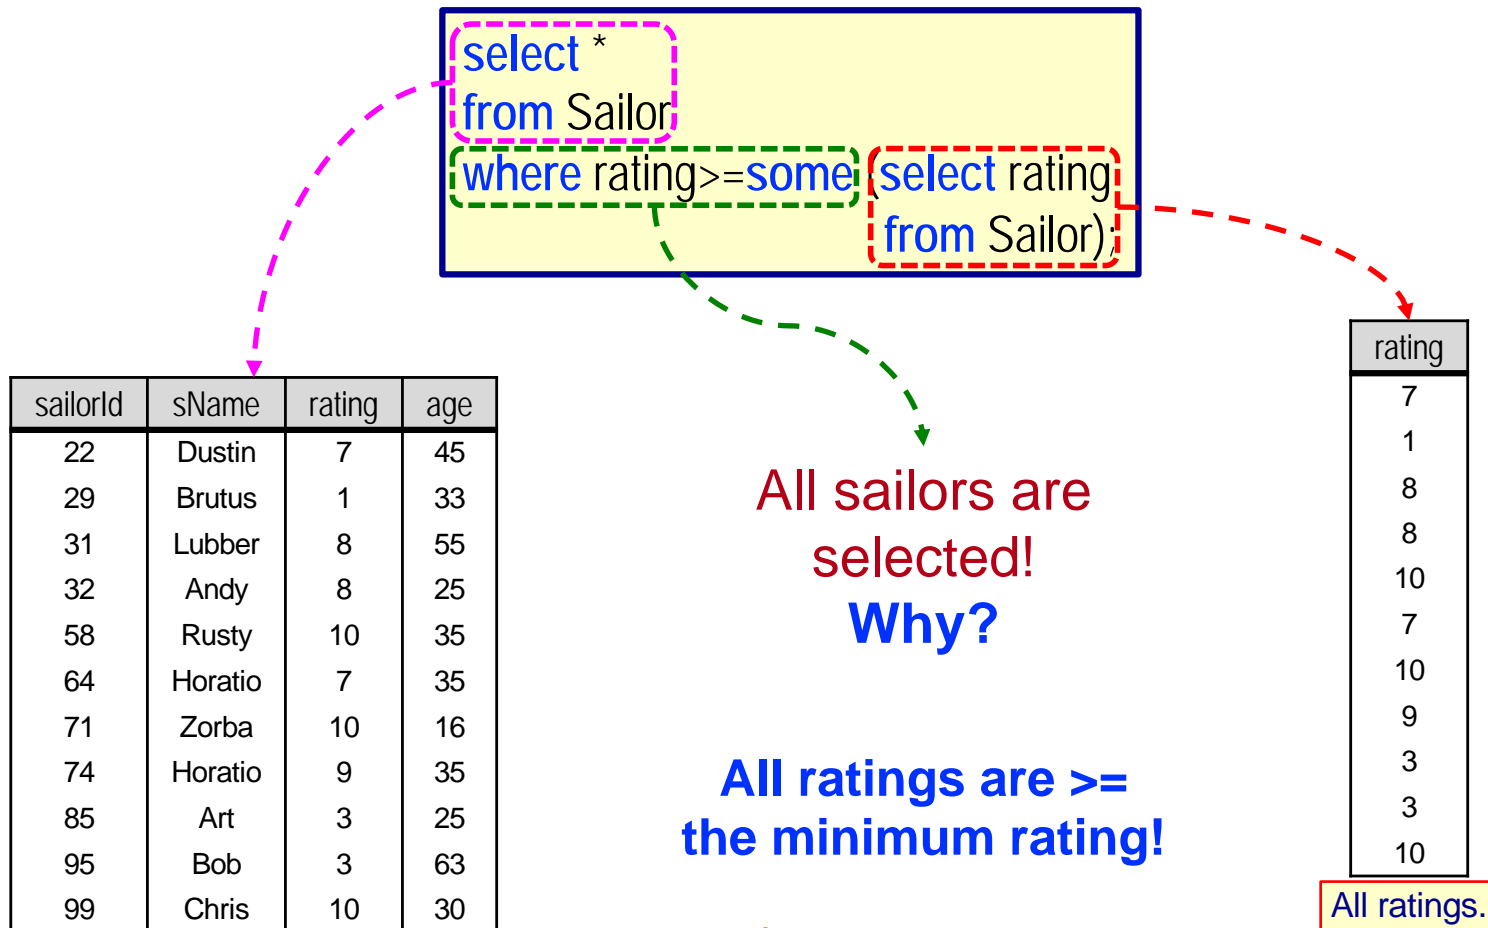
EXERCISE 3 (CONT'D)

What is the result if we replace “>=all” with “>all”?



EXERCISE 3 (CONT'D)

What is the result if we replace “>=all” with “>some”?



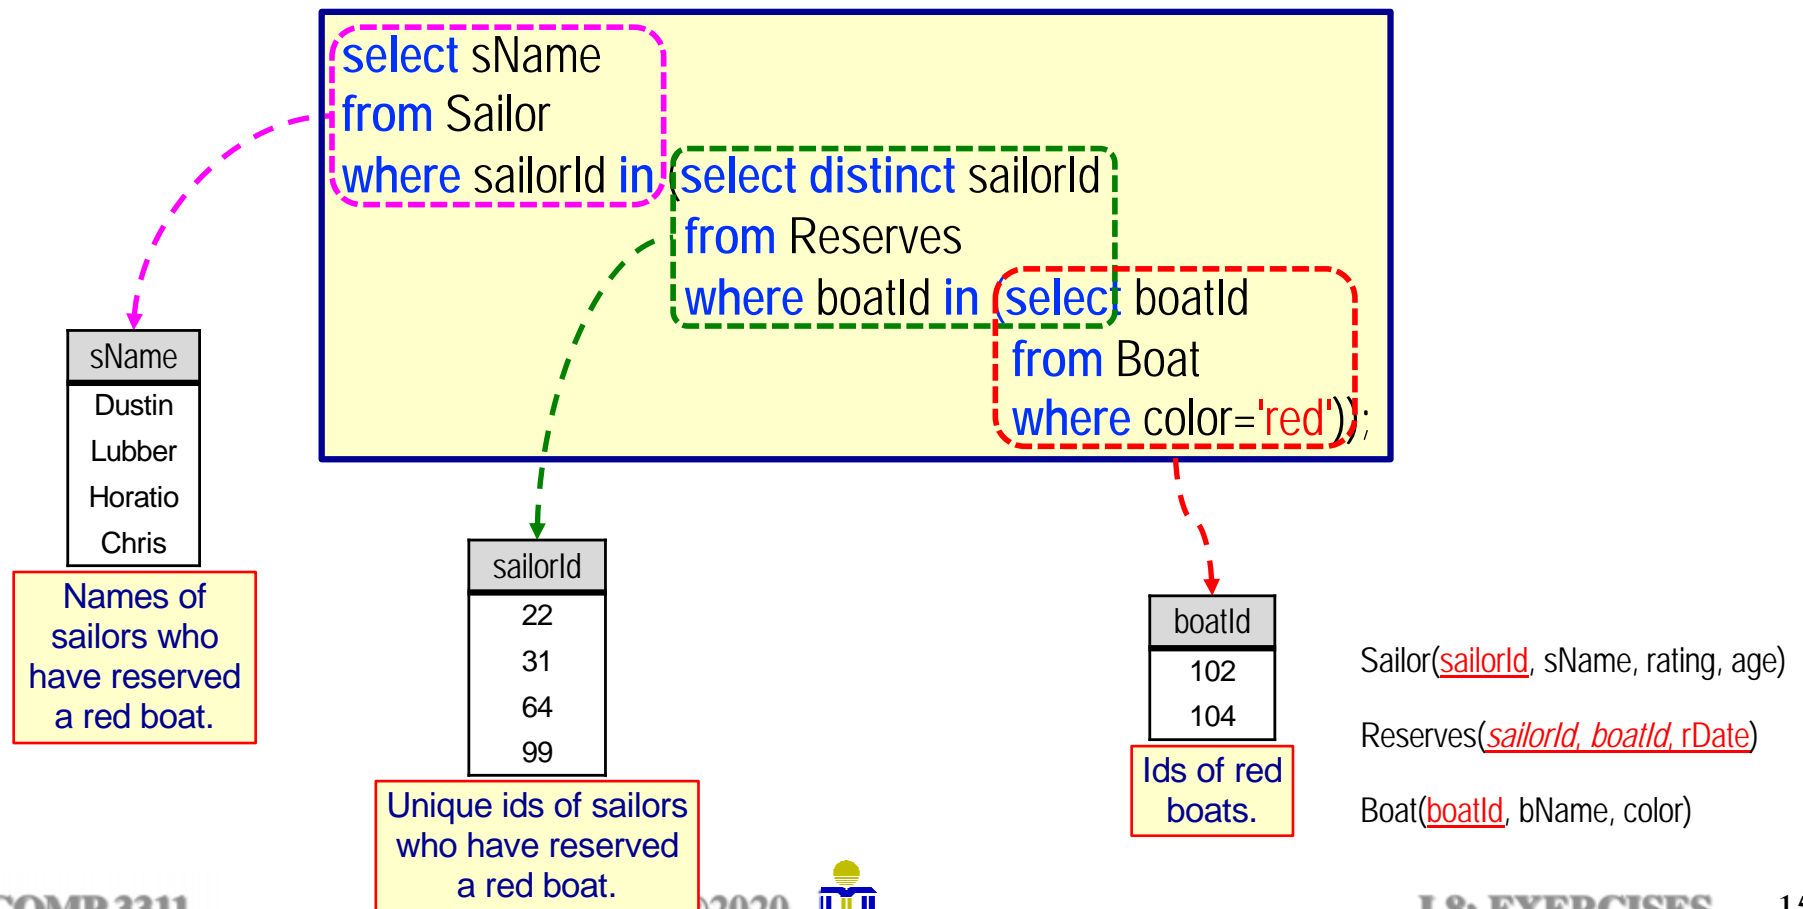
EXERCISE 4

DO NOT
use JOIN

**Find the names of sailors who
have reserved a red boat.**

Use *only* set
membership

☞ **Dustin, Lubber, Horatio, Chris**



EXERCISE 4 (cont'd)

What if we replace the first **in** with **not in**?

Stated in words, what does this result represent?

```
select sName
from Sailor
where sailorId not in (select distinct sailorId
                       from Reserves
                       where boatId in (select boatId
                                       from Boat
                                       where color='red'))
```

sName
Brutus
Andy
Rusty
Zorba
Horatio
Art
Bob

Names of sailors who have not reserved a red boat (including reserved no boat).

sailorId
22
31
64
99

Unique ids of sailors who have reserved a red boat.

boatId
102
104

Ids of red boats.

Sailor(sailorId, sName, rating, age)

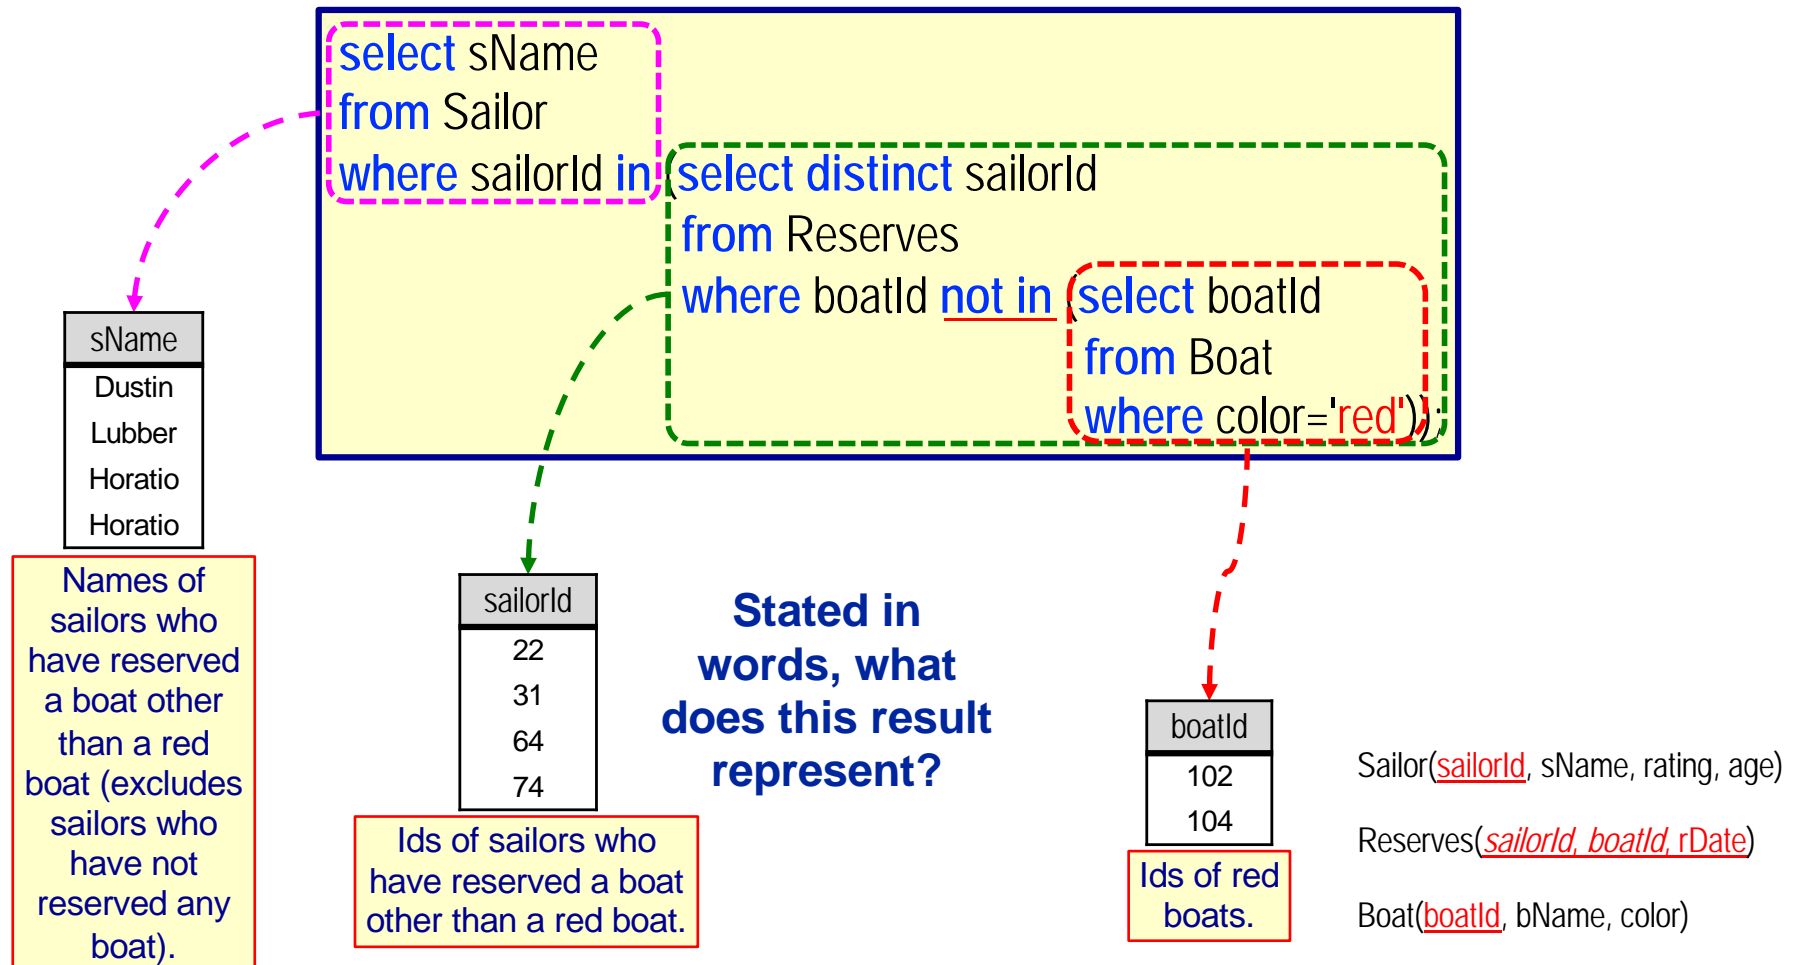
Reserves(sailorId, boatId, rDate)

Boat(boatId, bName, color)



EXERCISE 4 (cont'd)

What if we replace the second **in** with **not in**?



EXERCISE 4 (cont'd)

What if we replace both in's with not in?

Stated in words, what does this result represent?

```
select sName
from Sailor
where sailorId not in (select distinct sailorId
                       from Reserves
                       where boatId not in (select boatId
                                           from Boat
                                           where color='red'));
```

sName
Brutus
Andy
Rusty
Zorba
Art
Bob
Chris

Names of sailors who have reserved only a red boat (i.e., Chris) or have reserved no boat.

sailorId
22
31
64
74

Ids of sailors who have reserved a boat other than a red boat.

boatId
102
104

Ids of red boats.

Sailor(sailorId, sName, rating, age)

Reserves(sailorId, boatId, rDate)

Boat(boatId, bName, color)



EXERCISE 5

Find the ratings and the average age of the ratings where a rating's average age is equal to the minimum average age of all ratings.

 (10, 27)

Is this a
correct
solution?
NO! Why?

```
select rating
from Sailor
where avg(age)=min(select avg(age)
                    from Sailor
                    group by rating);
```

Cannot use “where
avg(age)=” since
avg(age) is not an
attribute of Sailor!

Cannot use “min(...)”.
Illegal SQL!

Is this a
correct
solution?
NO! Why?

```
select rating
from Sailor
group by rating
having age=(select avg(age)
            from Sailor
            group by rating);
```

avgAge
33
44
40
40
35
27

Cannot use “having
age=” since age is
not in the select or
group by clauses.
Illegal SQL!

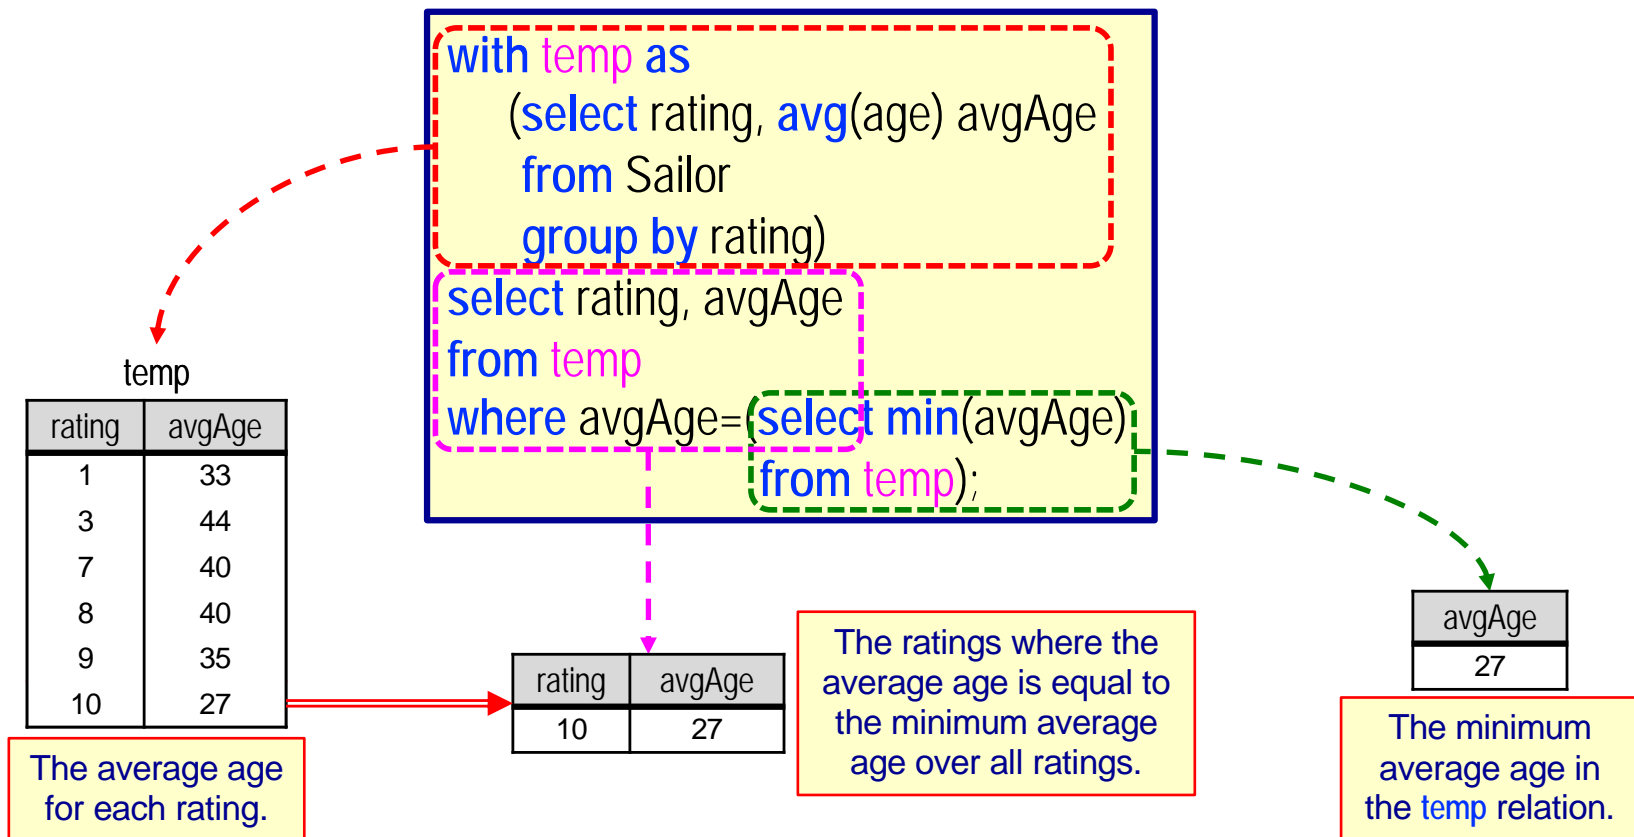
Subquery returns
multiple values.



EXERCISE 5 (CONTD)

Find the ratings and the average age of the ratings where a rating's average age is equal to the minimum average age of all ratings.

👉 (10, 27)




EXERCISE 5 (CONTD)

Find the ratings and the average age of the ratings where a rating's average age is equal to the minimum average age of all ratings.

 (10, 27)

```
select rating, avgAge
from (select rating, avg(age) avgAge
      from Sailor
      group by rating) temp
where avgAge=(select min(avgAge)
              from temp);
```

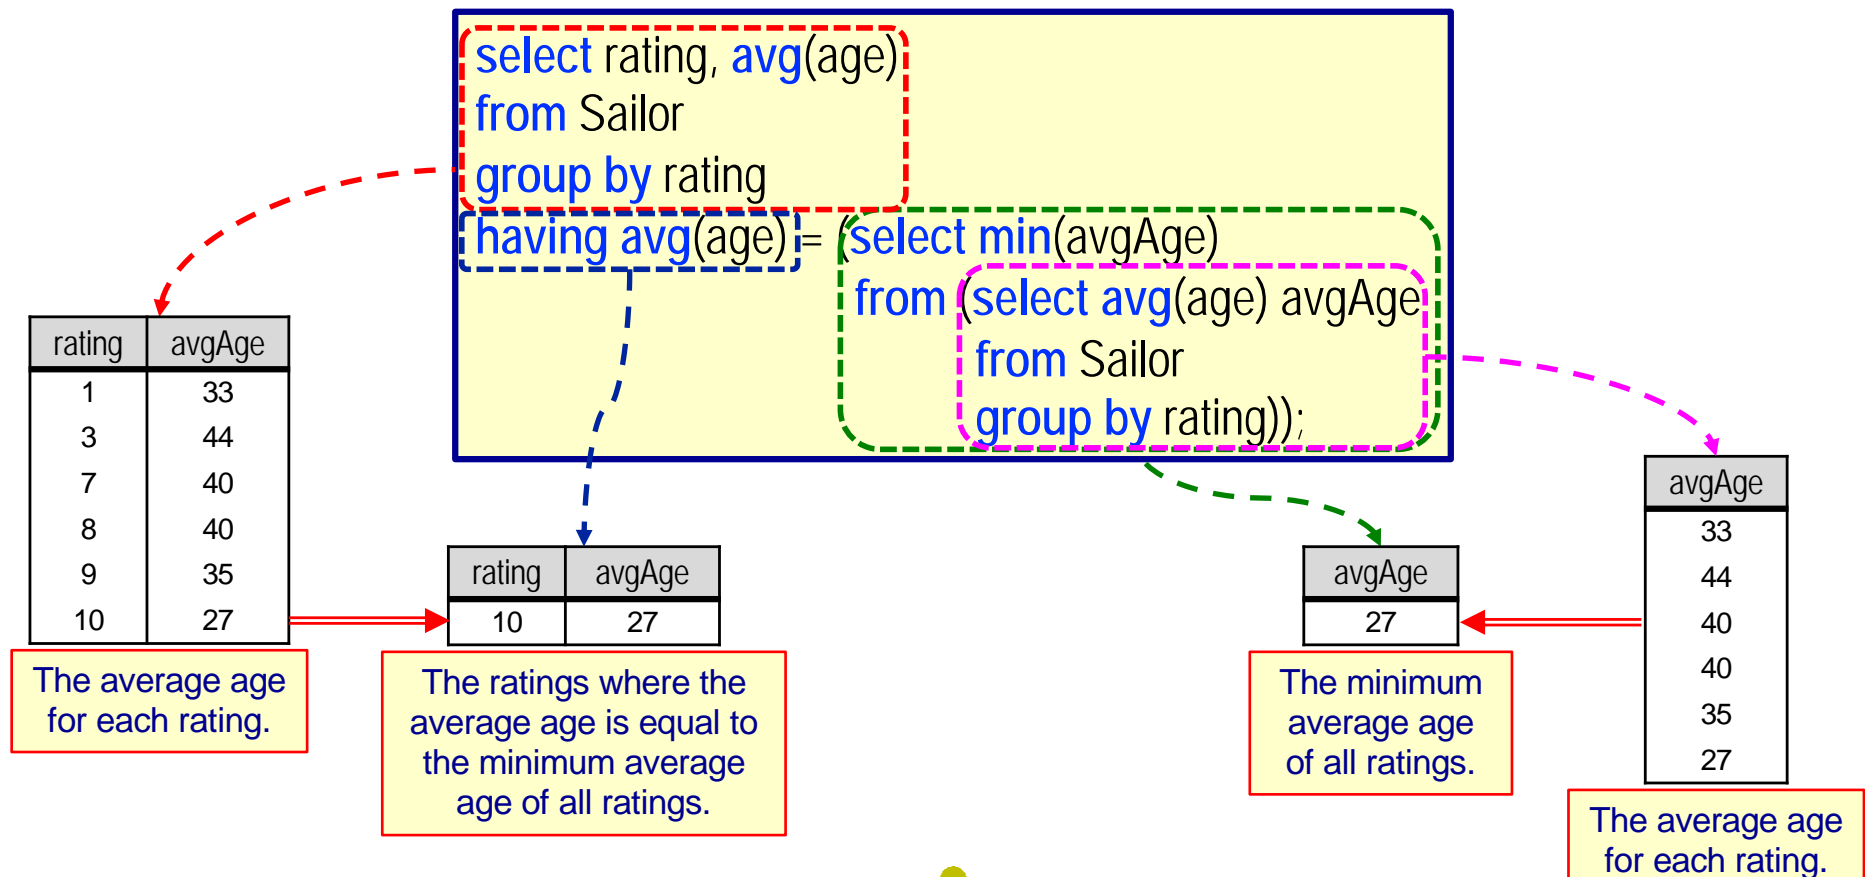
- This query is correct SQL, but will not execute in Oracle.
 - Returns the error “table or view does not exist”.

 **Oracle restricts the scope of the alias `temp` to the outer select.**

EXERCISE 5 (CONTD)

Find the ratings and the average age of the ratings where a rating's average age is equal to the minimum average age of all ratings.

🔍 (10, 27)



EXERCISE 6

Do not create any derived relations.

Find the boat name and number of reservations for each boat.

☞ (Clipper, 3), (Interlake, 2), (Interlake, 3), (Marine, 3), (Serenity, 0)

```
select bName, count(bName) reservationCount
from Boat natural left outer join Reserves
group by bName;
```

bName	reservation Count
Clipper	3
Interlake	5
Marine	3
Serenity	1

The count for Serenity is incorrect.
Interlake should have two separate counts.

What's the problem?

**How about group on boatId, bName;
count boatId?**

```
select bName, count(boatId) reservationCount
from Boat natural left outer join Reserves
group by boatId, bName;
```

bName	reservation Count
Clipper	3
Interlake	2
Interlake	3
Marine	3
Serenity	1

The count for Serenity is still incorrect!

What's the problem?

EXERCISE 6 (CONT'D)

Do not create any derived relations.

Find the boat name and number of reservations for each boat.

☞ (Clipper, 3), (Interlake, 2), (Interlake, 3), (Marine, 3), (Serenity, 0)

```
select bName, count(boatId) reservationCount
from Boat natural left outer join Reserves
group by boatId, bName;
```

Boat natural left outer join Reserves				
boatId	bName	color	sailorId	rDate
101	Interlake	blue	64	05/09/17
101	Interlake	blue	22	10/10/17
102	Interlake	red	22	10/10/17
102	Interlake	red	64	08/09/17
102	Interlake	red	31	10/11/17
103	Clipper	green	22	08/10/17
103	Clipper	green	31	06/11/17
103	Clipper	green	74	08/09/17
104	Marine	red	22	07/10/17
104	Marine	red	99	08/08/17
104	Marine	red	31	12/11/17
105	Serenity	cyan	(null)	(null)

☞ **We need to count sailorId or rDate!**


EXERCISE 6 (CONT'D)

Do not create any derived relations.

Find the boat name and number of reservations for each boat.

☞ (Clipper, 3), (Interlake, 2), (Interlake, 3), (Marine, 3), (Serenity, 0)

```
select bName, count(sailorId) reservationCount
from Boat natural left outer join Reserves
group by boatId, bName;
```



bName	reservation Count
Clipper	3
Interlake	2
Interlake	3
Marine	3
Serenity	0

EXERCISE 7

Do not create any derived relations.

Find the age of the youngest adult sailor (i.e., $\text{age} \geq 18$) for each rating for which there are **at least 2 adult** sailors with the same rating.

☞ (3, 25), (7, 35), (8, 25), (10, 30)

Is this a correct solution?
NO! Why?

```
select rating, min(age)
from Sailor S
group by rating
having 1 < (select count(*)
from Sailor
where S.rating=rating
and age >= 18);
```

sailorId	sName	rating	age
29	Brutus	1	33
85	Art	3	25
95	Bob	3	63
22	Dustin	7	45
64	Horatio	7	35
31	Lubber	8	55
32	Andy	8	25
74	Horatio	9	35
58	Rusty	10	35
71	Zorba	10	16
99	Chris	10	30

Select those groups having both more than one sailor and a sailor whose age is greater than or equal to 18 .

rating	min(age)
3	25
7	35
8	25
10	16

Select the rating and minimum age for each selected group.

X

EXERCISE 7 (CONTD)

Do not create any derived relations.

Find the age of the youngest adult sailor (i.e., $\text{age} \geq 18$) for each rating for which there are **at least 2 adult sailors with the same rating.**

✎ (3, 25), (7, 35), (8, 25), (10, 30)

Sailors whose age is greater than or equal to 18.

sailorId	sName	rating	age
22	Dustin	7	45
29	Brutus	1	33
31	Lubber	8	55
32	Andy	8	25
58	Rusty	10	35
64	Horatio	7	35
74	Horatio	9	35
85	Art	3	25
95	Bob	3	63
99	Chris	10	20

Group the result by rating.

```
select rating, min(age)
from Sailor
where age >= 18
group by rating
having count(*) >= 2
```

sailorId	sName	rating	age
29	Brutus	1	33
85	Art	3	25
95	Bob	3	63
22	Dustin	7	45
64	Horatio	7	35
31	Lubber	8	55
32	Andy	8	25
74	Horatio	9	35
58	Rusty	10	35
99	Chris	10	30

Select the rating and minimum age for each selected group.

rating	min(age)
3	25
7	35
8	25
10	30

Select those groups having at least 2 sailors.

