

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 9 EXERCISES

STRUCTURED QUERY LANGUAGE (SQL)

BOOK STORE RELATIONAL SCHEMA

Book(bookId, title, subject, quantityInStock, price, *authorId*)

Author(authorId, firstName, lastName)

Customer(customerId, firstName, lastName)

BookOrder(orderId, *customerId*, orderYear)

OrderDetails(orderId, bookId, quantity)

Attribute names in
italics are foreign
key attributes.

Assumptions

- Each author has authored at least one book in the store.
- Each book has exactly one Author.
- Each order is made by exactly one customer and has one or more associated tuples in OrderDetails (e.g., one order may contain several different books).

EXERCISE 1

Find all distinct book titles of the author whose last name is Piper.

Relational Algebra

$$\pi_{\text{title}}(\sigma_{\text{lastName}='Piper'}(\text{Author} \text{ JOIN } \text{Book}))$$
$$\pi_{\text{title}}((\sigma_{\text{lastName}='Piper'}\text{Author}) \text{ JOIN } \text{Book})$$

SQL

```
select distinct title
from Book natural join Author
where lastName='Piper';
```



EXERCISE 2

Find the last name and first name of all authors who wrote books in both the subjects of Art and Business.

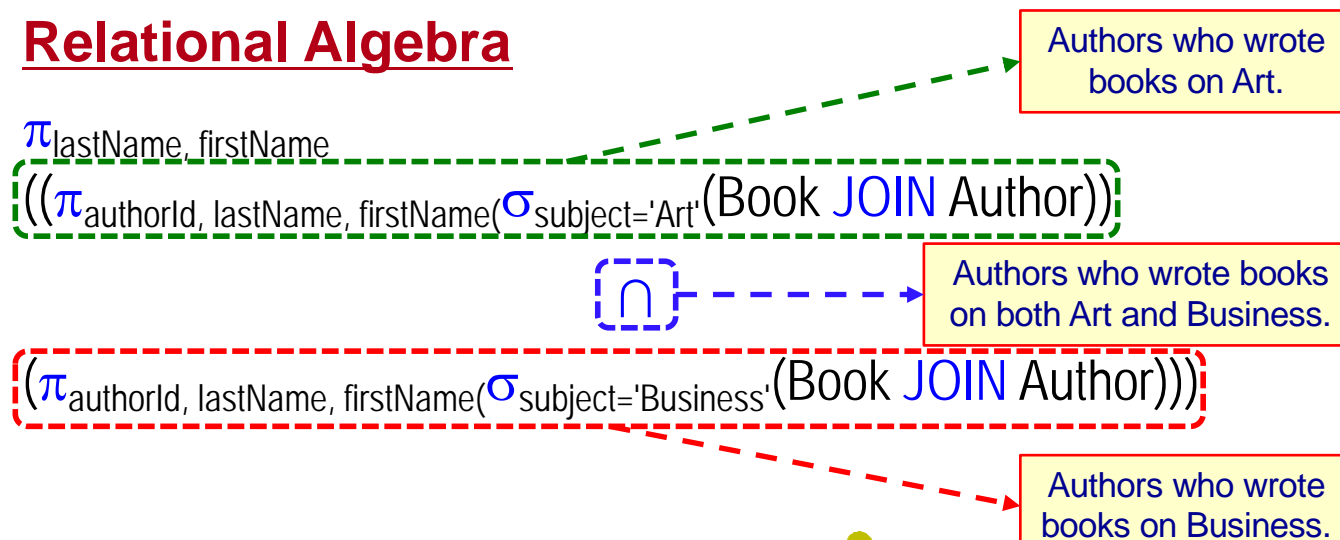
Can we say \Rightarrow where subject='Art' and subject='Business'? **No. Why?**

 **Selects nothing.**

Can we say \Rightarrow where subject='Art' or subject='Business'? **No. Why?**

 **Selects authors who wrote either Art or Business books, but not necessarily on both subjects.**

Relational Algebra



EXERCISE 2 (CONT'D)

Find the last name and first name of all authors who wrote books in both the subjects of Art and Business.

SQL

```
select lastName, firstName
from (select authorId, lastName, firstName
      from Author natural join Book
      where subject='Art'
     intersect
     select authorId, lastName, firstName
      from Author natural join Book
      where subject='Business');
```

Authors who wrote books on Art.

Select only those authors in the Art set who are also in the Business set.

Authors who wrote books on Business.



EXERCISE 2 (CONT'D)

Use *only* set membership

Find the last name and first name of all authors who wrote books in both the subjects of Art and Business.

SQL

```
select lastName, firstName
from Author natural join Book
where subject='Art'
and authorId in (select authorId
from Author natural join Book
where subject='Business');
```

Authors who wrote books on Art (Art set).

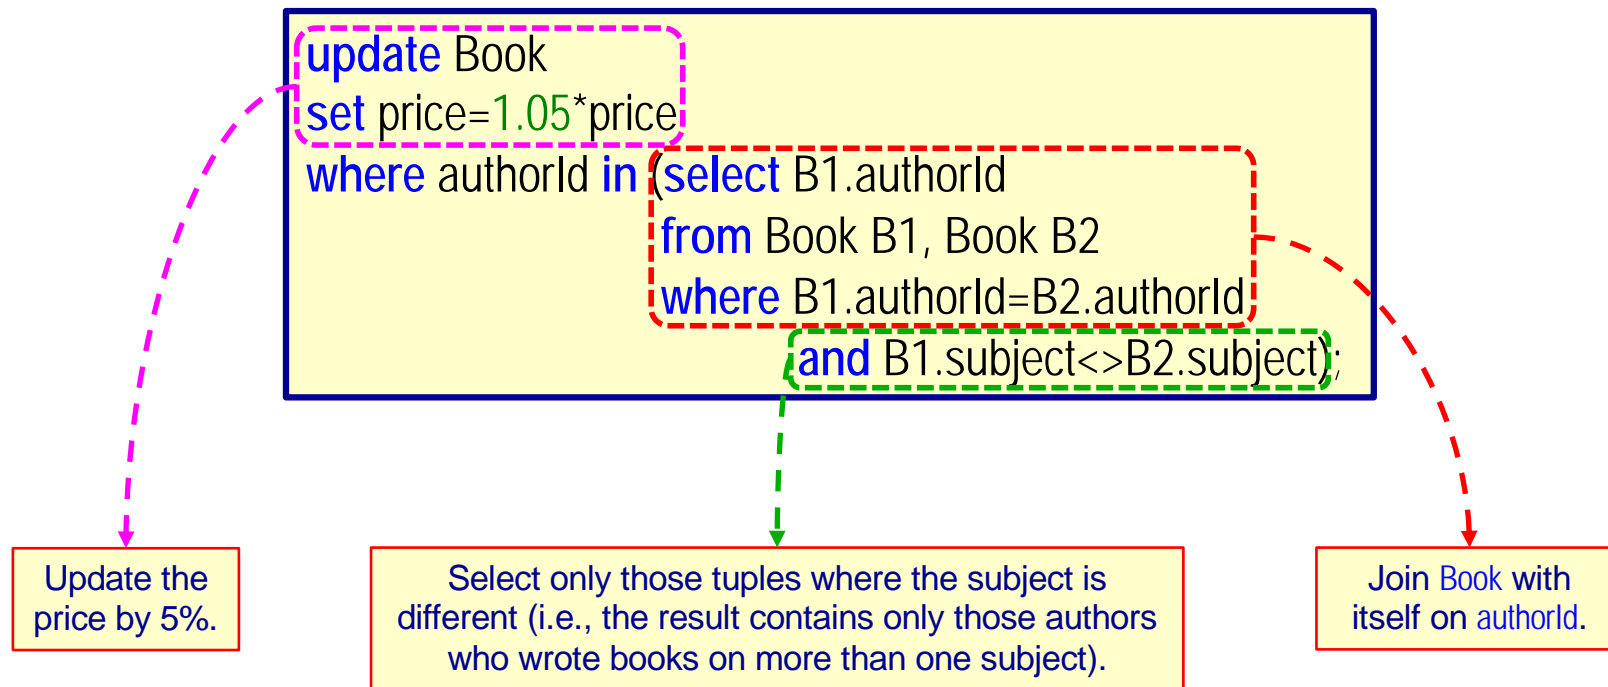
Authors who wrote books on Business (Business set).

Select only those authors in the Art set who are also in the Business set.



EXERCISE 3

For all authors who wrote books on at least two subjects, increase the price of all their books by 5%.



Note: Natural join cannot be used if self join is required. **Why?**



EXERCISE 3 (CONT'D)

Do not
use join

For all authors who wrote books on at least two subjects, increase the price of all their books by 5%.

```
update Book
set price=1.05*price
where authorId in (select authorId
from Book
group by authorId
having count(distinct subject)>=2);
```

Update the
price by 5%.

Authors who wrote books
on more than one subject.



Do not use
subqueries.

EXERCISE 4

Do not create any
derived tables.

Find the last name and first name of all authors who wrote books on exactly ten different subjects.

Is this a
correct
solution?
NO! Why?

```
select lastName, firstName, count(distinct subject) numSubjects
from Author natural join Book
group by authorId, lastName, firstName
having numSubjects=10;
```

Select only those groups
having exactly ten subjects.

Group the result by authorId,
lastName and firstName.

Join Customer and
Book on authorId.

✋ **numSubjects is not an attribute of either Author or Book!**
(It is an attribute only of the **final result**.)



Do not use
subqueries.

EXERCISE 4 (cont'd)

Do not create any
derived tables.

Find the last name and first name of all authors who wrote books on exactly ten different subjects.

Is this a
correct
solution?
NO! Why?

```
select lastName, firstName
from Author A
where 10 = (select count(*)
           from Book
           where A.authorId=Book.authorId
           group by A.authorId);
```

Selects last name of author
if the `where` clause is true.

Counts the number of books written by each author and
returns true if the author wrote exactly ten books.

👉 **Selects authors who wrote exactly ten books.**
(But the subject could be the same!)

👉 **How to fix this?**

Change `select count(*)` **to** `select count(distinct subject)`.

**Do not use
subqueries.**

EXERCISE 4 (cont'd)

**Do not create any
derived tables.**

Find the last name and first name of all authors who wrote books on **exactly ten different subjects.**

Is this a
correct
solution?
YES!

```
select lastName, firstName
from Author natural join Book
group by authorId, lastName, firstName
having count(distinct subject)=10;
```

Join Author and
Book on authorId.

Group the result by
authorId and lastName.

Select only those groups having
exactly ten different subjects.

Is authorId needed in the group by clause?

👉 **YES**, otherwise the count for two different authors with the same last and first name will be incorrect resulting in an incorrect result.



Do not use
subqueries.

EXERCISE 4 (cont'd)

Do not create any
derived tables.

Find the last name and first name of all authors who wrote books on exactly ten different subjects.

Is this a
correct
solution?

YES!

(But should
not use
subquery!)

```
select lastName, firstName
from Author
where authorId in (select authorId
                   from Book
                   group by authorId
                   having count(distinct subject)=10)
```

Last and first names of authors who
are in the result of the inner select.

Ids of authors who have written
books on ten different subjects.



Do not use
subqueries.

EXERCISE 5

Do not create any
derived tables.

For each customer who made more than 10 orders in 2018, find the customer id, last name and the number of orders in 2018.

```
select customerId, lastName, count(*)  
from Customer natural join BookOrder  
where orderYear='2018'  
group by customerId, lastName  
having count(*)>10
```

Select only those groups having
more than 10 orders.

Group the result by
customerId and lastName.

Join Customer and Order on
customerId for orders in 2018.

Are both customerId and lastName needed in the group by clause?

👉 **YES**, since they are both present in the select clause.

EXERCISE 6

Find the customer id, last name and total quantity ordered for those customers who ordered the **largest total quantity** of books.

```
with temp as
  (select customerId, lastName, sum(quantity) totalQuantity
   from Customer natural join BookOrder natural join OrderDetails
   group by customerId, lastName)
select customerId, lastName, totalQuantity
from temp
where totalQuantity = (select max(totalQuantity)
                      from temp);
```

The customers who ordered the largest total quantity of books.

Customer(customerId, firstName, lastName)

The largest total quantity of books ordered (from the temp relation).

BookOrder(orderId, customerId, orderYear)

The total quantity of books ordered by each customer.

OrderDetails(orderId, bookId, quantity)

EXERCISE 6 (cont'd)

Find the customer id, last name and total quantity ordered for those customers who ordered the **largest total quantity** of books.

```
select customerId, lastName, sum(quantity) totalQuantity
from Customer natural join BookOrder natural join OrderDetails
group by customerId, lastName
having sum(quantity) = (select max(sum(quantity))
                       from BookOrder natural join OrderDetails
                       group by customerId);
```

The customers who ordered the largest total quantity of books.

Select those customers whose total quantity is the largest.

The total quantity of books ordered by each customer.

Customer(customerId, firstName, lastName)

BookOrder(orderId, customerId, orderYear)

OrderDetails(orderId, bookId, quantity)

EXERCISE 6 (cont'd)

Find the customer id, last name and total quantity ordered for those customers who ordered the **largest total quantity** of books.

```
select customerId, lastName, sum(quantity) totalQuantity
from Customer natural join BookOrder natural join OrderDetails
group by customerId, lastName
having sum(quantity) >= all (select sum(quantity)
                             from BookOrder natural join OrderDetails
                             group by customerId);
```

The customers who ordered the largest total quantity of books.

Select those customers whose total quantity is the largest.

The total quantity of books ordered by each customer.

Customer(customerId, firstName, lastName)

BookOrder(orderId, customerId, orderYear)

OrderDetails(orderId, bookId, quantity)

