

# COMP 3311

# DATABASE MANAGEMENT

# SYSTEMS

## LECTURE 12 EXERCISES

## INDEXING: INTRODUCTION

$$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$$
$$\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$$

Film records: 30,000  
Actor records: 100,000  
Page size: 512 bytes  
Pointer size: 6 bytes

## EXERCISE 1

A movie database has the following files and sizes of each field:

84 bytes/record Film(title: 40 bytes, director: 20 bytes, releaseYear: 4 bytes, company: 20 bytes)

28 bytes/record Actor(id: 4 bytes, name: 20 bytes, dateOfBirth: 4 bytes)

There exist 30,000 films in the database and 100,000 actors.  
Each page is 512 bytes and each pointer is 6 bytes.

a) What is the blocking factor  $bf_F$  for the Film file and  $bf_A$  for the Actor file?

$$bf_F: \lfloor 512 \text{ bytes per page} / 84 \text{ bytes per Film record} \rfloor$$
$$= 6 \text{ records/page}$$

$$bf_A: \lfloor 512 \text{ bytes per page} / 28 \text{ bytes per Actor record} \rfloor$$
$$= 18 \text{ records/page}$$

$$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$$
$$\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$$

Film records: 30,000  
Actor records: 100,000  
Page size: 512 bytes  
Pointer size: 6 bytes  
Film record size: 84 bytes;  $bf_F = 6$   
Actor record size: 28 bytes;  $bf_A = 18$

## EXERCISE 1 (CONTD)

b) Assuming that the Film file is **ordered on title** and there is **no index**, what is the cost (in terms of page I/Os) for:

i. Finding the film with title "**Titanic**"?

**Pages needed:**  $\lceil (30,000 \text{ Film records} / 6 \text{ Film records per page}) \rceil$   
= 5000 pages

**Search cost:**  $\lceil \log_2 5000 \rceil = \underline{13}$  page I/Os (binary search)

ii. Finding all the films directed by "**John Woo**"?

**Search cost:** 5000 page I/Os

A sequential scan is needed since the file is not ordered based on director.

$$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$$
$$\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$$

Film records: 30,000  
Actor records: 100,000  
Page size: 512 bytes  
Pointer size: 6 bytes  
Film record size: 84 bytes;  $bf_F = 6$   
Actor record size: 28 bytes;  $bf_A = 18$

## EXERCISE 2

Assume that the Actor file is **ordered on the name** and we want to create an **ordered index on id (4 bytes)** where each index entry has the form **<id, pointer>**.

a) What is  $bf_{A_{index}}$  if the index is single-level?

$$bf_{A_{index}}: \lfloor 512 \text{ bytes per page} / (4 + 6) \text{ bytes per index entry} \rfloor = \underline{51}$$

b) How many index entries are needed?

**Index entries:** 100,000 index entries **Why?**

A dense index is needed (i.e., one entry per Actor record) since the file is ordered on name, not on id.

c) How many pages are required for the Actor index entries?

$$\text{Pages needed: } \lceil 100,000 \text{ Actor records} / 51 \text{ index entries per page} \rceil = \underline{1961}$$

$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$   
 $\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$

Film records: 30,000  
 Actor records: 100,000  
 Page size: 512 bytes  
 Pointer size: 6 bytes  
 Film record size: 84 bytes;  $bf_F = 6$   
 Actor record size: 28 bytes;  $bf_A = 18$

## EXERCISE 2 (CONT'D)

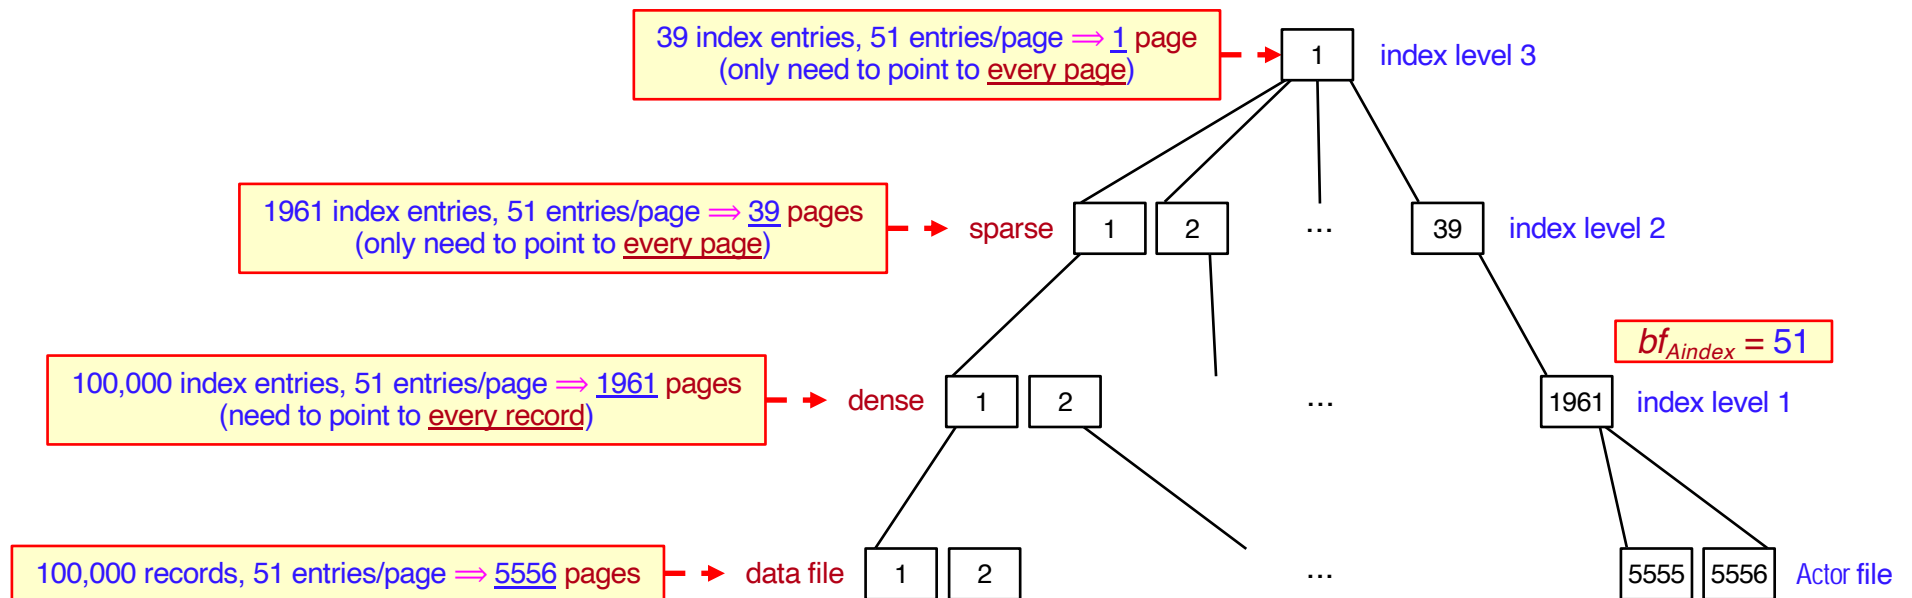
- d) What is the cost of retrieval based on a single id value using the Actor index (e.g., “Find actor with id 100”)?

Actor file ordered on name.

**Search cost:**  $\lceil \log_2 1961 \rceil + 1 = 12$  page I/Os

- e) If we convert the single-level index into a multi-level index, how many levels are needed (assuming full pages)?

**Levels needed: 3**



$$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$$

$$\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$$

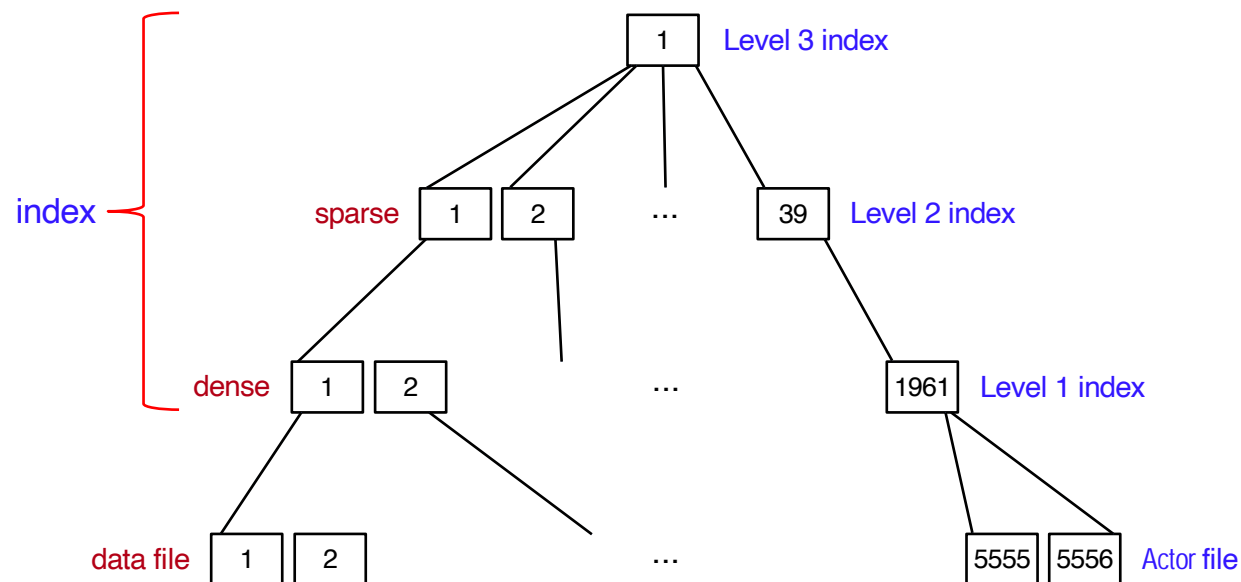
Film records: 30,000  
 Actor records: 100,000  
 Page size: 512 bytes  
 Pointer size: 6 bytes  
 Film record size: 84 bytes;  $bf_F = 6$   
 Actor record size: 28 bytes;  $bf_A = 18$

## EXERCISE 2 (CONT'D)

f) Using the multi-level index, what is the cost of answering the query “Find the actor with id 100”?

Search cost: 4 page I/Os **Why?**

3 page I/Os for the index plus 1 page I/O to retrieve the record.



## EXERCISE 3

A company database has the following file and sizes of each field

Employee(employeeId: 6 bytes, employeeName: 10 bytes, departmentId: 4 bytes)

where departmentId is the id of the department where the employee works.

There are 100,000 employee records and 1,000 departments (each department has 100 employees).

A page is 1,000 bytes and a pointer is 4 bytes.

Assume that the file is ordered on departmentId and there is no index.

$$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$$
$$\# \text{ pages} = \lceil \# \text{ records} / bf_r \rceil$$

Employee records: 100,000  
Departments: 1,000  
Page size: 1000 bytes  
Record size: 20 bytes  
Pointer size: 4 bytes

## EXERCISE 3 (CONT'D)

a) What is the blocking factor for the Employee file?

$$bf_{\text{Employee}} = \lfloor 1000 \text{ bytes per page} / 20 \text{ bytes per record} \rfloor = 50$$

b) How many pages are needed to store the Employee file?

$$\text{Pages needed} = \lceil 100,000 \text{ records} / 50 \text{ records per page} \rceil = 2000$$



## EXERCISE 3 (CONTD)

Employee records: 100,000  
Departments: 1,000  
Page size: 1000 bytes  
Record size: 20 bytes  
Pointer size: 4 bytes  
Pages: 2000

- c) What is the cost (in terms of page I/Os) for retrieving the records of all employees working in a department with a given **departmentId** (e.g., **departmentId** = 64)?

**Search cost:** 12 page I/Os **Why?**

Finding the first record requires  $\lceil \log_2 2000 \rceil = 11$  page I/Os plus 1 more page access to search the remaining records. Since each department has 100 employees and a page can hold 50 records, these records are distributed in at least 2 pages.

The answer  $\lceil \log_2 2000 \rceil$  plus 2 can also be considered correct if a department's employee records are distributed across three pages (e.g., the first page contains 25 records – the second one 50 – and the third one 25). **Total page I/Os** = 13.

## EXERCISE 4

Employee records: 100,000  
Departments: 1,000  
Page size: 1000 bytes  
Record size: 20 bytes  
Pointer size: 4 bytes  
Pages: 2000

For the Employee file of Exercise 3, assume that we add a **single-level ordered index on employeeld** (6 bytes) where each entry has the form  $\langle \text{employeeld}, \text{pointer} \rangle$  and **the number of pointers is the same as the number of search keys**.

**a) How many index entries are needed? (Explain briefly.)**

**Index entries needed: 100,000 Why?**

Since the file is ordered on departmentId, the **index is secondary** and therefore it **must be dense**. Thus, one index entry is needed for each employee.

$bf = \lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$   
 $\# \text{ pages} = \lceil \# \text{ records} / bf \rceil$

Employee records: 100,000  
Departments: 1,000  
Page size: 1000 bytes  
Record size: 20 bytes  
Pointer size: 4 bytes

## EXERCISE 4 (CONT'D)

b) How many pages are required for these index entries?

$$bf_{\text{employeeidindex}}: \lfloor 1000 \text{ bytes per page} / 10 \text{ bytes per index entry} \rfloor = 100$$

$$\text{Index pages: } \lceil 100,000 \text{ records} / 100 \text{ index entries per page} \rceil = 1000$$



c) What is the cost of retrieving the record of an employee with a given **employeeid**?

$$\text{Search cost: } \lceil \log_2 1000 \rceil + 1 = 10 + 1 = 11 \text{ page I/Os}$$

## EXERCISE 4 (CONT'D)

Employee records: 100,000  
Departments: 1,000  
Page size: 1000 bytes  
Record size: 20 bytes  
Pointer size: 4 bytes

d) If we convert the single-level index into a **multi-level index**, how many levels are needed (assuming full pages)?

**Levels needed: 3 Why?**

At the next level we index 1000 index pages (i.e., the index contains  $\lceil 1000 \text{ pages} / 100 \text{ index entries per page} \rceil = 10 \text{ pages}$ . We also need an additional top level with 1 page.

