

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 14 EXERCISES

INDEXING:

HASH INDEX & BITMAP INDEX

EXERCISE 1

A company database has the following file and sizes of each field

Employee(employeeId: 6 bytes, employeeName: 10 bytes, departmentId: 4 bytes)

where departmentId is the id of the department where the employee works.

There are 100,000 employee records and 1,000 departments (each department has 100 employees).

A page is 1,000 bytes and a pointer is 4 bytes.

Assume that the file is sorted on departmentId and there is no index.

We want to build a hash index on employeeId on the above file where each entry has the form <employeeId, pointer>.

EXERCISE 1 (CONTD)

Employee records: 100,000
Departments: 1,000
Page size: 1000 bytes
Record size: 20 bytes
Index entry size: 10 bytes

a) How many index entries are needed?

Index entries needed: 100,000 **Why?**

Since a hash index is always secondary, it must be dense.

b) How many pages are required for these index entries (assuming full pages)?

Pages $\lceil 100,000 \text{ index entries} / 100 \text{ index entries per page} \rceil$
needed: = 1000

c) Using the hash index, what is the cost of retrieving the record of an employee with a given **employeeid**, assuming that there are **no overflow pages**?

Search cost: 2 page I/Os **Why?**

1 access for the hash index + 1 access to retrieve the record.

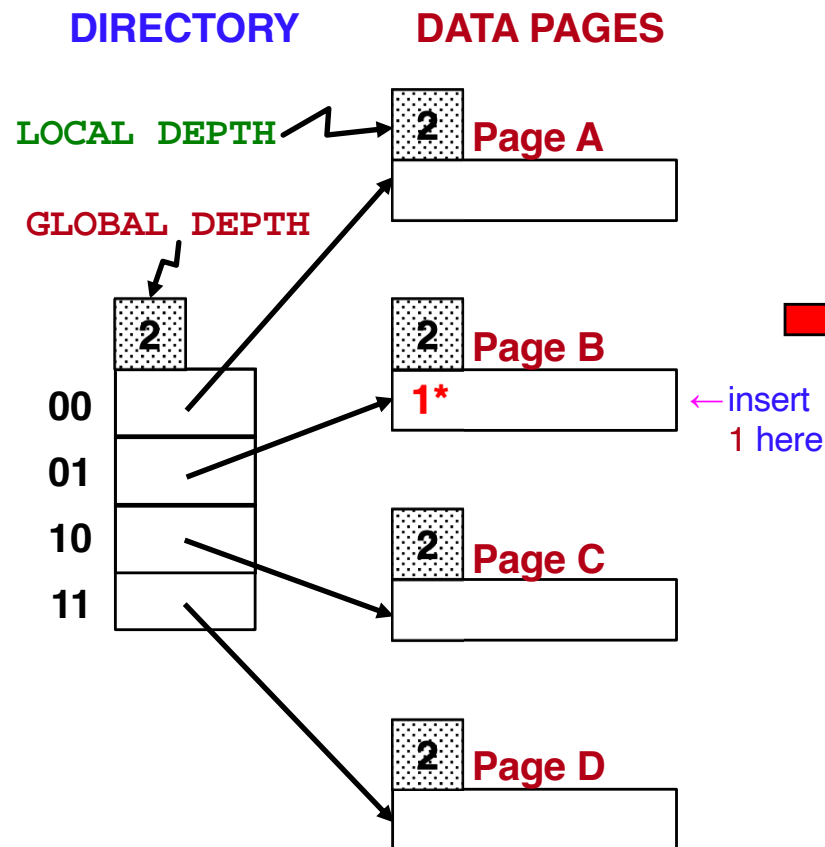
EXERCISE 2

Using the template below, construct a file that contains records with the given search-key values using extendable hashing. Use the hash function $h(x) = x \bmod 8$ and insert records one at a time in order into an empty file. Assume data pages can hold 4 records.

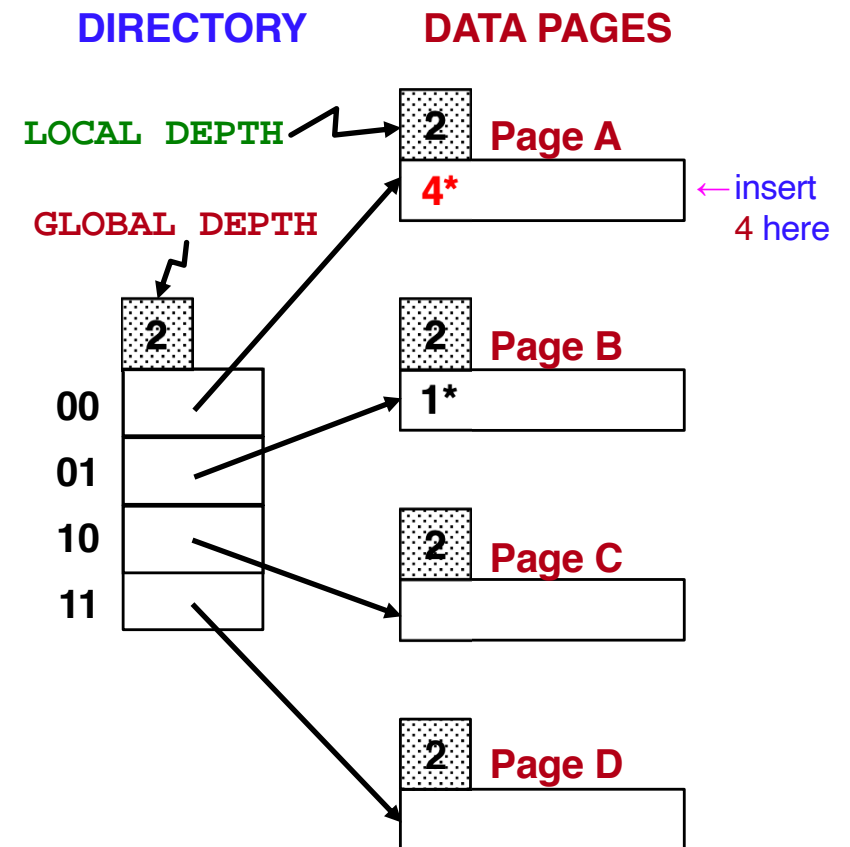
key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

EXERCISE 2 (CONT'D)

Insert: 1



Insert: 4

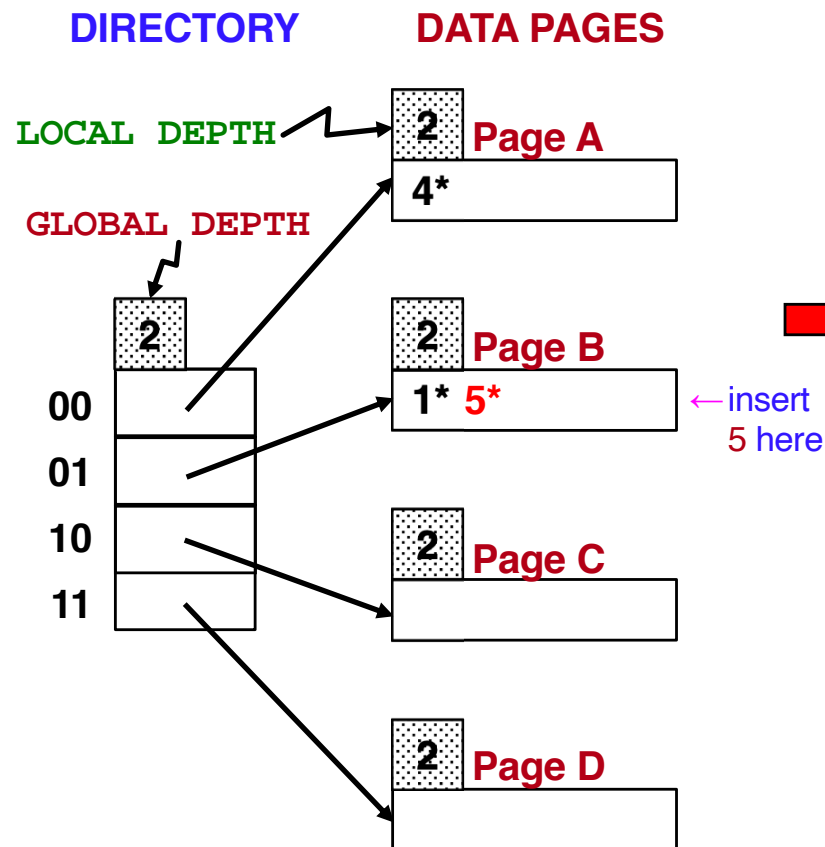


key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

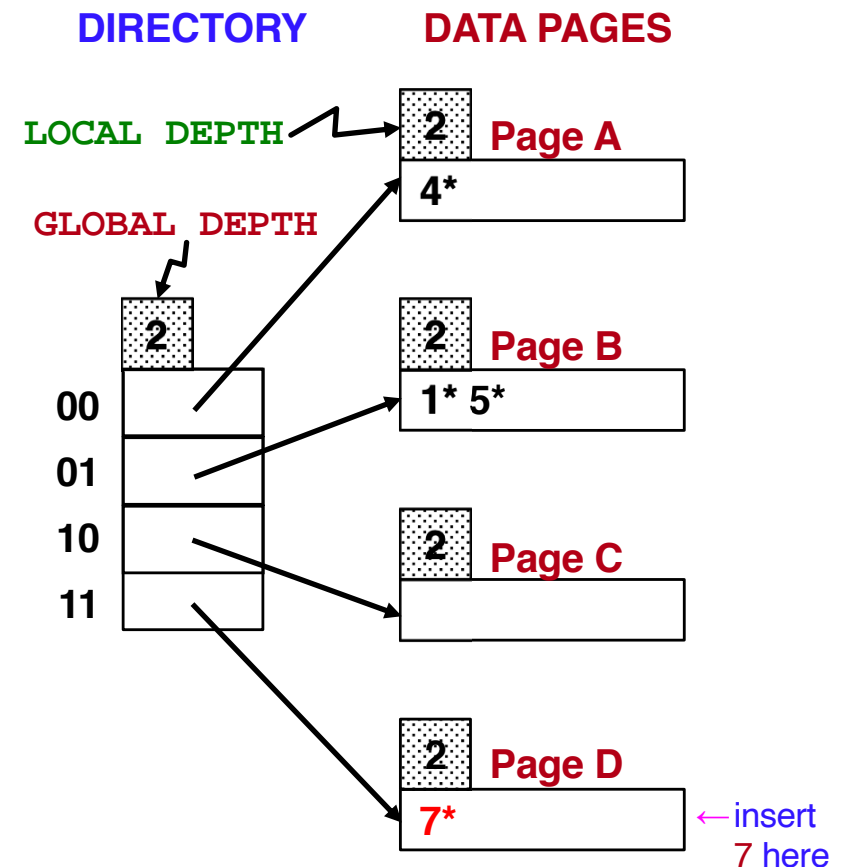


EXERCISE 2 (CONT'D)

Insert: 5



Insert: 7

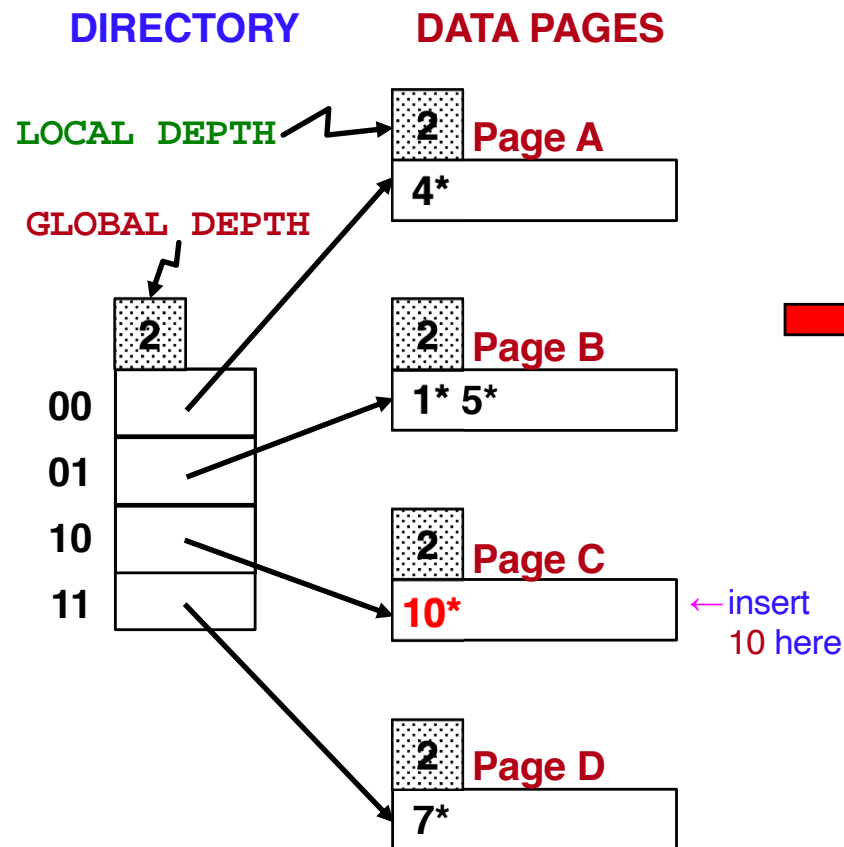


key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

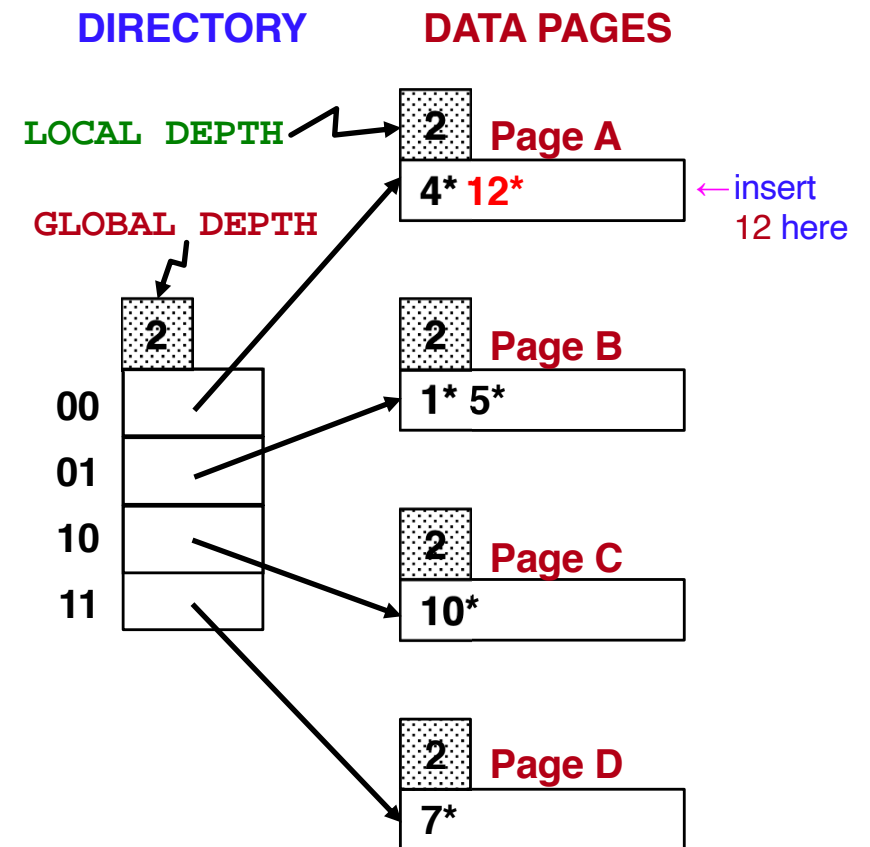


EXERCISE 2 (CONT'D)

Insert: 10



Insert: 12

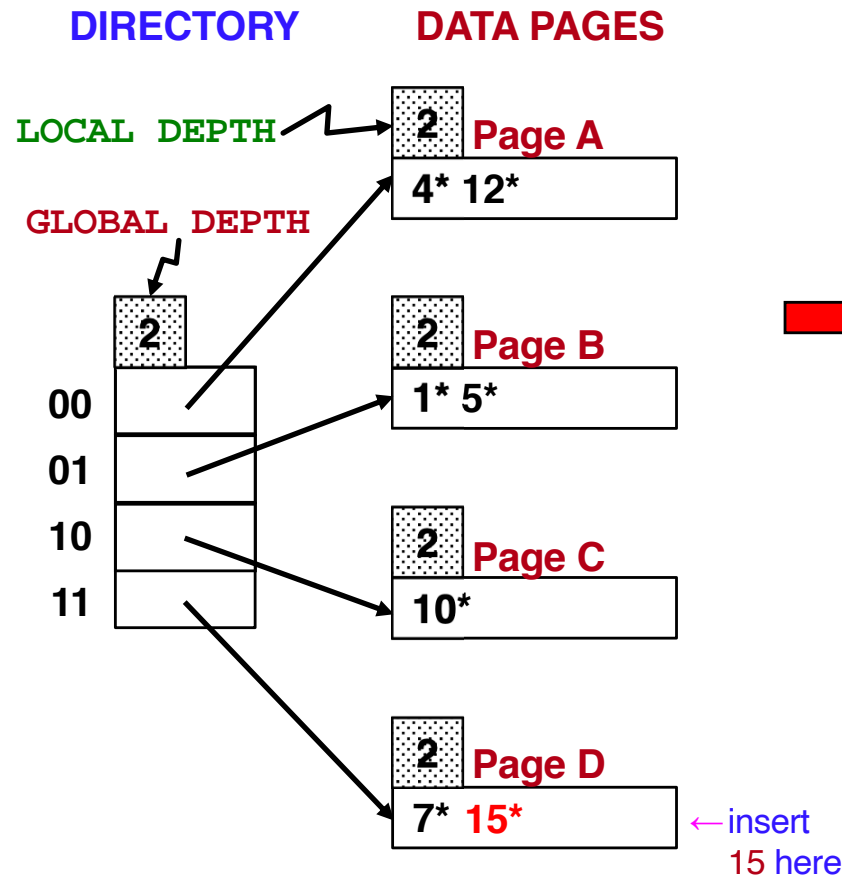


key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

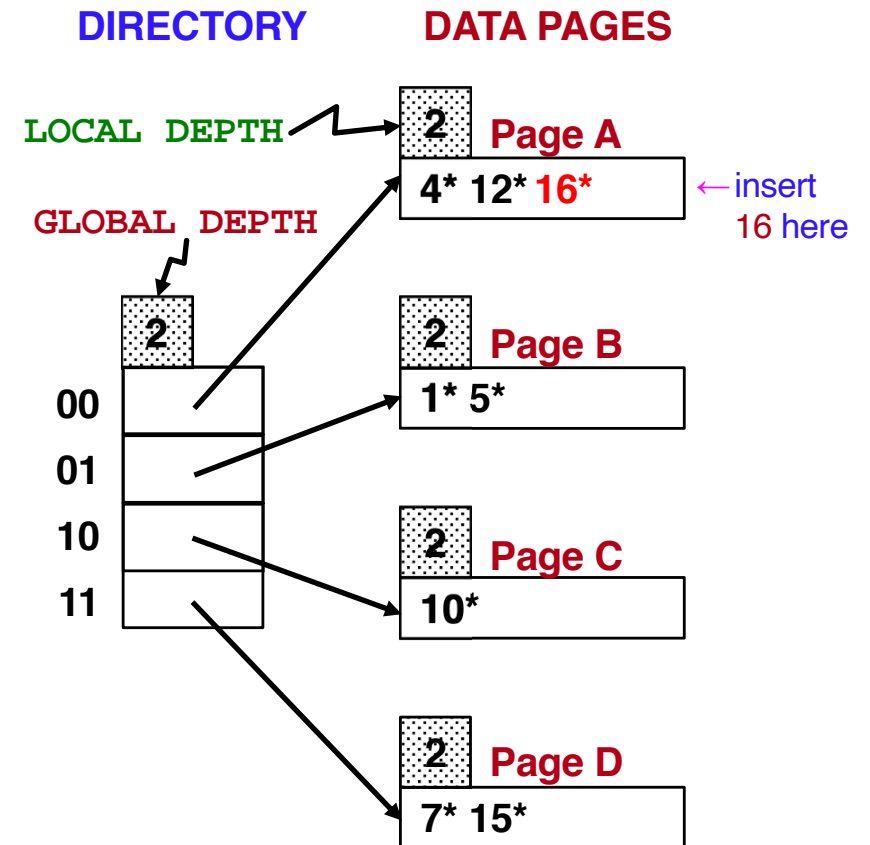


EXERCISE 2 (CONT'D)

Insert: 15



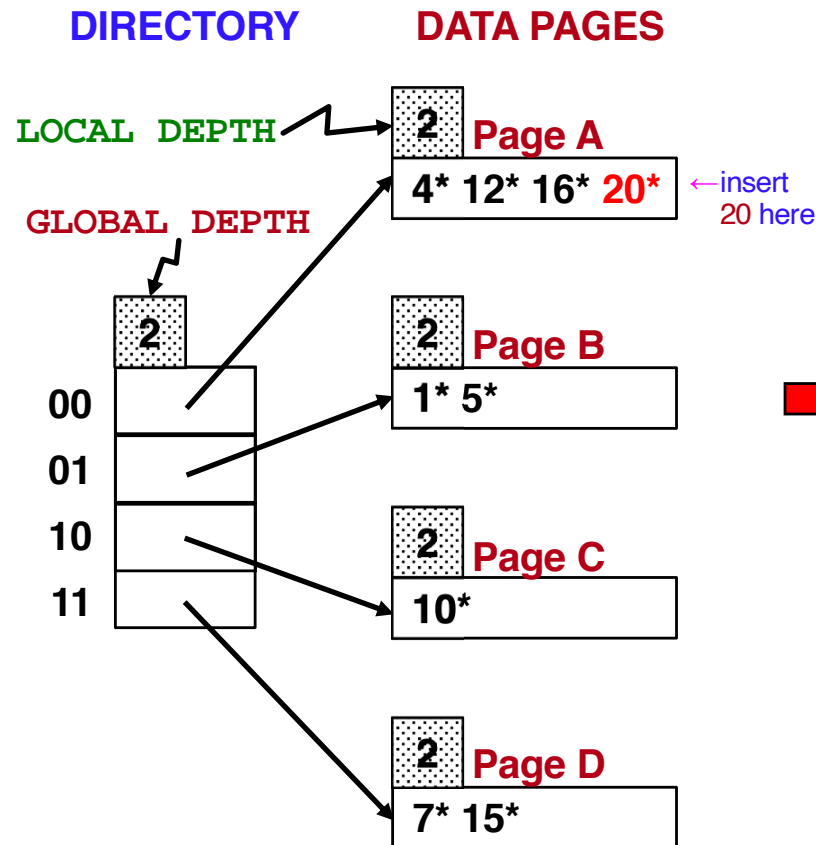
Insert: 16



key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

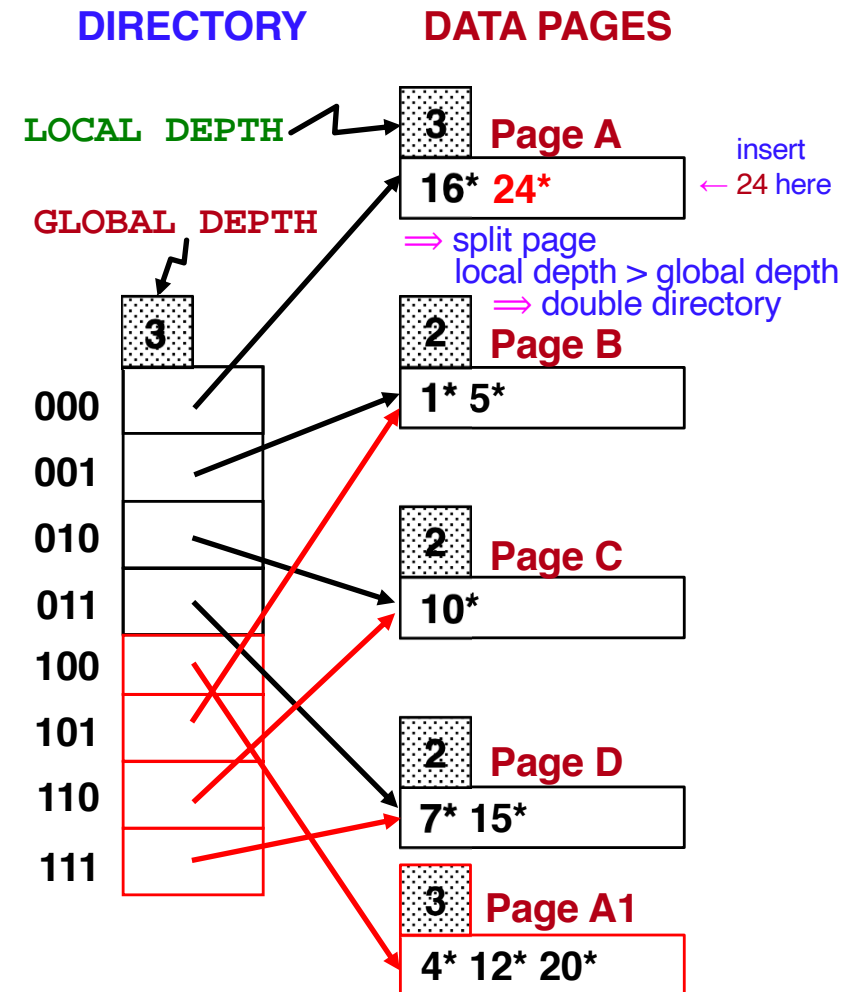
EXERCISE 2 (CONT'D)

Insert: 20



key value	1	4	5	7	10	12	15	16	20	24
$h(x)$	1	4	5	7	2	4	7	0	4	0
binary value	001	100	101	111	010	100	111	000	100	000

Insert: 24




EXERCISE 3

A global e-commerce website maintains a Customer file with attributes customerId, name, address, email and country. The file is organized as a hash file on customerId. There is also a secondary hash index on country. For each country there is only one index entry in the hash index. Assume that on average there are 9,000 customers for each country, there are 90 countries and that a page can hold 100 record pointers. How many page I/Os are required to retrieve the records of all the customers for a given country using the hash index on country?

EXERCISE 3 (CONTD)

From the problem description, the secondary hash index uses a level of indirection to store all the record pointers to the records for a given country as shown in the figure.

Consequently, the number of page I/Os required to retrieve the records for all the customers in a given county is:



Hash index: 1 page I/O

Indirection level: $\lceil 9000 / 100 \rceil = 90$ page I/Os

Record retrieval: 9000 page I/Os (one for each indirection pointer)

Total page I/Os: $1 + 90 + 9000 = \underline{9091}$


EXERCISE 4

Construct a file that contains records with the given search-key values using extendable hashing. Use the hash function $h(x) = x \bmod 8$ and insert records one at a time in order into an empty file. Assume data pages can hold 3 records.

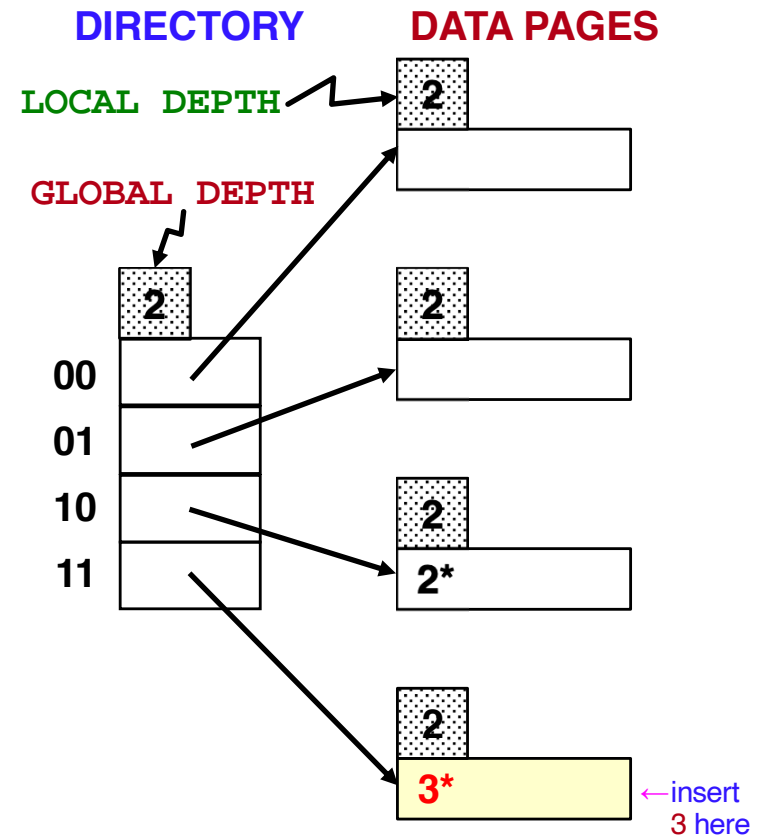
key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111

EXERCISE 4 (CONTD)

Insert: 2



Insert: 3

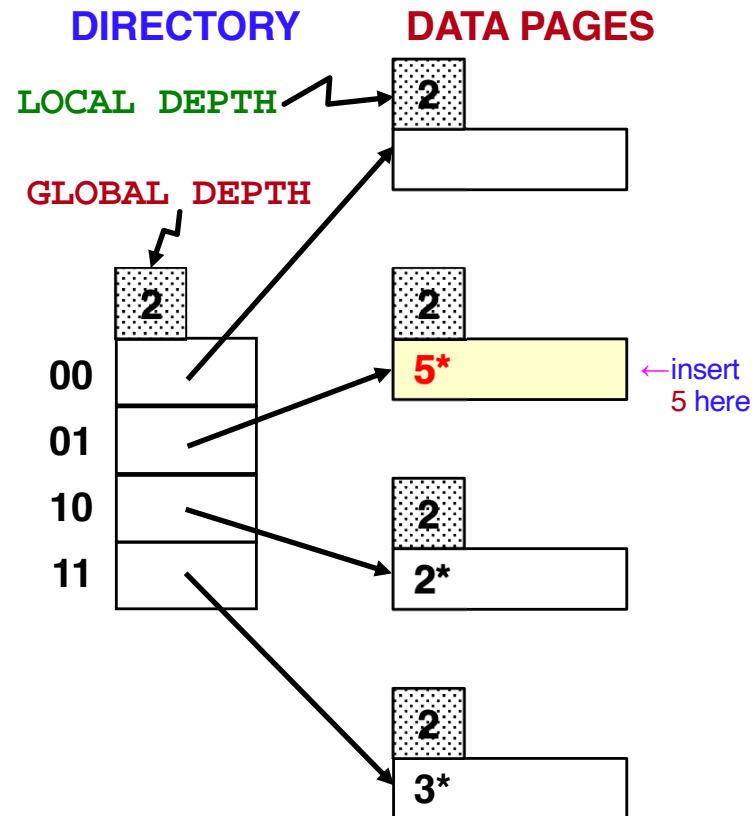


key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111




EXERCISE 4 (CONTD)

Insert: 5



Insert: 7

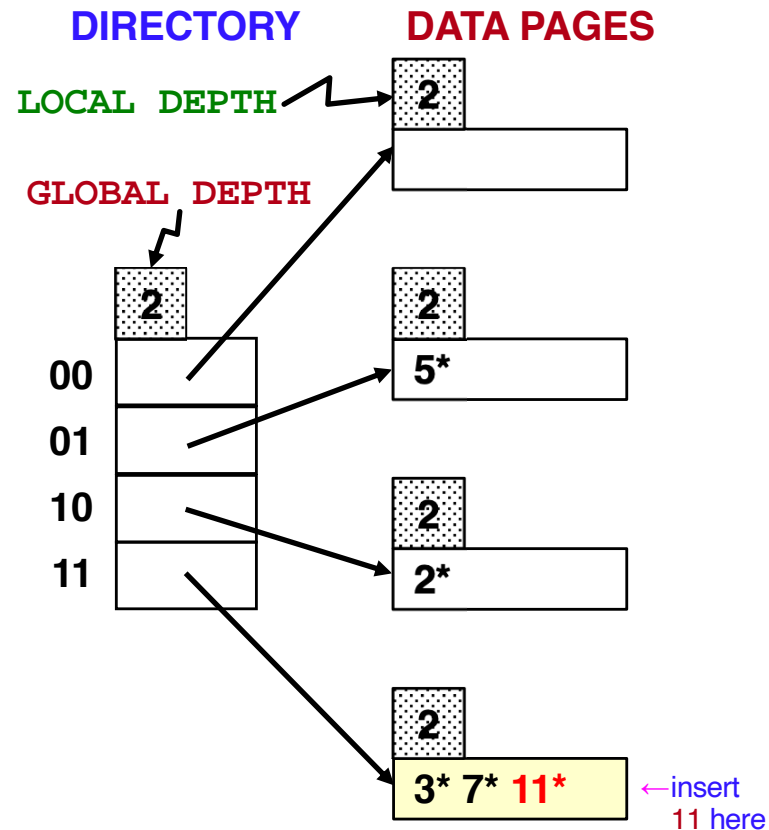


key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111

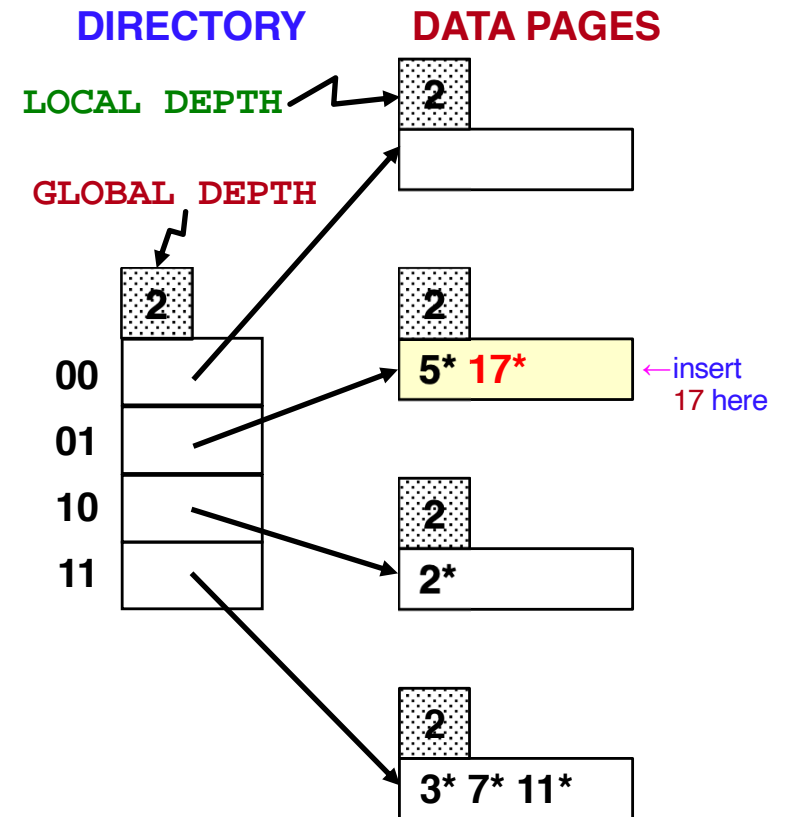


EXERCISE 4 (CONTD)

Insert: 11



Insert: 17

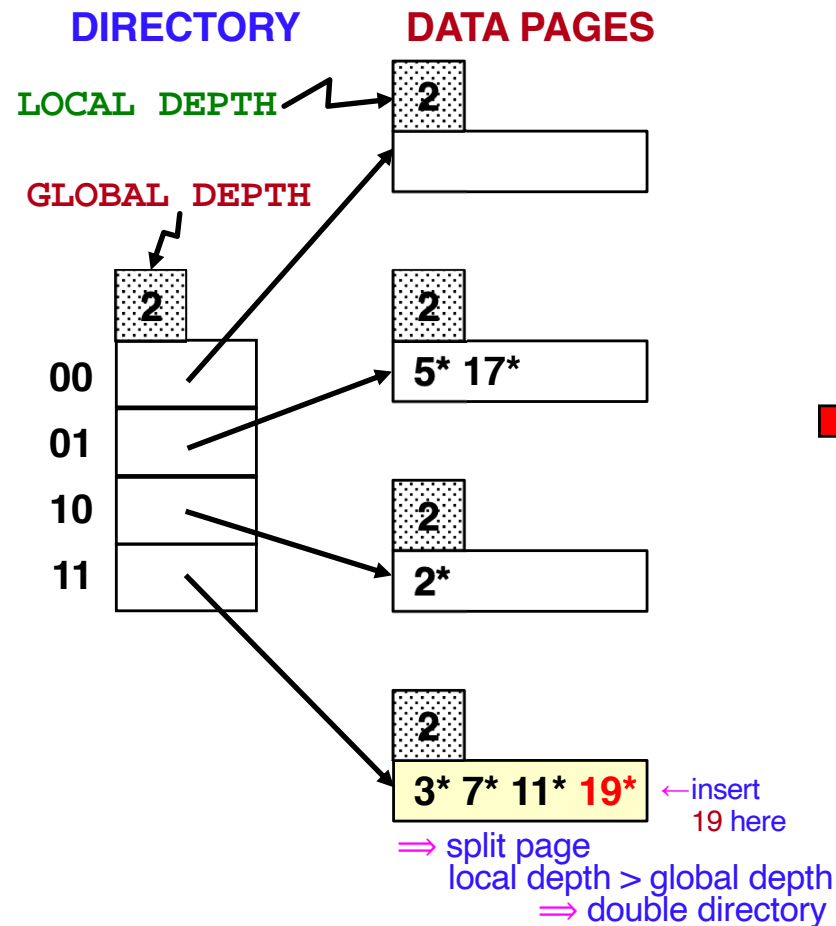


key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111




EXERCISE 4 (CONTD)

Insert: 19




key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111

Insert: 23




EXERCISE 4 (CONTD)

Insert: 29




Insert: 31



key value	2	3	5	7	11	17	19	23	29	31
$h(x)$	2	3	5	7	3	1	3	7	5	7
binary value	010	011	101	111	011	001	011	111	101	111



EXERCISE 4 (CONTD)



EXERCISE 5

Given the Customer relation and only the two bitmap indexes on gender and rating shown below, explain how you would use the bitmap indexes to answer the following queries? If the bitmap indexes are not useful, explain why.

Do not calculate the result of a query, but explain how to obtain the result using the bitmaps.

gender index

male	female
1	0
1	0
0	1
1	0

Customer

id	name	gender	rating
112	Joe	m	2
115	Ram	m	5
119	Sue	f	5
112	Woo	m	1

rating index

1	2	3	4	5
0	1	0	0	0
0	0	0	0	1
0	0	0	0	1
1	0	0	0	0

EXERCISE 5 (CONTD)

a) How many customers with a rating less than 3 are male?

- or the rating 1 and 2 bitmaps.
- and the result with the male bitmap.
- count the number of 1 bits in the result.

b) What percentage of customers are male?

- count the number of 1 bits in the male bitmap.
- divide the result by the length of the bitmap.
- convert to a percentage.

gender index

male	female
1	0
1	0
0	1
1	0

Customer

id	name	gender	rating
112	Joe	m	2
115	Ram	m	5
119	Sue	f	5
112	Woo	m	1

rating index

1	2	3	4	5
0	1	0	0	0
0	0	0	0	1
0	0	0	0	1
1	0	0	0	0

EXERCISE 5 (CONTD)

c) How many customer there are?

count the total number of bits in any bitmap *or*
use the length of any bitmap.

d) How many customer are named Woo?

The bitmaps are not useful for this query.

gender index

male	female
1	0
1	0
0	1
1	0

Customer

id	name	gender	rating
112	Joe	m	2
115	Ram	m	5
119	Sue	f	5
112	Woo	m	1

rating index

1	2	3	4	5
0	1	0	0	0
0	0	0	0	1
0	0	0	0	1
1	0	0	0	0