

# COMP 3311

# DATABASE MANAGEMENT

# SYSTEMS

## LECTURE 20 EXERCISES


## TRANSACTIONS

# EXERCISE 1

Indicate which of the following schedules involving  $T_1$  and  $T_2$  is **serial**, **serializable** or **not serializable**.  $r_i$  denotes a read (of transaction  $T_i$ ) and  $w_i$  is a write (of transaction  $T_i$ ).

a)  $r_1(A) w_1(A) r_2(A) w_2(A)$

**Serial**




$T_1$	$T_2$
read(A) write(A)	read(A) write(A)

b)  $r_1(A) r_2(A) w_1(A) w_2(B)$

**Serializable**

What is the equivalent serial schedule?

$T_2 T_1$




$T_1$	$T_2$
read(A) write(A)	read(A) write(B)

## EXERCISE 1 (CONTD)

c)  $r_1(A) \ r_2(A) \ w_1(A) \ w_2(A)$

**Not Serializable**



$T_1$	$T_2$
read(A)	read(A)
write(A)	write(A)

Red arrows indicate dependencies: from  $T_2$  read(A) to  $T_1$  read(A), from  $T_1$  write(A) to  $T_2$  write(A), and from  $T_2$  write(A) to  $T_1$  write(A).

d)  $r_2(A) \ r_1(A) \ w_2(B) \ w_1(A)$

**Serializable**

**What is the equivalent serial schedule?**

$T_2 T_1$



$T_1$	$T_2$
read(A)	read(A)
write(A)	write(B)


A red arrow indicates a dependency from  $T_2$  read(A) to  $T_1$  read(A).

## EXERCISE 2

For each of the following schedules, state whether it is **serializable**, **recoverable** and **cascadeless**. Justify your answers.  $r_i$  denotes a read (of transaction  $T_i$ ) and  $w_i$  a write (of transaction  $T_i$ ).

a)  $w_1(X) \ r_2(X) \ w_1(X) \ c_2 \ a_1$

$T_1$	$T_2$
write(X)	
	read(X)
write(X)	
	commit
abort	



Serializable?

No The precedence graph has a cycle.

Recoverable?

No  $T_2$  reads data item  $X$  and commits, but then  $T_1$  aborts.

Cascadeless?

No If  $T_1$  aborts,  $T_2$  must also abort.

## EXERCISE 2 (CONTD)

b)  $r_2(X) w_3(X) c_3 w_1(Y) c_1 r_2(Y) w_2(Z) c_2$

$T_1$	$T_2$	$T_3$
	read(X)	write(X) commit
write(Y) commit	read(Y) write(Z) commit	

Serializable?


Yes The equivalent serial schedule is  
 $T_1 T_2 T_3$

Recoverable?

Yes  $T_2$  reads a data item written by  $T_1$   
and commits after  $T_1$ .

Cascadeless?


Yes The commit of  $T_1$  appears before  
the commit of  $T_2$ .



## EXERCISE 2 (CONTD)

c)  $r_1(X) w_2(X) c_2 w_1(X) c_1 r_3(X) c_3$

$T_1$	$T_2$	$T_3$
read(X)		
	write(X)	
	commit	
write(X)		
commit		
		read(X)
		commit



### Serializable?

**No** The precedence graph has a cycle.

### Recoverable?

**Yes**  $T_3$  reads a data item written by  $T_1$  and commits after  $T_1$ .

### Cascadeless?

**Yes** The commit of  $T_1$  appears before the commit of  $T_3$ .

### **Why recoverable and cascadeless?**


The  $T_2$  write is **useless** (called a **blind write**). It is not preceded by a read and it is immediately overwritten by  $T_1$ . So, its value is never seen by anyone.

## EXERCISE 3

For each of the following schedules, answer the questions.

a) Is the following schedule **serializable**?

$T_1$	$T_2$	$T_3$
read(X)	read(Y)	
	write(Y)	
write(X)		write(Z)
	read(X)	
	write(X)	
		read(Y)
		write(Y)
write(Z)		




The schedule is **not serializable** since there is a **cycle**  $T_1$ ,  $T_2$ ,  $T_3$ ,  $T_1$  in the precedence graph.

There is no equivalent serial schedule because of the cycle.

## EXERCISE 3 (CONTD)

- b) Is the following schedule **serializable**, **recoverable** and **cascadeless**?

$T_1$	$T_2$
read(A) write(A)  commit	write(B)  read(A)  commit



### Serializable?

**Yes** The equivalent serial schedule is  $T_1 T_2$ .

### Recoverable?

**Yes**  $T_2$  reads A, written by  $T_1$  and commits after  $T_1$ .

### Cascadeless?

**No** If  $T_1$  aborts,  $T_2$  must also abort.

**How would you place the commit statements in order to make the schedule cascadeless?**

$T_1$	$T_2$
read(A) write(A) commit	write(B)  read(A) commit



## EXERCISE 3 (CONTD)

c) Is the following schedule **recoverable** and **cascadeless**?

$T_1$	$T_2$
read(A) write(A)	read(A)  write(B) commit
commit	

Recoverable?

**Yes** No transaction reads data items after they have been written by the other.

Cascadeless?

**Yes** No transaction reads data items after they have been written by the other.

## EXERCISE 4

Consider the following schedule consisting of three transactions  $T_1$ ,  $T_2$  and  $T_3$ .  $r_i$  denotes a read (of transaction  $T_i$ ) and  $w_i$  is a write (of transaction  $T_i$ ).

**Schedule:**  $r_3(Z)$   $w_3(Z)$   $r_1(X)$   $r_2(Y)$   $w_2(Y)$   $w_1(X)$   $r_1(Y)$   $r_3(X)$

- Show that the schedule is serializable by constructing the precedence graph.
- Which is the equivalent serial schedule?
- Modify the original schedule so it becomes **recoverable, but not cascadeless** by adding commit operations in the appropriate locations in the schedule.
- Modify the schedule so it becomes **both recoverable and cascadeless** by adding commit operations in the appropriate locations in the schedule.


## EXERCISE 4 (CONTD)

**Schedule:**  $r_3(Z)$ ,  $w_3(Z)$ ,  $r_1(X)$ ,  $r_2(Y)$ ,  $w_2(Y)$ ,  $w_1(X)$ ,  $r_1(Y)$ ,  $r_3(X)$ ,  $c_1$ ,  $c_2$ ,  $c_3$ ;

- a) Show that the schedule is serializable by constructing the precedence graph.

$T_1$	$T_2$	$T_3$
$read(X)$  $write(X)$ $read(Y)$  $commit$	$read(Y)$ $write(Y)$    $commit$	$read(Z)$ $write(Z)$    $read(X)$   $commit$

Precedence Graph



- b) What is the equivalent serial schedule?

$T_2$ ,  $T_1$ ,  $T_3$

## EXERCISE 4 (CONTD)

- c) Modify the original schedule so it becomes **recoverable, but not cascadeless** by adding commit operations in the appropriate locations in the schedule

**Recall:** A schedule is recoverable if the **commit** of a transaction  $T_j$  that reads data items *previously written* by a transaction  $T_i$  appears **after** the commit operation of  $T_i$ .

$T_1$  reads  $Y$  written by  $T_2$

$\Rightarrow T_1$  must commit after  $T_2$ .

$T_2$  does not read any data items written by other transactions.

$T_3$  reads  $X$  written by  $T_1$

$\Rightarrow T_3$  must commit after  $T_1$ .

### Schedule:

$r_3(Z), w_3(Z), r_1(X), r_2(Y), w_2(Y),$   
 $w_1(X), r_1(Y), r_3(X), c_2, c_1, c_3$

$T_1$	$T_2$	$T_3$
		read(Z) write(Z)
read(X)		
	read(Y) write(Y)	
write(X) read(Y)		
		read(X)
	commit	
commit		commit

## EXERCISE 4 (CONTD)

- d) Rewrite the schedule so it becomes **both recoverable and cascadeless** by adding commit operations in the appropriate locations in the schedule.

**Recall:** A schedule is cascadeless if, for each pair of transactions  $T_i, T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the **commit** operation of  $T_i$  appears before the **read** operation of  $T_j$ .

- The commit of  $T_1$  must appear before the **read(X)** of  $T_3$ .
- The commit of  $T_2$  must appear before the **read(Y)** of  $T_1$ .
- The commit of  $T_3$  must appear at the end of  $T_3$ .

### Schedule:

$r_3(Z), w_3(Z), r_1(X), r_2(Y), w_2(Y),$   
 $c_2, w_1(X), r_1(Y), c_1, r_3(X), c_3$

$T_1$	$T_2$	$T_3$
read(X)		read(Z) write(Z)
	read(Y) write(Y) commit	
write(X) read(Y) commit		
		read(X) commit

