

# COMP 3311 DATABASE MANAGEMENT SYSTEMS

FINAL REVIEW

# EXTERNAL SORTING

Consider the relation Sailor(sailorId, name, rating, age), which is not sorted. Assume each attribute is 25 bytes, the page size is 1,000 bytes and there are 11,000 tuples. For the following questions, apply external sorting using a buffer of 11 pages.

A. How many sorted runs will be produced in pass 0?

a) 10

b) 11

c) 100

d) 110

e) None of the above.

tuple size: 100 bytes  
blocking factor: 10 tuples/page  
file size: 1100 pages  
memory size: 11 pages

In pass zero, 11 pages at a time are sorted in memory.

Therefore, in pass zero  $1100/11 = 100$  sorted runs are produced.

## EXTERNAL SORTING (CONT'D)

B. What is the total number of passes required to sort the result completely (including pass 0)?

- a) 1
- b) 2
- c) 3**
- d) 4
- e) 5

tuple size: 100 bytes  
blocking factor: 10 tuples/page  
file size: 1100 pages  
memory size: 11 pages

After pass zero there are 100 sorted runs.

In pass 1, 10 runs at a time are merged, producing 10 sorted runs.

In pass 2, these 10 sorted runs are merged to produce the final output.

Therefore, the total number of passes = 3.

## EXTERNAL SORTING (CONT'D)

C. What is the total page I/O cost of sorting?

a) 4,400

b) 5,500

c) 6,600

d) 7,770

e) None of the above.

tuple size: 100 bytes  
blocking factor: 10 tuples/page  
file size: 1100 pages  
memory size: 11 pages

There are a total of three passes.

In each pass the whole relation (1100 pages) is read and written.

Therefore, the total cost is  $3 * 2 * 1100 = 6,600$  page I/Os.

# QUERY PROCESSING

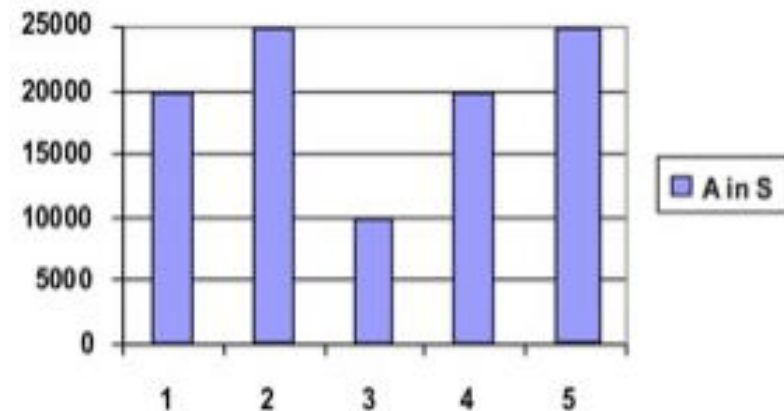
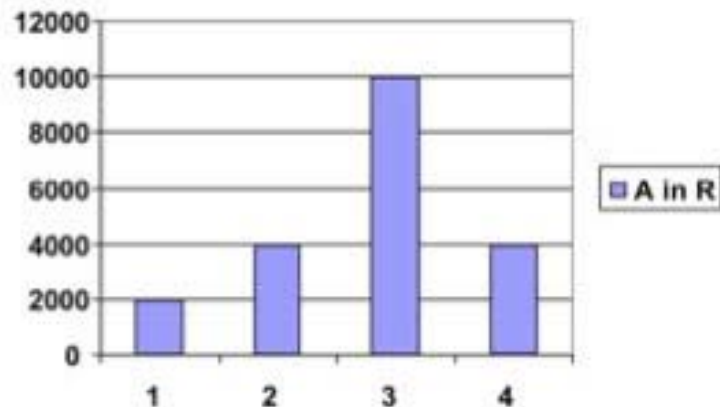
Consider the two relations  $R(A, \underline{B}, C)$  and  $S(A, \underline{B}, Y)$ .

$R$  contains 20,000 tuples and  $S$  contains 100,000 tuples.

Assume that for both relations, 10 tuples fit per page (i.e., the size of  $R$  is 2,000 pages and that of  $S$  is 10,000 pages).

The possible values of attribute  $A$  in  $R$  are  $\{1, 2, 3, 4\}$ , whereas the possible values of  $A$  in  $S$  are  $\{1, 2, 3, 4, 5\}$ .

The following histograms present statistical information about the occurrences of values for  $A$  in  $R$  and  $S$  (e.g., there are 2,000 tuples with  $A=1$  in  $R$  and 20,000 tuples with  $A=1$  in  $S$ ).



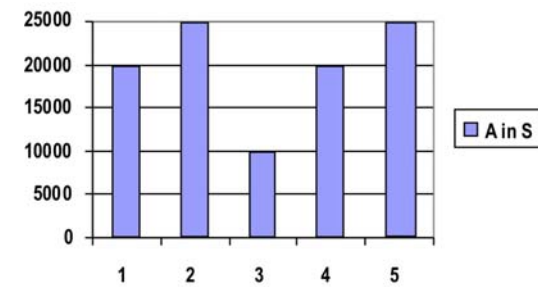
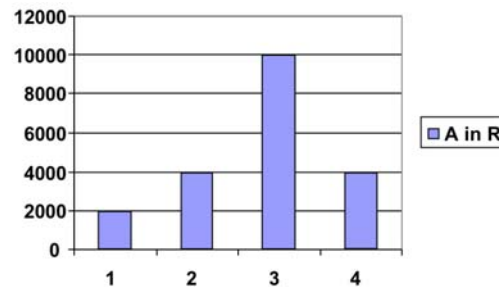
$R(A, \underline{B}, C)$   
 $S(A, \underline{B}, Y)$

## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

A. How many **tuples** are there in the result of the query  
 $(R \text{ JOIN}_{R.A=S.A} S)$  (i.e., what is the cardinality of the output)?

- a) 20,000
- b) 100,000
- c) 120,000
- d) 320,000
- e) 320,000,000**



According to the histograms, each tuple of  $R$  with value  $R.A=1$  (of which there are 2,000) can join with 20,000 tuples of  $S$  with  $S.A=1$ . That is, the join result will contain  $2,000 * 20,000 = 40 * 10^6$  tuples where  $R.A=S.A=1$ .

Performing the same computation for  $A = 2, 3, 4$  and  $5$  the final result will contain  $(40 + 100 + 100 + 80 + 0) * 10^6 = 320,000,000$  tuples.

R(A, B, C)  
S(A, B, Y)

## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

**B. What is the minimum cost (in terms of page I/Os) for computing  $(R \text{ JOIN}_{R.A=S.A} S)$  using block nested-loop join and how many main memory pages are needed?**

- a) Minimum cost is 12,000 and I need 2,000 main memory pages.
- b) Minimum cost is 12,000 and I need 2,002 main memory pages.**
- c) Minimum cost is 12,000 and I need 12,000 main memory pages.
- d) Minimum cost is 320,000 and I need 2,002 main memory pages.
- e) Minimum cost is 320,000 and I need 12,000 main memory pages.

Since there is no index, at least both relations have to be read (i.e., the minimum cost is 12,000 page I/Os, *independent of the algorithm*).

For block nested-loop join, the smaller relation **R** (2,000 pages) needs to be able to be kept in memory, so that **S** is scanned only once.

Since 1 page is needed for reading **S** and 1 page for holding the output, a total of 2,002 main memory pages are needed.

$R(A, \underline{B}, C)$   
 $S(A, \underline{B}, Y)$

## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

C. We want to compute  $(R \text{ JOIN}_{R.A=S.A} S)$  using block nested-loop join with *R as the outer relation*. What is the **minimum number of main memory pages needed** in order to achieve a cost of 42,000 page I/Os?

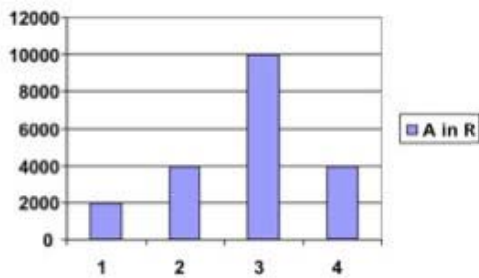
- a) 502
- b) 736
- c) 1,000
- d) 1,002
- e) 2,000

Since  $R$  (2,000 pages) is the outer relation and the total cost is 42,000 page I/Os,  $S$  needs to be scanned  $(42,000 - 2,000) / 10,000 = 4$  times.

If  $S$  is scanned 4 times, this means that 4 “blocks” of  $R$  need to be read, and each “block” should be at least  $2,000 / 4 = 500$  pages.

Since, 1 page is needed for reading  $S$  and 1 page for the output, in total 502 main memory pages are needed.





## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

**D. We want to compute  $(R \text{ JOIN}_{R.A=S.A} S)$  using hash join with *R as the build input*. How many buckets should be used for partitioning and what is the minimum main memory requirement?**

- a) 4 buckets and at least 202 main memory pages.
- b) 4 buckets and at least 1,002 main memory pages.**
- c) 10 buckets and at least 202 main memory pages.
- d) 10 buckets and at least 502 main memory pages.
- e) 20 buckets and at least 102 main memory pages.

The join condition is  $R.A=S.A$  and there are only four values of  $A$  in  $R$ . Therefore, 4 buckets should be used.

According to the histograms, half of the tuples (10,000 tuples) of the build input  $A$ , have value  $R.A=3$ . Thus, the bucket size should be 1,000 pages so that each bucket of  $R$  can fit into main memory.

In addition, we need 1 page for reading  $S$  and 1 for the output.

Therefore, 1,002 main memory pages are needed.

R(A, B, C)  
S(A, B, Y)

## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

E. Given that R.B is a **NOT NULL** foreign key referencing S.B, **how many tuples** are in the result of the query  $(R \text{ JOIN}_{R.B=S.B} S)$ ?

- a) 20,000
- b) 100,000
- c) 120,000
- d) 320,000
- e) 320,000,000

Each tuple of R joins with exactly one tuple in S.

Therefore, the size of the join result is the same as the size of R, namely, 20,000 tuples.

R(A, B, C)  
S(A, B, Y)

## QUERY PROCESSING (CONTD)

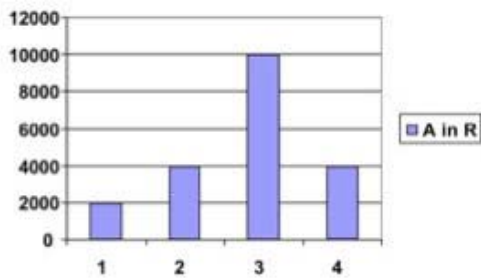
Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

F. We want to compute  $(R \text{ JOIN}_{R.B=S.B} S)$  using indexed nested-loop join with *R as the outer relation*. Assume that there is a hash index on S.B with no overflow buckets (i.e., finding an index entry has cost 1). What is the **total page I/O cost** of the join?

- a) 6,000
- b) 22,000
- c) 40,000
- d) 42,000**
- e) None of the above.

For each tuple of R, the index entry with the corresponding value of B is found and the pointer followed to the tuple of S for a cost of **2 page I/Os** per tuple. This is repeated for all **20,000** tuples of R, for a total cost of  $20,000 * 2 = 40,000$  page I/Os.

In addition, the tuples of R have to be read for a total cost of **40,000 + 2,000 = 42,000** page I/Os.



## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

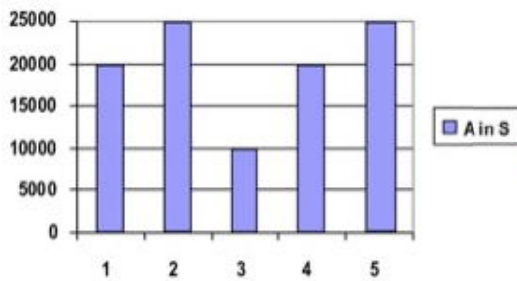
**G. How many tuples are expected in the result of  $((\sigma_{A=1}R) \text{ JOIN}_{R.B=S.B} S)$  and what is the minimum page I/O cost of processing the query using indexed nested-loop join with *R as the outer relation*. Assume that the only index is a hash index on *S.B* with no overflow buckets.**

- a) The result has 2,000 tuples and the cost is 6,000 page I/Os.
- b) The result has 2,000 tuples and the cost is 22,000 page I/Os.
- c) The result has 20,000 tuples and the cost is 40,000 page I/Os.
- d) The result has 20,000 tuples and the cost is 42,000 page I/Os.
- e) None of the above.

According to the histograms, there are only 2,000 tuples in *R* with  $R.A=1$ . Each of these tuples, matches exactly 1 tuple in the join with *S* ( $R.B=S.B$ ). Thus, the result contains 2,000 tuples.

Finding the matching tuple in *S*, has cost 2 page I/Os (see previous question) and so finding all matches for the 2,000 tuples has cost 4,000 page I/Os.

Adding the cost of reading *R*, the total cost is 6,000 page I/Os.



## QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

**H. How many tuples are expected in the result of  $((\sigma_{A=1}R) \text{ JOIN}_{R.B=S.B} (\sigma_{A=3}S))$  and what is the minimum page I/O cost of processing the query using index nested-loop join with *R as the outer relation*. Assume that the only index is a hash index on S.B with no overflow buckets.**

- a) Expected number of tuples is 200 with cost 600 page I/Os.
- b) Expected number of tuples is 200 with cost 6,000 page I/Os.**
- c) Expected number of tuples is 2,000 with cost 6,000 page I/Os.
- d) Expected number of tuples is 2,000 with cost 42,000 page I/Os.
- e) None of the above.

According to the histogram of S, among the 2,000 tuples in the result of  $((\sigma_{A=1}R) \text{ JOIN}_{R.B=S.B} S)$ , only  $10,000/100,000=0.1$  (i.e., 10%) are expected to satisfy the condition  $S.A=3$ . Thus, the result is expected to contain  $(2,000 * 0.1) = 200$  tuples.

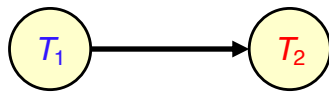
The query processing is the same as in the previous question with cost 6,000 page I/Os.

# TRANSACTION MANAGEMENT

3. Consider the following schedules of two transactions  $T_1$  and  $T_2$ . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a READ and w a WRITE operation.

a) Schedule:  $r_1(A) w_1(A) r_2(A) w_2(B)$

**Serial:**  $T_1, T_2$



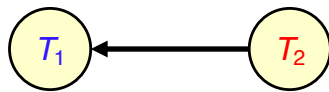
$T_1$	$T_2$
read(A) write(A)	read(A) write(B)

# TRANSACTION MANAGEMENT (CONTD)

3. Consider the following schedules of two transactions  $T_1$  and  $T_2$ . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a READ and w a WRITE operation.

b) Schedule:  $r_1(A) r_2(A) w_1(A) w_2(B)$

**Serializable:**  $T_2, T_1$



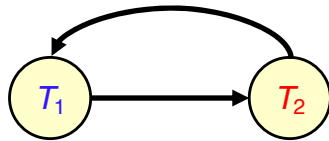
$T_1$	$T_2$
read(A)	
	read(A)
write(A)	
	write(B)

# TRANSACTION MANAGEMENT (CONTD)

3. Consider the following schedules of two transactions  $T_1$  and  $T_2$ . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a READ and w a WRITE operation.

c) Schedule:  $r_1(A) r_2(A) w_2(A) w_1(A)$

**Not Serializable**



$T_1$	$T_2$
read(A)	read(A)
write(A)	write(A)

Diagram illustrating the schedule operations for transactions  $T_1$  and  $T_2$ . Red arrows indicate dependencies: one from  $T_1$ 's read(A) to  $T_2$ 's read(A), and another from  $T_2$ 's write(A) to  $T_1$ 's write(A). The crossing arrows demonstrate a cycle in the dependency graph, proving the schedule is not serializable.



# TRANSACTION MANAGEMENT (CONTD)

4. Consider the schedule  $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$  (where  $c_1$  and  $c_2$  indicate the commit statements).

a) Is the schedule recoverable? Why?

It is recoverable ( $T_2$  reads database item **A** written by  $T_1$ , but  $T_1$  commits before  $T_2$ ).

Is the schedule cascadeless? Why?

It is not cascadeless ( $T_2$  read database item **A** before  $T_1$  commits).

$T_1$	$T_2$
read(A)	
write(A)	
	read(A)
	write(B)
commit	
	commit

# TRANSACTION MANAGEMENT (CONTD)

4. Consider the schedule:  $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$  (where  $c_1$  and  $c_2$  indicate the commit statements).

b) Change the time of the commits ( $c_1, c_2$ ) in the schedule in a) so that it becomes a cascadeless schedule.

$r_1(A) w_1(A) c_1 r_2(A) w_2(B) c_2$

Transaction  $T_2$  reads the value of **A** written by  $T_1$ , after  $T_1$  commits.

$T_1$	$T_2$
read(A) write(A) commit	read(A) write(B) commit

## TRANSACTION MANAGEMENT (CONTD)

4. Consider the schedule:  $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$  (where  $c_1$  and  $c_2$  indicate the commit statements).

(c) Is the schedule  $r_2(A) r_1(A) w_1(A) w_2(B) c_2 c_1$  recoverable and cascadeless?

**Both recoverable and cascadeless – no transaction reads items after they have been written by the other.**

$T_1$	$T_2$
read(A) write(A)  commit	read(A)  write(B) commit

## TRANSACTION MANAGEMENT (CONTD)

5. Rewrite the schedule  $r_2(A) r_1(A) w_1(A) w_2(B)$  according to the 2PL protocol (i.e., add lock-S, lock-X, unlock statements below). Explain briefly whether the schedule is allowed by 2PL.

The schedule is allowed by 2PL .  $T_1$  has to wait until  $T_2$  unlocks  $A$ .

$T_1$	$T_2$
   lock-S(A) read(A) lock-X(A) $\Rightarrow$ wait write(A)	 lock-S(A) read(A)    lock-X(B) write(B)

# TRANSACTION MANAGEMENT (CONTD)

6. Rewrite the schedule  $r_2(A)$   $r_1(A)$   $w_1(A)$   $w_2(B)$  according to timestamp-ordering protocol (i.e., add RTS (read timestamp) and WTS (write timestamp) statements). Assume that the timestamps of  $T_1$  and  $T_2$  are 2 and 1, respectively. The initial read and write timestamps of A and B are both 0.

$T_1$ [TS=2]	$T_2$ [TS=1]
read(A) RTS(A)=2; WTS(A)=0	read(A) RTS(A)=1; WTS(A)=0
write(A) RTS(A)=2; WTS(A)=2	
	write(B) RTS(B)=0; WTS(B)=1

## Read

If  $TS(T_i) < WTS(Q)$  **rollback**

If  $TS(T_i) \geq WTS(Q)$

$RTS(Q) = \max(TS(T_i), RTS(Q))$

## Write

If  $TS(T_i) < RTS(Q)$  **rollback**

If  $TS(T_i) < WTS(Q)$  **ignore**

Otherwise  $WTS(Q) = TS(T_i)$

# TRANSACTION MANAGEMENT (CONTD)

7. Rewrite the schedule  $r_2(A) r_1(A) w_1(A) w_2(B)$  according to the multi-version timestamp-ordering protocol (i.e., add RTS (read timestamp) and WTS (write timestamp) statements and specify the versions of the items). Assume that the timestamps of  $T_1$  and  $T_2$  are 1 and 2, respectively and that the initial versions of items are  $A_0$  and  $B_0$ . Complete the correct version numbers (e.g.,  $\text{read}(A_0)$  instead of  $\text{read}(A)$ ).

$T_1$ [TS=1]	$T_2$ [TS=2]
$\text{read}(A_0)$ RTS( $A_0$ )=2; WTS( $A_0$ )=0	$\text{read}(A_0)$ RTS( $A_0$ )=2; WTS( $A_0$ )=0
$\text{write}(A)$ <span style="border: 1px solid black; padding: 2px;">TS(<math>T_1</math>)=1 &lt; RTS(<math>A_0</math>)=2 <math>\Rightarrow</math> rollback</span>	$\text{write}(B)$

## Read

**Reads always succeed**

set RTS( $Q$ ) = TS( $T_i$ )

## Write

If TS( $T_i$ ) < RTS( $Q$ ) **rollback**

If TS( $T_i$ ) = WTS( $Q$ )

**overwrite contents**

If TS( $T_i$ ) > WTS( $Q$ )

**create new version**

set R/WTS( $Q$ )=TS( $T_i$ )