

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 3

ENTITY-RELATIONSHIP (E-R) MODEL

AND DATABASE DESIGN

E-R MODEL: CONSTRAINTS

A constraint is a logical restriction or property of data that for any set of data values:

- we can determine whether the constraint is **true** or **false**;
- we expect the constraint to be **always true**;
- we can **enforce** the constraint.

Examples:


on attributes

salary is between
\$0 and \$100,000




on relationships

every employee works in
at most one department



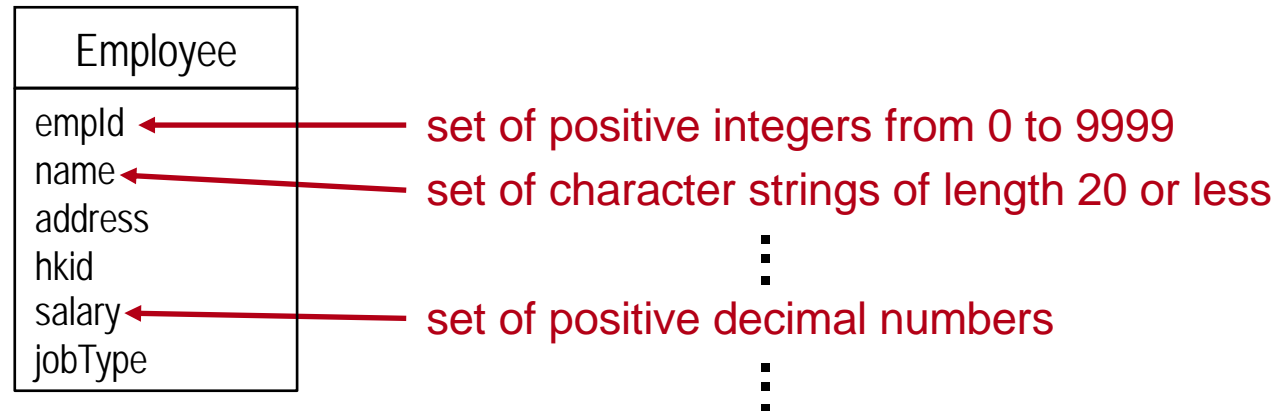
not every employee
is a manager



👉 **Constraints add additional semantics (meaning) to data**
(so as to more accurately reflect the application requirements).

ATTRIBUTE CONSTRAINTS: DOMAIN

A domain constraint restricts an attribute to only have certain values.



A domain constraint can be specified as a type for the attribute and/or a logical predicate that restricts the values.

ATTRIBUTE CONSTRAINTS: KEY

- If the values of some attributes **uniquely identify an entity instance** then these attributes are a **key** of the entity.
- An entity may have **more than one key**.
 - A **candidate key** is a **minimal set of attributes** (i.e., all attributes are needed) that **uniquely identifies** an entity instance.
- One candidate key is selected to be the **primary key**.

Employee
<u>empld</u>
name
address
hkid
salary
jobType

 **This has enforcement implications for implementation.**

primary key \Rightarrow uniqueness is **automatically enforced** by a DBMS

other candidate keys \Rightarrow uniqueness is **not automatically enforced** by a DBMS

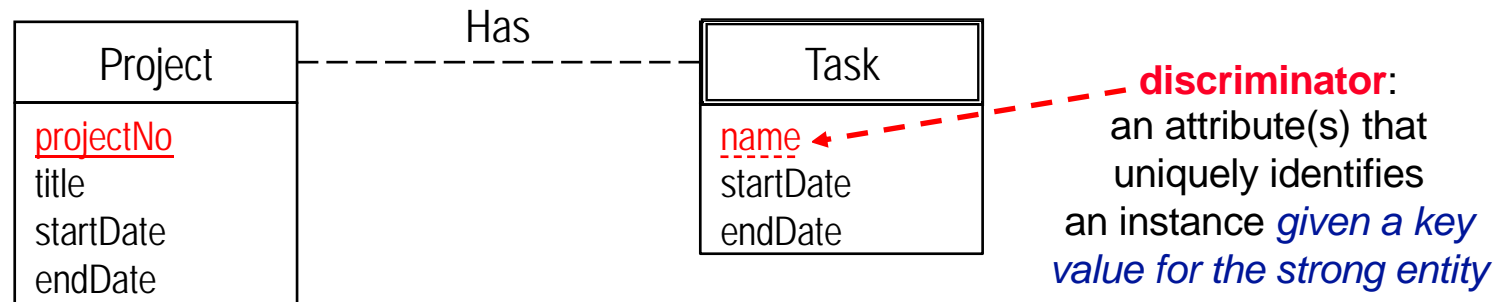
- A candidate/primary key can be composed of a set of attributes \Rightarrow **composite key**.

EnrollsIn
<u>studentId</u>
<u>courseCode</u>
grade

STRONG ENTITY VS. WEAK ENTITY: KEY

Strong entity: An entity that has a primary key.

Weak entity: An entity that does not have a primary key.



- A weak entity must be associated with a strong entity, called the **identifying entity**, to be meaningful.
 - ✎ A weak entity depends on its identifying entity for its existence.
- The relationship associating the weak entity to the strong entity is called the **identifying relationship**.
- A **discriminator**, *if present*, uniquely identifies a weak entity instance within its identifying relationship.


GENERALIZATION CONSTRAINTS: COVERAGE

Disjointness

(a) overlapping

A superclass instance can relate to **more than one** subclass.


E.g., a given company can be both a customer and a supplier at the same time.



b) disjoint

A superclass instance can relate to **at most one** subclass.

E.g., a given employee can be either a secretary or an engineer, but not both at the same time.




GENERALIZATION CONSTRAINTS: COVERAGE (CONT'D)

Completeness

a) partial

A superclass instance **does not need to** relate to any of the subclasses.


E.g., a loan does not need to be a mortgage (loan) or a tax (loan) or a car (loan)—there are other kinds of loans.



(b) total


A superclass instance **must** relate to at least one of the subclasses.

E.g., a given customer must be either a personal or a business customer.




GENERALIZATION CONSTRAINTS: **COVERAGE** (CONT'D)


overlapping, partial




disjoint, partial



overlapping, total




disjoint, total



RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION

Cardinality specifies the maximum number and **participation** specifies the minimum number of relationship instances in which an entity may participate.




For a given project, how many employees can manage it?

☞ Each project is managed by one and only one employee.

For a given employee, how many projects can he/she manage?

☞ An employee does not have to manage any project, but may manage several (i.e., an unknown number of) projects.

RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION



minimum cardinality (min-card) \Rightarrow participation constraint

$\text{min-card}(E_1, R)$: The *minimum* number of relationship instances in which *each entity* of E_1 *must* participate in R .

$\text{min-card}(E_1, R) = 0 \Rightarrow$ partial participation

$\text{min-card}(E_1, R) > 0 \Rightarrow$ total participation

maximum cardinality (max-card) \Rightarrow cardinality constraint

$\text{max-card}(E_1, R)$: The *maximum* number of relationship instances in which *each entity* of E_1 *may* participate in R .

RELATIONSHIP CONSTRAINTS:

SPECIAL MAXIMUM CARDINALITIES

$\text{max-card}(E_1, R) = 1$ and $\text{max-card}(E_2, R) = 1$

👉 **one-to-one relationship (1:1)**

$\text{max-card}(E_1, R) = 1$ and $\text{max-card}(E_2, R) = N$

👉 **one-to-many relationship (1:N)**


$\text{max-card}(E_1, R) = N$ and $\text{max-card}(E_2, R) = 1$

👉 **many-to-one relationship (N:1)**

$\text{max-card}(E_1, R) = N$ and $\text{max-card}(E_2, R) = N$


👉 **many to many relationship (N:M)**

RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION




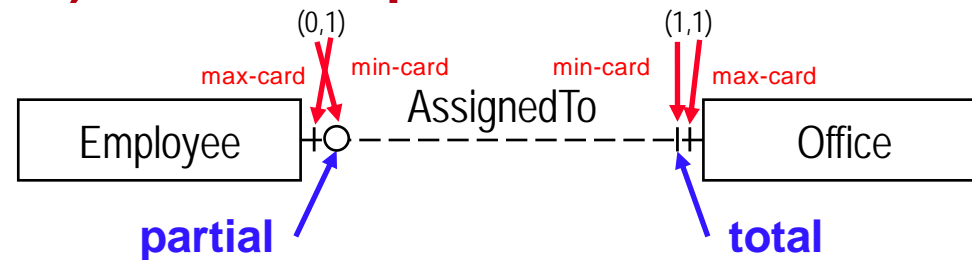
An employee does not have to manage a project, but can manage several projects.

Every project must be managed by an employee and at most one employee.



RELATIONSHIP CONSTRAINTS: EXAMPLE CARDINALITY & PARTICIPATION

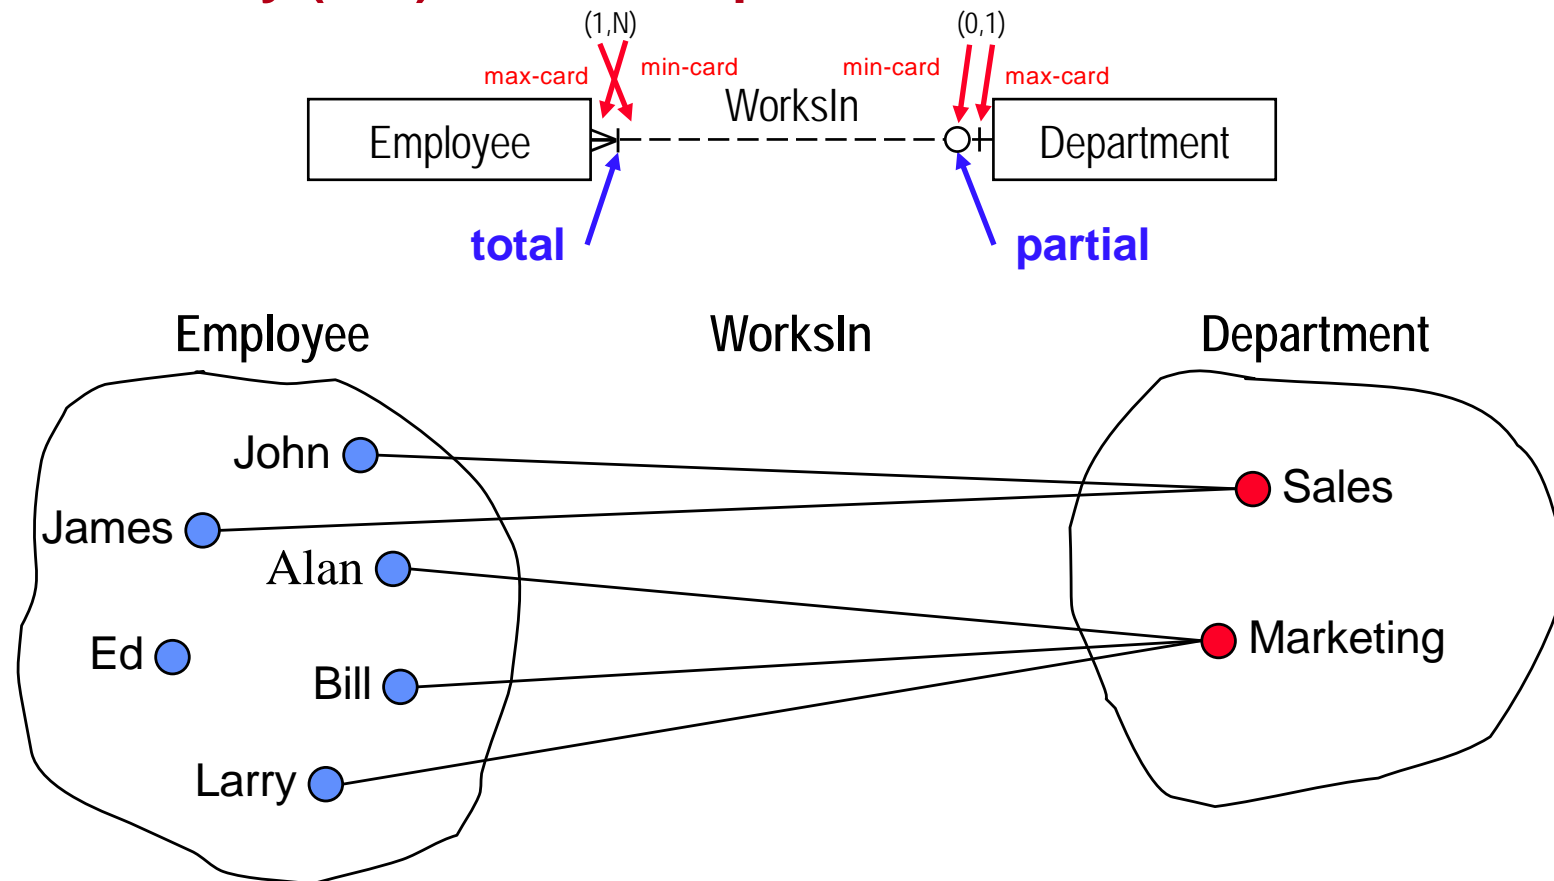
one-to-one (1:1) relationship



RELATIONSHIP CONSTRAINTS:

EXAMPLE CARDINALITY & PARTICIPATION



one-to-many (1:N) relationship



RELATIONSHIP CONSTRAINTS:

EXAMPLE CARDINALITY & PARTICIPATION

many-to-many (N:M) relationship

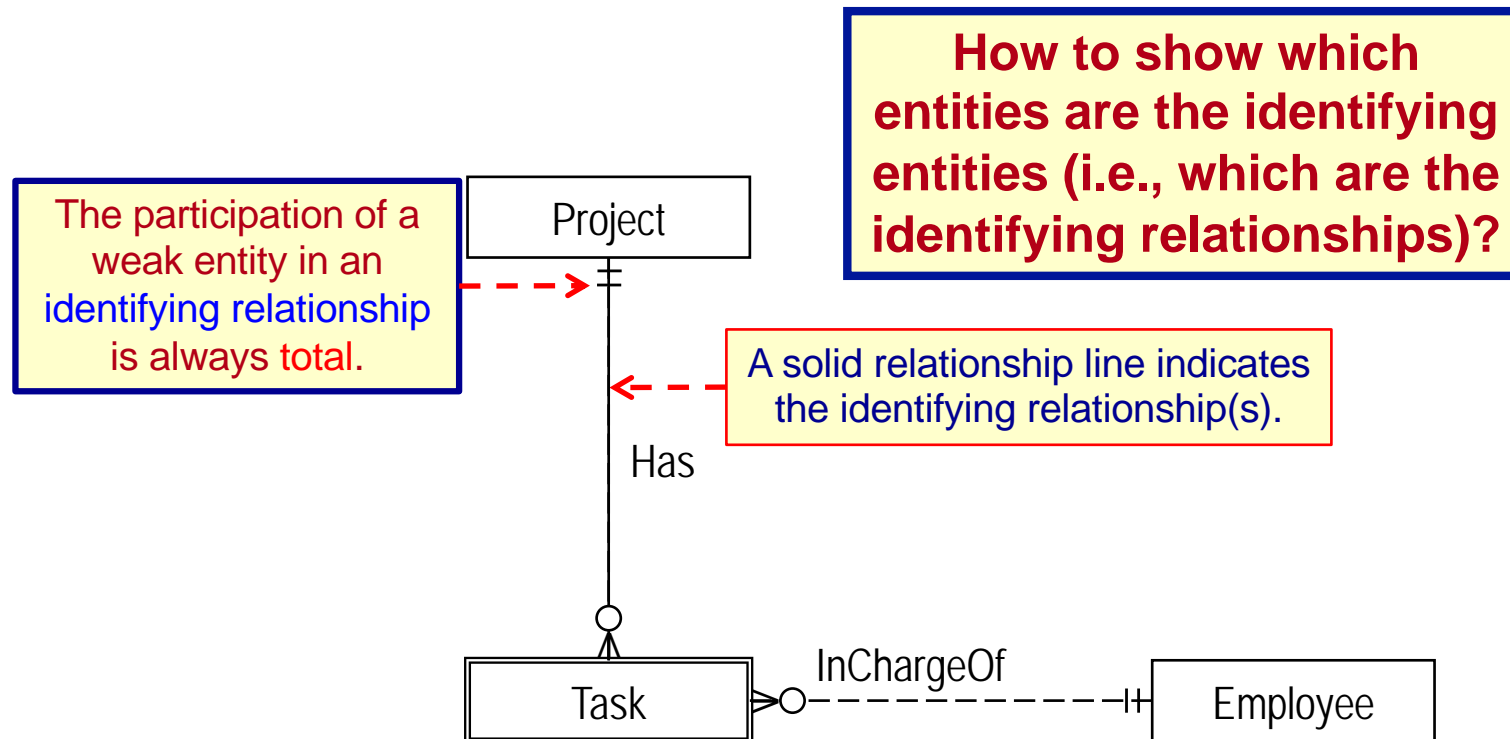


RELATIONSHIP CONSTRAINTS:

WEAK ENTITY PARTICIPATION

- A weak entity may be related to more than one strong entity, but may depend on only some of these for its **existence**.


👉 Only some of the strong entities are identifying entities.



RELATIONSHIP CONSTRAINTS: EXCLUSION

An *exclusion (XOR)* constraint specifies that at most one entity instance, among several entity types, can participate in a relationship with a single “root” entity.


Example: A task can be related to *either* an internal project *or* an external project, *but not both*.



RELATIONSHIP CONSTRAINTS: **EXCLUSION** (CONTD)

An **exclusion (XOR)** constraint specifies that **at most one entity instance, among several entity types, can participate in a relationship with a single “root” entity.**

Example: A task can be related to *either* an internal project *or* an external project, *but not both*.



ANALYZING APPLICATION REQUIREMENTS

1. Identify entities

- What are the major concepts about which data needs to be **permanently** stored?
- Focus on the “**big picture**”, not the details.
 - E.g., student, course not name, address, email, description, credits, etc.

2. Identify relationships between entities

- How are the major concepts related? How do they interact?
- What interactions need to be **permanently** stored.
 - E.g., students enroll in courses not students browse courses

3. Identify properties of entities and relationships


- For each entity and relationship, what information needs to be **permanently** stored.

DESIGN CHOICE: ENTITY VERSUS ATTRIBUTE


entity: When **several properties** can be associated with the concept.

attribute: When the concept has a **simple atomic structure** or **no property** of interest.

Office as attribute



Office as entity




DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP


entity: When the concept **represents something distinct** in the application domain **with several properties**.

relationship: When the concept is **not a distinct application domain concept** and/or has **no property of interest**.

Account as entity



Account as relationship

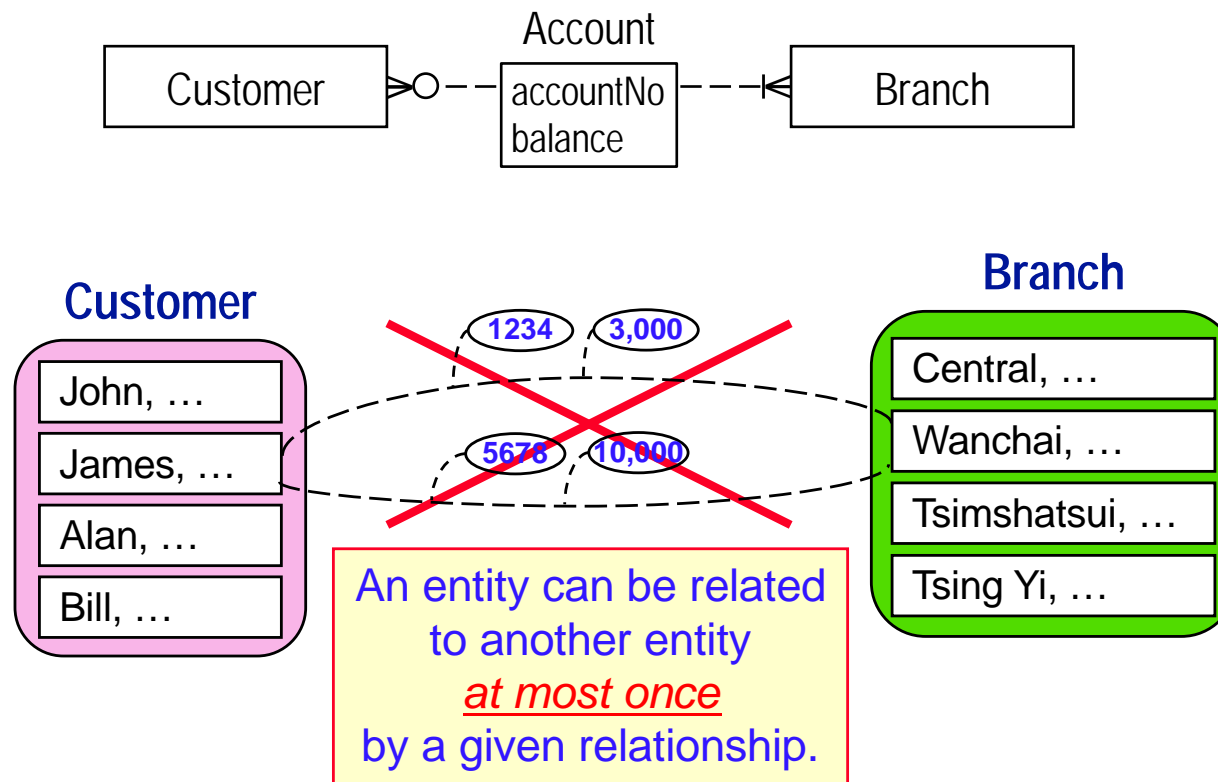


What if you want to have several accounts for a customer?

DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP

(CONT'D)

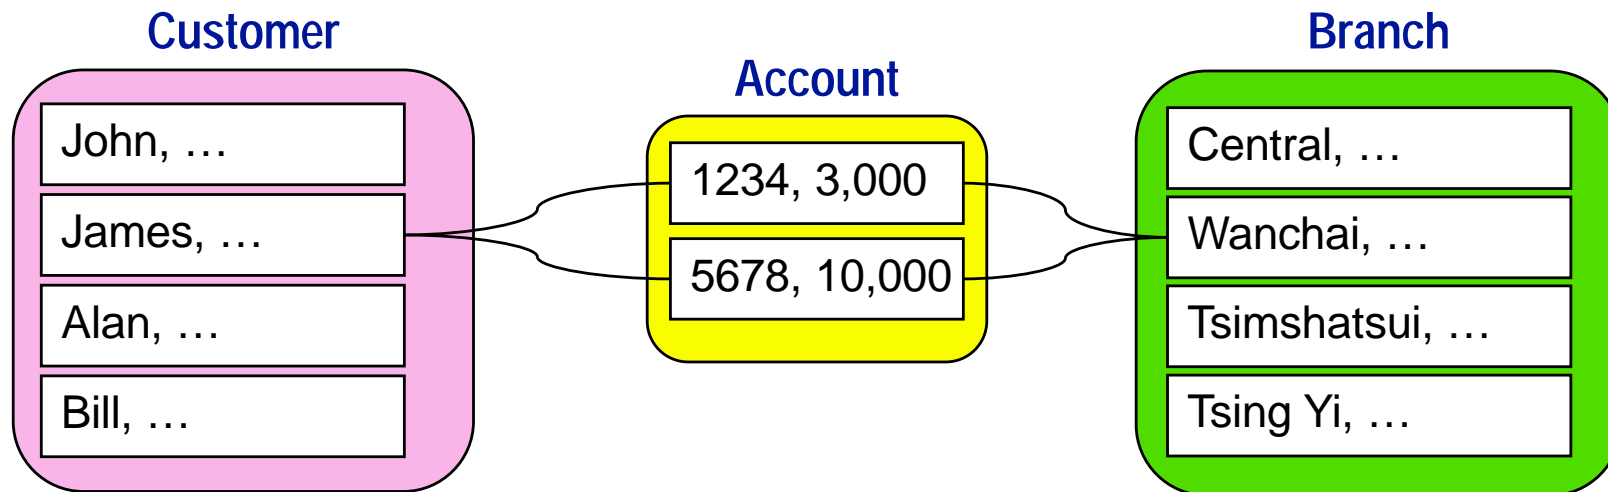

We want to represent the fact that James has two accounts at the same branch.



DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP

(CONT'D)

We need to use an entity for Account!



There can be only one relationship instance of a given relationship type between the same two entity instances.

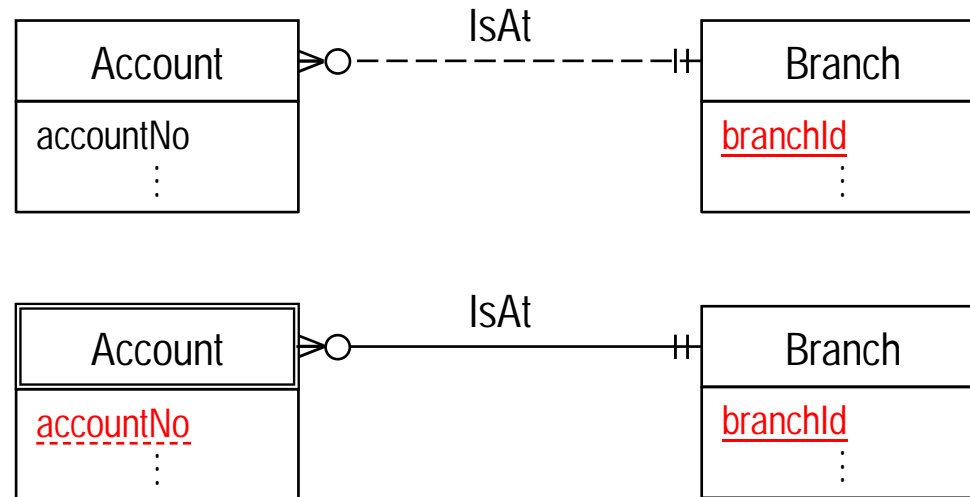
DESIGN CHOICE: STRONG VERSUS WEAK ENTITY

strong entity: When the concept can be **uniquely identified in the application domain** (i.e., it has a **key**).

weak entity: When the concept has **no unique identifier**.


Suppose an account must be associated with exactly one branch and two different branches are allowed to have accounts with the same number.

Should Account be a strong or weak entity?



DESIGN CHOICE: PLACING AN ATTRIBUTE

? salary



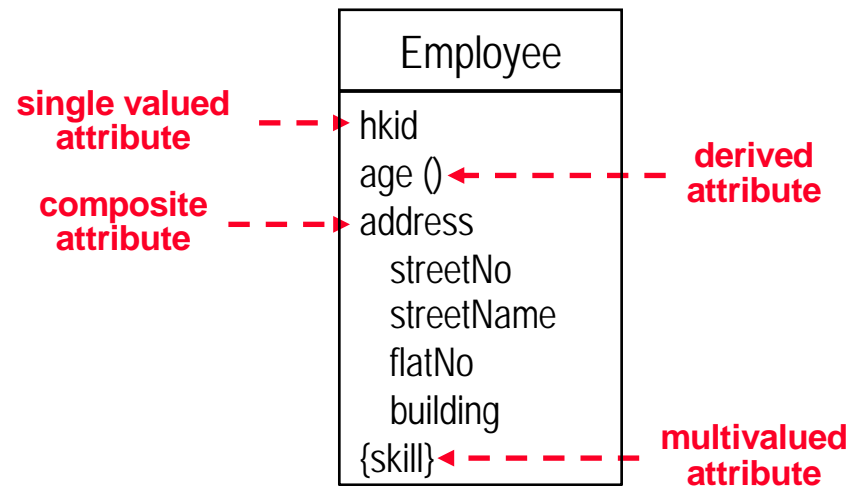
Where to place salary?

Relationship attributes are usually needed only for
many to many relationships!
(But can also be used in one to one and one to many relationships.)

E-R MODEL & DB DESIGN: SUMMARY

Entity-Relationship Model


entity and
attribute




relationship

○ - partial (min-card=0) | - total (min-card=1)

(a) one-to-one
(1:1)




(b) one-to-many
(1:N)




E-R MODEL & DB DESIGN: SUMMARY (CONT'D)


(c) many-to-many
(N:M)




(d) explicit
cardinality



(e) relationship
attribute




(f) recursive
(with role names)



E-R MODEL & DB DESIGN: SUMMARY (CONT'D)

(h) generalization/specialization



(i) exclusion constraint

