

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 9

STRUCTURED QUERY LANGUAGE (SQL)

EXAMPLE BANK RELATIONAL SCHEMA

Branch(branchName, branchCity, assets)

Client(clientName, clientStreet, clientCity)

Loan(loanNo, amount, *branchName*)

Account(accountNo, balance, *branchName*)

Borrower(*clientName*, *loanNo*)

Depositor(*clientName*, *accountNo*)

Attribute names in
italics are foreign
key attributes.

DATA DEFINITION LANGUAGE (DDL)

The SQL DDL allows the specification of:

- The **schema** for each relation (attributes).
- The **types of values** associated with each attribute (i.e., the domain of values the attribute, such as string, number, date, etc.).
- **User-defined types** and **domains**.
- **Integrity constraints** (ICs).
 - domain, key, foreign key, general
- The **physical storage structure** of each relation on disk.
- The **set of indices** to be maintained for each relation.
- **Security and authorization** information for each relation.

BASIC TYPES

<code>char(<i>n</i>)</code>	Fixed length character string with length <i>n</i> .
<code>varchar2(<i>n</i>)</code>	Variable-length character string with maximum length <i>n</i> .
<code>int</code>	An integer (a finite subset of the integers that is machine-dependent).
<code>smallint</code>	A small integer (a machine-dependent subset of the integer domain type).
<code>number(<i>p</i>,<i>d</i>)</code>	A fixed point number with a total of <i>p</i> digits (the precision) and <i>d</i> digits to the right of the decimal point.
<code>float(<i>n</i>)</code>	Floating point number, with user-specified precision of at least <i>n</i> digits.
<code>date</code>	A date containing a (4 digit) year, month and day of month.
<code>time</code>	The time of day, in hours, minutes and seconds.
<code>timestamp</code>	A combination of date and time.

- **Null values** are allowed in all the domain types.
- Declaring an attribute to be **not null** prohibits null values for that attribute.

 **Some relational systems also allow user-defined types.**

USER-DEFINED TYPES

- The `create type` clause is used to define a new type.

```
create type PersonName as char(25) final;
```

```
create type Dollars as numeric(12,2) final;
```

```
create type Pounds as numeric(12,2) final;
```

The keyword `final` is required by the SQL:1999 standard, but has no real meaning in this context; some implementations allow the `final` keyword to be omitted.

 **Not all relational systems support user-defined types.**

USER-DEFINED DOMAINS

- The `create domain` clause is used to define a new domain.

```
create domain hourlyWage numeric(5,2);
```

- Differences between user-defined types and domains:
 - Domains can have constraints, such as `not null`, specified on them and can have default values defined for variables of the domain; types cannot.
 - Domains are not strongly typed; types are strongly typed.

 **Not all relational systems support user-defined domains.**

CREATING RELATIONS

- The **create table** command is used to define and create a relation.
- The **domain type** of each attribute needs to be specified.
 - **Basic domain types:** **char**(*n*), **varchar2**(*n*), **int**, **smallint**, **number**(*p*,*d*), **float**(*n*), **date**, **time**, **timestamp**
 - A **default** value can be specified for an attribute.
 - **Null values** are allowed in all the basic domain types.

👉 **The domain type of an attribute is enforced by the DBMS whenever tuples are added or modified.**

```
create table Student (
  studentId  char(8) not null,
  name       varchar2(45) not null,
  email      varchar2(30),
  birthdate  date not null,
  cga        number(3,2));
```

```
create table EnrollsIn (
  studentId  char(8) not null,
  courseId   char(8) not null,
  grade      number(4,1) default 0 not null);
```

ALTERING AND DESTROYING RELATIONS

- The **alter table** command is used to **add** attributes to, **modify** attributes in or **drop** attributes from an existing relation.

Example:

```
alter table Student  
add firstYear int;
```

The schema is altered by adding a new attribute and extending every tuple in the current instance with a null value for the new attribute.

```
alter table Student  
drop column firstYear;
```

The schema is altered by dropping the attribute from the relation and deleting its value in every tuple.

- The **drop table** command deletes **all** information about a relation (both **data and schema**).

Example:

```
drop table Student;
```


INTEGRITY CONSTRAINTS (IC)

An integrity constraint (IC) ensures that authorized changes to the database do not result in a loss of data consistency.

👉 **An IC guards against accidental damage to the database.**

- ICs are based upon the requirements of the real-world application that is being described in the database relations.
 - An IC is a statement about *all possible instances*!
 - For the **Student** relation, we know, from common knowledge, that **name** is not a key, but the constraint that an attribute, such as **studentId**, is a key has to be given to us by the client.
- We can check a database instance to see if an IC is violated, but we can **never** infer that an IC is true by looking at a database instance. **Why?**

DOMAIN CONSTRAINTS

- Domain constraints define **valid values** for attributes.
- They **test values** inserted into the database and **test queries** to ensure that the comparisons make sense.
- Besides a basic domain type, additional constraints can be specified on attributes in the **create table** command.

not null specifies that null values are **not allowed**

primary key specifies a key for a relation (the value of a key attribute **cannot be null** \Rightarrow no need to specify **not null**)

unique specifies that an attribute or a set of attributes is a candidate key (the attribute value(s) **can be null**)

foreign key specifies that one or more attributes refer to a primary key attribute in another relation

check specifies a predicate that the values in every tuple must satisfy

PRIMARY KEY, UNIQUE CONSTRAINTS

- A relation can possibly have many *candidate keys* one of which is chosen as the *primary key*.

```
create table Student (  
  studentId  char(8),  
  name       varchar2(45) not null,  
  email      varchar2(30) not null,  
  birthdate  date not null,  
  cga        number(3,2),  
  primary key (studentId),  
  unique (email));
```

```
create table EnrollsIn (  
  studentId  char(8),  
  courseId   char(8) not null,  
  grade      number(4,1) default 0 not null,  
  primary key (studentId),  
  unique (courseId, grade));
```

X

Used carelessly, an IC can prevent the storage of database instances that should be allowed!

FOREIGN KEY CONSTRAINT

Recall: A **foreign key** is a set of attributes in one relation whose values must match the values of the primary key in another relation or be null.

☞ **A foreign key must reference the primary key of the referenced relation.**

Example: Only students listed in the Student relation should be allowed to enroll for courses.

```
create table EnrollsIn (  
    studentId char(8),  
    courseId char(8),  
    grade number(4,1) default 0 not null,  
    primary key (studentId, courseId),  
    foreign key (studentId) references Student(studentId));
```

☞ **Every studentId value in the EnrollsIn relation must reference a tuple in the Student relation with a matching studentId value.**

FOREIGN KEY: ENFORCING REFERENTIAL INTEGRITY

- What should be done if an **EnrollsIn** tuple with a non-existent student id is inserted?

👉 **Reject it!**

- What should be done if a **Student** tuple is deleted?
 1. **Disallow deletion** of a **Student** tuple that is referred to by an **EnrollsIn** tuple (*default action*).
 2. Alternatively, delete all **EnrollsIn** tuples that refer to it (**on delete cascade**).
 3. Set **studentId** in **EnrollsIn** tuples that refer to it to a *default value* (**on delete set default**).
 4. Set **studentId** in **EnrollsIn** tuples that refer to it to a *null value* (**on delete set null**).

👉 3 and 4 are not applicable in the example since **studentId** is part of the primary key.

FOREIGN KEY:

ENFORCING REFERENTIAL INTEGRITY (cont'd)

- What should be done if the primary key **student id** of a tuple in **Student** is updated?

👉 **Reject it!**

- Alternatively, propagate the update to the tuples in the **EnrollsIn** relation with matching student ids (**on update cascade**).

```
create table EnrollsIn (  
  studentId char(8),  
  courseId char(10),  
  grade number(4,1) default 0 not null,  
  primary key (studentId, courseId),  
  foreign key (studentId) references Student(studentId)  
    on delete cascade  
    on update cascade);
```

The referential integrity actions in the **referencing relation** (**EnrollsIn**) are triggered when a tuple in the **referenced relation** (**Student**) is deleted or updated.

Oracle Note
Oracle does not support
on update cascade.

CHECK CLAUSE: ATTRIBUTES

- The **check** clause is used to add an integrity constraint for an attribute and can contain an **arbitrary predicate**.

☞ The predicates are similar to those allowed in a **where** clause.

- It is specified in the definition of a relation and checked whenever there is an update to the relation.

Example: Ensure that semester can have only specified values and that year is between 2020 and 2024.

```
create table Section (  
    courseId    char(8),  
    sectionId   char(2),  
    semester    char(6),  
    year        char(4) check (year between '2020' and '2024'),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer')));
```

CHECK CLAUSE: DOMAINS

- The **check** clause can be used in a **create domain** clause to add an integrity constraint to the domain.

Example: Ensure that an **hourlyWage** domain allows only values greater than a specified value.

```
create domain hourlyWage numeric(5, 2)  
    constraint wageTest check (value >= 40.00);
```

- The new domain **hourlyWage** is declared to be a decimal number with 5 digits, 2 of which are after the decimal point.
- The domain has a constraint named **wageTest** that ensures that **hourlyWage** is greater than 40.00.

 **The constraint name is optional, but useful to indicate which constraint an update violated.**

VIEWS

- Views provide a way to hide certain data from certain users.
- The **create view** command creates a view.

create view *view-name* **as** *<query expression>*

where: *<query expression>* is any legal SQL query.

Example: Create a view from the Loan relation that hides the amount.

```
create view BranchLoan as
select loanNo, branchName
from Loan;
```

- The **drop view** command deletes a view.

Example: **drop view** BranchLoan;

Loan

loanNo	amount	branch Name
L-170	30000	Central
L-155	170000	Tsimshatsui
L-230	40000	Central



BranchLoan

loanNo	branch Name
L-170	Central
L-155	Tsimshatsui
L-230	Central




VIEWS (cont'd)

- A view can be queried the same as any other relation.

Example: Find all loans in the Central branch.

```
select loanNo
from BranchLoan
where branchName='Central';
```

loanNo	branch Name
L-170	Central
L-155	Tsimshatsui
L-230	Central



loanNo
L-170
L-230

- A user who has access to the **BranchLoan** view, but not the **Loan** relation, cannot see the amount.
- Some database systems allow view relations to be stored and kept up-to-date if the relations used to define the view change.

☞ **Such views are called materialized views.**

VIEWS (cont'd)

- Not all views can allow update operations.
- A view is **updateable** (i.e., tuples can be inserted, updated and deleted) if **all** of the following **conditions** are **satisfied** by the query that defines the view.
 - The **from** clause has only one relation.
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
 - Any attribute not listed in the **select** clause can be set to **null** (i.e., it does not have a **not null** constraint and is not part of a primary key.
 - The query does not have a **group by** or **having** clause.

 **The specific rules for view update are often system dependent.**

TUPLE DELETION

- The **delete** command deletes *zero or more tuples* from a relation.

Example: Delete all account tuples at the Central branch.

```
delete from Account  
where branchName='Central';
```

- Conceptually, deletion is done in two steps.

1. Find the tuples to delete.

```
select * from Account  
where branchName='Central';
```

2. Delete the tuples found.

 Deletion can be done **only if** no integrity constraints are violated!



COMPLEX DELETION

- The **where** clause predicate in a **delete** statement can be as complex as in a **select** statement.

Example: Delete all accounts at the Mongkok branch.

```
delete from Depositor
where accountNo in (select accountNo
                    from Depositor natural join Account
                    where branchName= 'Mongkok');
```

👉 **Must also delete the accounts of these depositors!**



TUPLE INSERTION

- The **insert** command adds one or more tuples to a relation.

Example: Add a new Account tuple.

```
insert into Account values ('A-732', 1200 , 'Central');
```

Example: Add a new tuple to Account with balance set to null.

```
insert into Account values ('A-733', null, 'Central');
```

✎ **The order of the values must match the order of the attributes in the relation.**

- Attribute names need to be *specified explicitly* for order-independent insertion.

```
insert into Account (accountNo, branchName, balance)  
values ('A-734', 'Central', 1200);
```



COMPLEX INSERTION

- Insertion values can be obtained from the result of a query.

Example: Create a \$200 savings account for all loan clients of the Central branch. Let the loan number serve as the account number for the new savings account.

```
insert into Account  
(select loanNo, 200, branchName  
from Loan  
where branchName='Central');
```

```
insert into Depositor  
(select clientName, loanNo  
from Loan natural join Borrower  
where branchName='Central');
```

The order of the attributes in the **select** clause must match the order of the attributes in the table being inserted into.

Note: The keyword **values** is omitted when the values to insert are obtained from a **select** statement.



TUPLE UPDATE

- The **update** command is used to change a value in a tuple.

Example: Increase all accounts with balance over \$10,000 by 6%; all other accounts receive 5%.

```
update Account  
set balance=balance*1.06  
where balance>10000;
```

```
update Account  
set balance=balance*1.05  
where balance<=10000;
```

👉 **Need two update statements!**

👉 **The order is important! Why?**



CONDITIONAL UPDATE: CASE STATEMENT

- This update can be specified more easily using the **case** statement.

Example: Increase all accounts with balances over \$10,000 by 6%; all other accounts receive 5%.

```
update Account
set balance= case
    when balance<=10000 then balance*1.05
    else balance*1.06
end;
```

