

# COMP 3311

# DATABASE MANAGEMENT


# SYSTEMS

## LECTURE 10


## STRUCTURED QUERY LANGUAGE (SQL)

# DATABASE SYSTEM ARCHITECTURES

## Centralized

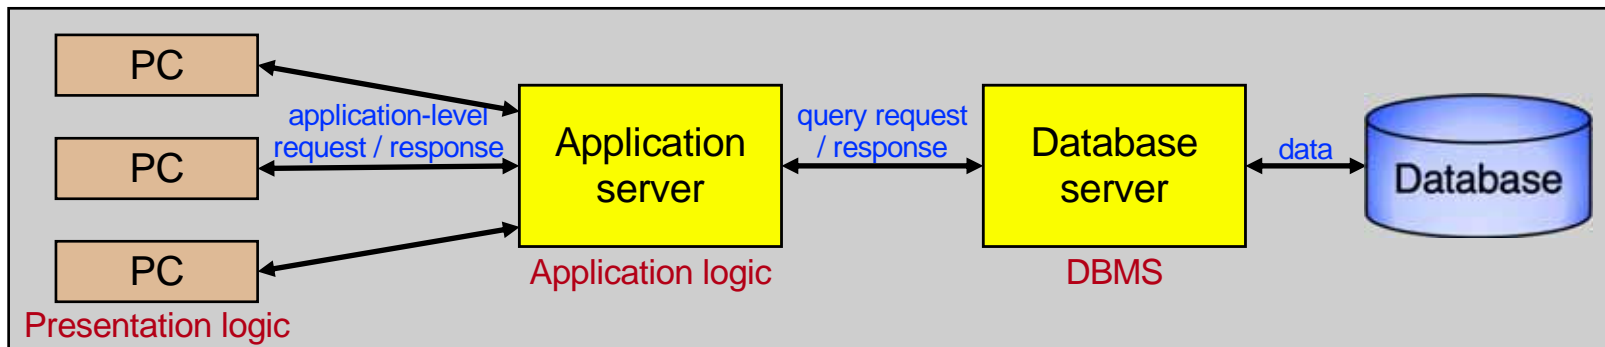


## Two-tier (client-server)

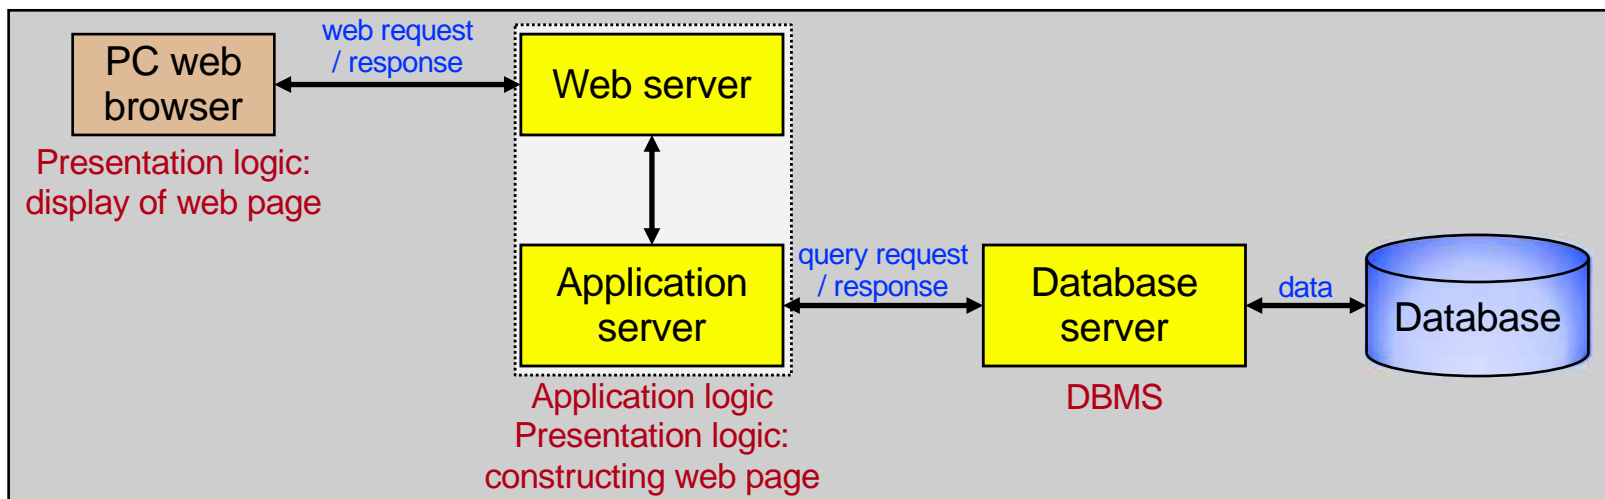


# DATABASE SYSTEM ARCHITECTURES (CONT'D)

## Three-tier



## n-tier (e.g., web database connectivity)



# DATABASE SYSTEM ARCHITECTURES (cont'd)

## Tiered system architectures

- The aim is to decouple the centralized architecture by combining a central computer's powerful computing capabilities with the flexibility of PCs.

Fat client: The presentation logic and application logic are handled by the client (i.e., the PC).

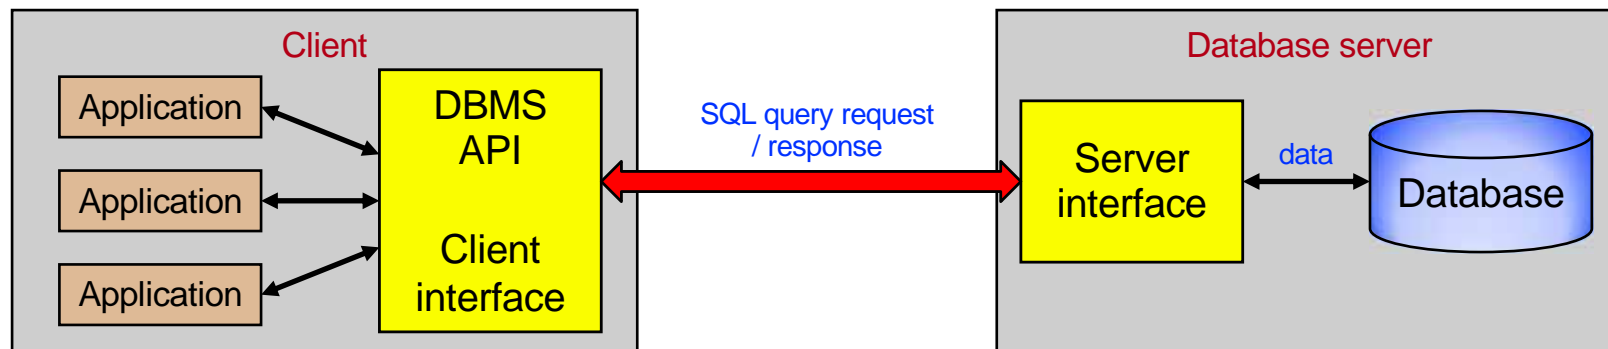
- Common in cases where it makes sense to couple an application's workflow (e.g., opening of windows, screens, and forms) with its look and feel (i.e., its front-end).

Thin client: Only the presentation logic is handled by the client.

- Common in cases where the application logic and database logic are tightly coupled (fat server/thin client architecture).
- It is also possible to **decouple** the **application logic** from the DBMS and place this in a separate layer (i.e., an application server) and to also **decouple** some of the **presentation logic** from the PCs (e.g., a web server).

# API BASICS

- To utilize DBMS services, client applications use a specific application programming interface (API) provided by the DBMS.
  - Facebook, Google, Instagram, etc. have such APIs.
- The DBMS API exposes an interface through which **the services provided by the DBMS** can be accessed.
  - The client and server interfaces often are implemented in the form of network sockets that use a specific port number on the server (e.g., port 1521 for the course Oracle Database server).



# PROPRIETARY VS UNIVERSAL API

## Proprietary, DBMS-specific API

- Provided by most vendors, but requires client applications to:
  - be aware of the DBMS that will be utilized on the server side.
  - be modified to interact with a new DBMS API.

## Generic, vendor-agnostic universal API

- Allows easy porting of applications to multiple DBMSs.
- Does not allow access to some vendor-specific optimizations.

### **Examples:**

- ODBC (Open Database Connectivity)
- JDBC (Java Database Connectivity)
- ADO.NET (ActiveX Data Objects for Microsoft's .NET framework)

# EMBEDDED VS CALL-LEVEL API

## Embedded API

- SQL statements are part of the host programming language source code.
- An SQL pre-compiler parses and checks the SQL instructions *before* the program is compiled and replaces these with source code instructions native to the host programming language used.

## Call-level API

- Passes SQL instructions to the DBMS by direct calls to a series of procedures, functions or methods provided by the API.
- The calls perform actions such as setting up a database connection, sending queries and iterating over the query result.

# EARLY VS LATE BINDING

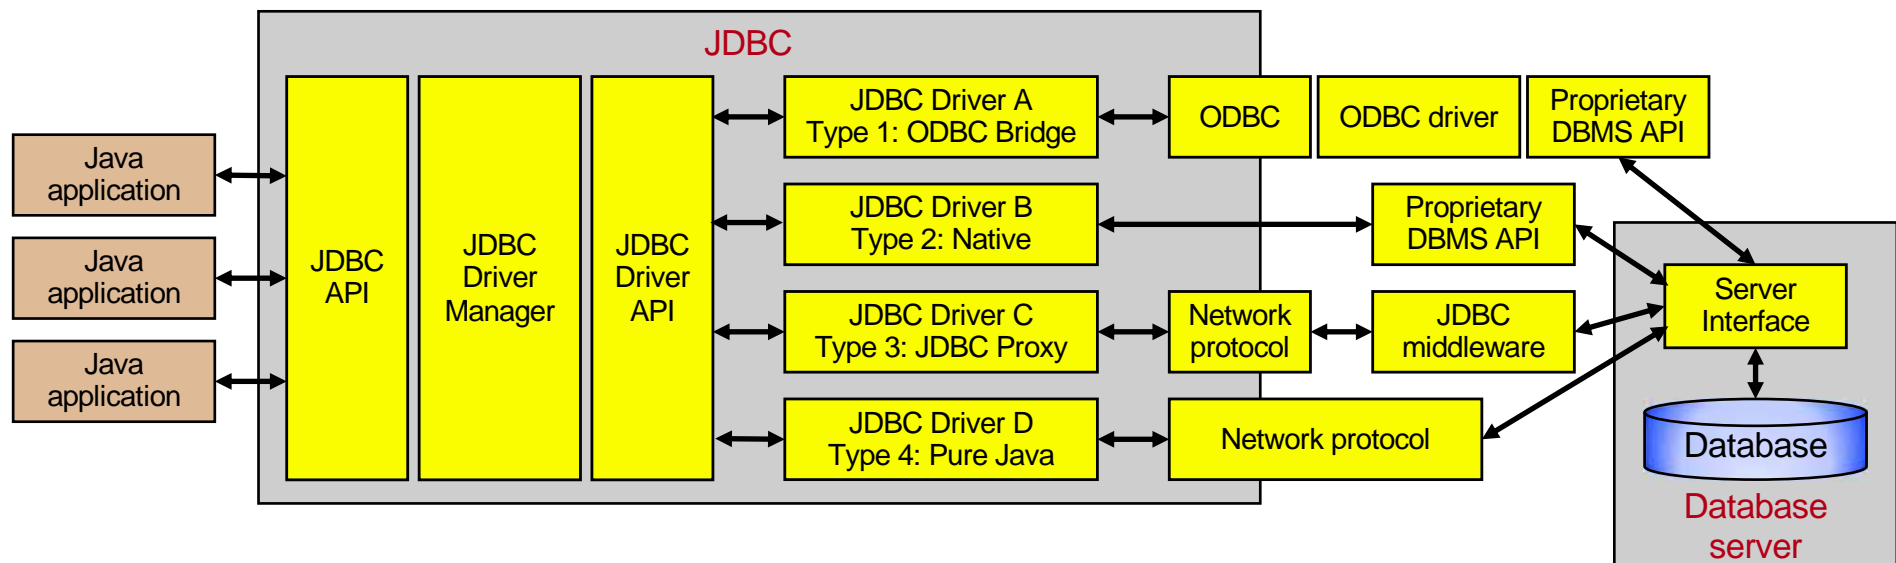
- **SQL binding** is the translation of SQL statements in a programming language into a form that can be executed by the DBMS.
  - Involves performing tasks such as validating table and attribute names, checking whether the user or client has sufficient access rights and generating an efficient query plan to access the data.
- **Early binding** performs these tasks only once *before program execution* (i.e., using a pre-compiler with an embedded API).
- **Late binding** performs these tasks every time *at runtime* (i.e., when using a call-level API).

👉 **It is still possible to do early binding using call-level APIs by using stored procedures in the DBMS.**



# JAVA DATABASE CONNECTIVITY (JDBC)

- A **call-level API** for Java that is **highly portable** and **object-oriented**.
  - Database connections, drivers, queries and results are all expressed as objects, based on uniform interfaces.
  - Exposes a uniform set of methods, no matter which DBMS is used.

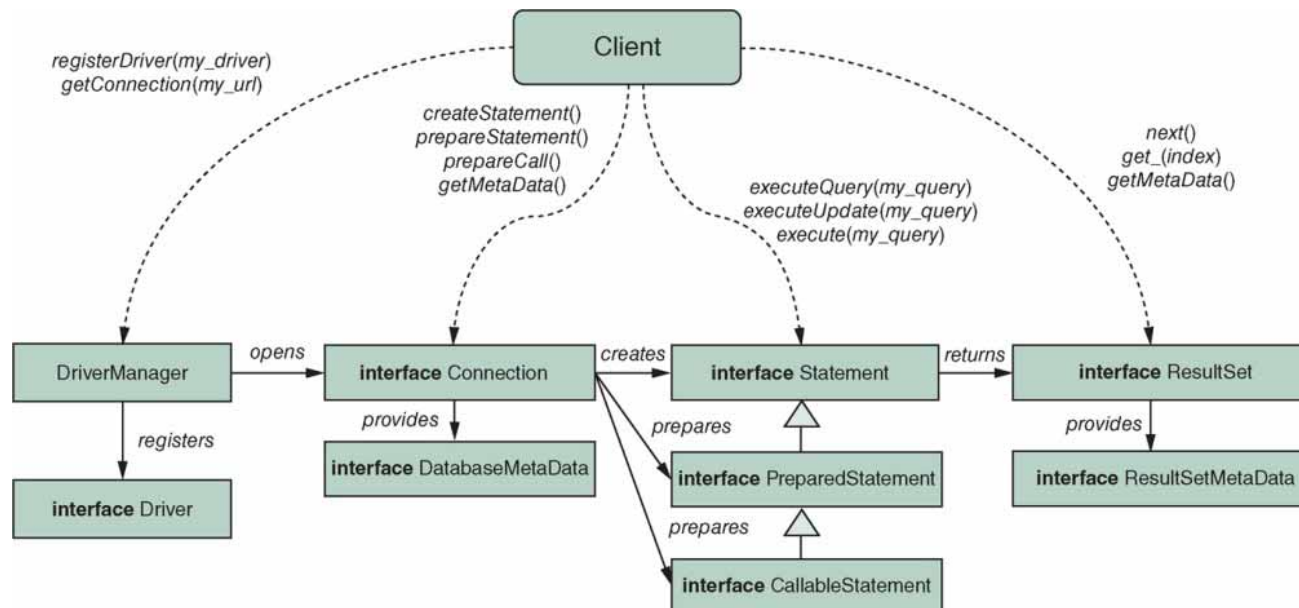


# JAVA DATABASE CONNECTIVITY (JDBC) (CONT)

- **DriverManager** is a singleton object which acts as the basic service to register and manage JDBC drivers.
- **Driver** objects implement the **Driver** interface and enable the communication between the **DriverManager** and the DBMS using one of four types of drivers.
  - **Type-1: ODBC Bridge** drivers do not communicate with a DBMS directly, but instead translate JDBC calls to corresponding ODBC calls.
  - **Type-2: Native** drivers are written in Java, but will communicate to a DBMS using its native database API.
  - **Type-3: JDBC Proxy** drivers are written in Java. The JDBC client uses standard networking sockets to communicate with an application server, which converts the calls into a native database API call or utilizes a different JDBC type-1, 2, or 4 driver on its end.
  - **Type-4: Pure Java** drivers are written in Java and use networking functionality to connect directly with the database server.
- The **getConnection** method creates a database connection using one of the registered drivers.

# JAVA DATABASE CONNECTIVITY (JDBC) (CONT)

```
DriverManager.registerDriver(new org.sqlite.JDBC());  
String dbURL = "jdbc:sqlite:my_database";  
Connection conn = DriverManager.getConnection(dbURL);  
if (conn != null) {  
    System.out.println("Connected to the database");  
    DatabaseMetaData dm = conn.getMetaData();  
    System.out.println("Driver name: " + dm.getDriverName);  
    conn.close();  
}
```



# JDBC: EXECUTING SQL STATEMENTS

```
Statement selectStatement = conn.createStatement("select * from Book");  
ResultSet selectResult = selectStatement.executeQuery();  
.  
.  
.
```

- SQL statements are executed and results returned within the context of a database connection.
- A `Statement` object represents an SQL instruction.
- An SQL statement is created with the `createStatement` method.
- The `executeQuery` method is used to execute an SQL `select` statement and return a `ResultSet` representing the returned data.
- The `executeUpdate` method is used to execute `insert`, `update` and `delete` statements.

# JDBC: CURSORS

```
Statement selectStatement = conn.createStatement("select * from Book");
ResultSet selectResult = selectStatement.executeQuery();
while (selectResult.next()) {
    String bookTitle = selectResult.getString("title"); // or: .getString(1);
    int bookQuantity = selectResult.getInt("quantityInStock"); // or: .getInt(2);
    System.out.println(bookTitle + " has " + bookQuantity + " books in stock.");
}
```

- Since SQL is a set-oriented language, the query result (the **ResultSet** object) will generally contain multiple tuples.
- Host languages, such as Java, are essentially record-oriented.
  - They cannot work on more than one record/tuple at a time.
- To overcome this **impedance mismatch**, JDBC uses a **cursor mechanism** to step through result sets.
  - A cursor is a programmatic control structure that enables **one-by-one traversal** over the records in a query result set.

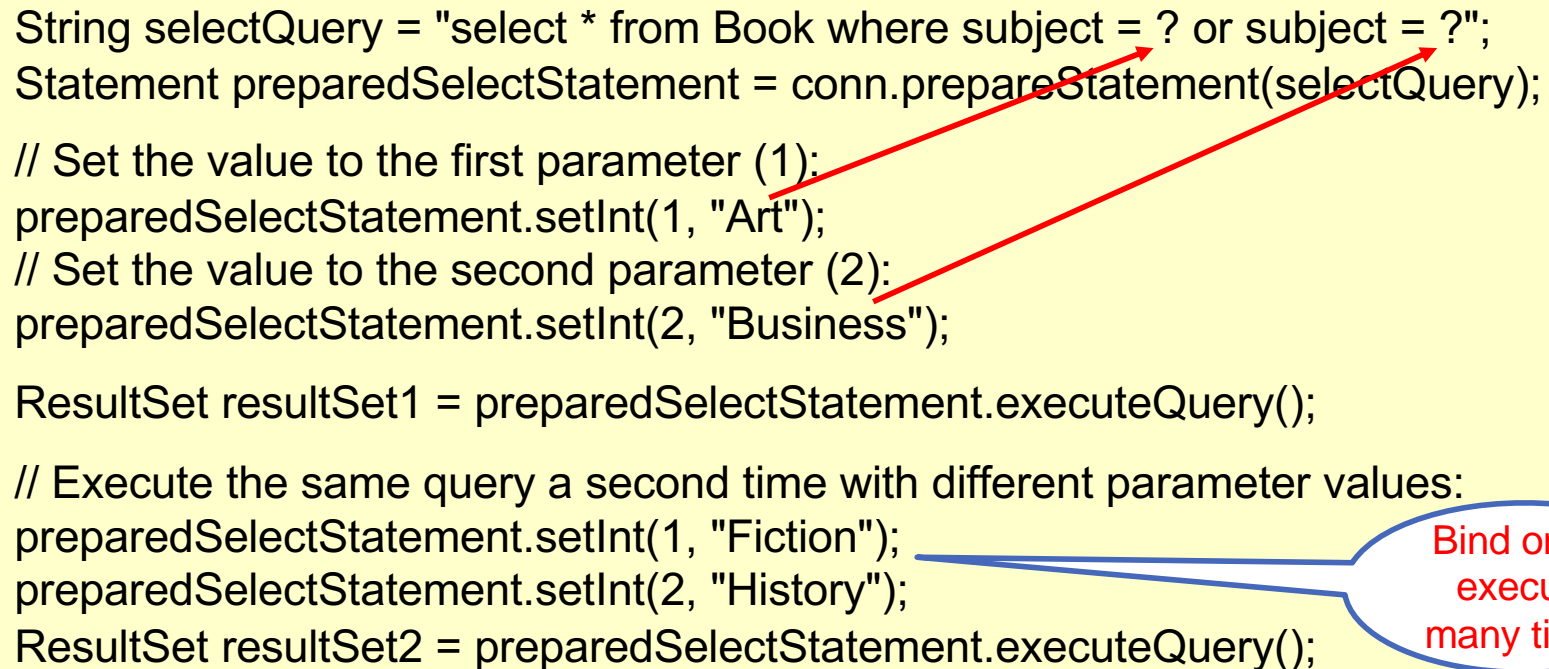
# JDBC: STORED PROCEDURES

```
String selectQuery = "select * from Book where subject = ? or subject = ?";
Statement preparedSelectStatement = conn.prepareStatement(selectQuery);

// Set the value to the first parameter (1):
preparedSelectStatement.setInt(1, "Art");
// Set the value to the second parameter (2):
preparedSelectStatement.setInt(2, "Business");

ResultSet resultSet1 = preparedSelectStatement.executeQuery();

// Execute the same query a second time with different parameter values:
preparedSelectStatement.setInt(1, "Fiction");
preparedSelectStatement.setInt(2, "History");
ResultSet resultSet2 = preparedSelectStatement.executeQuery();
```



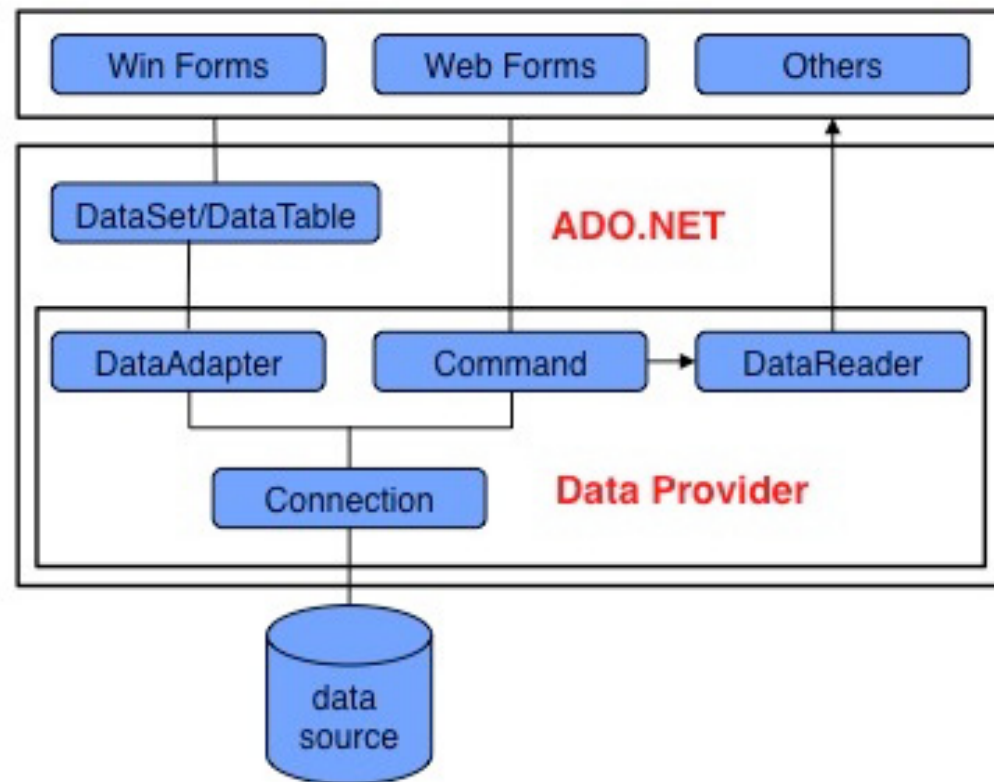
- The `prepareStatement` and `prepareCall` methods are used to create objects representing prepared (often parameterized) statements and stored procedure calls, respectively.

# JDBC: STORED PROCEDURES

- The `PreparedStatement` interface extends `Statement` with functionalities to bind a query once and then execute it multiple times in an efficient manner.
- Prepared statements also provide support for parameterized queries by passing query parameters, which are instantiated using setter methods such as `setInt`, `setString`, etc.
- Question marks (?) are used inside an SQL query to indicate that this represents a parameter value that will be bound later.
- `CallableStatement` extends `PreparedStatement` and offers support to execute **stored procedures**.

# ADO.NET: ARCHITECTURE

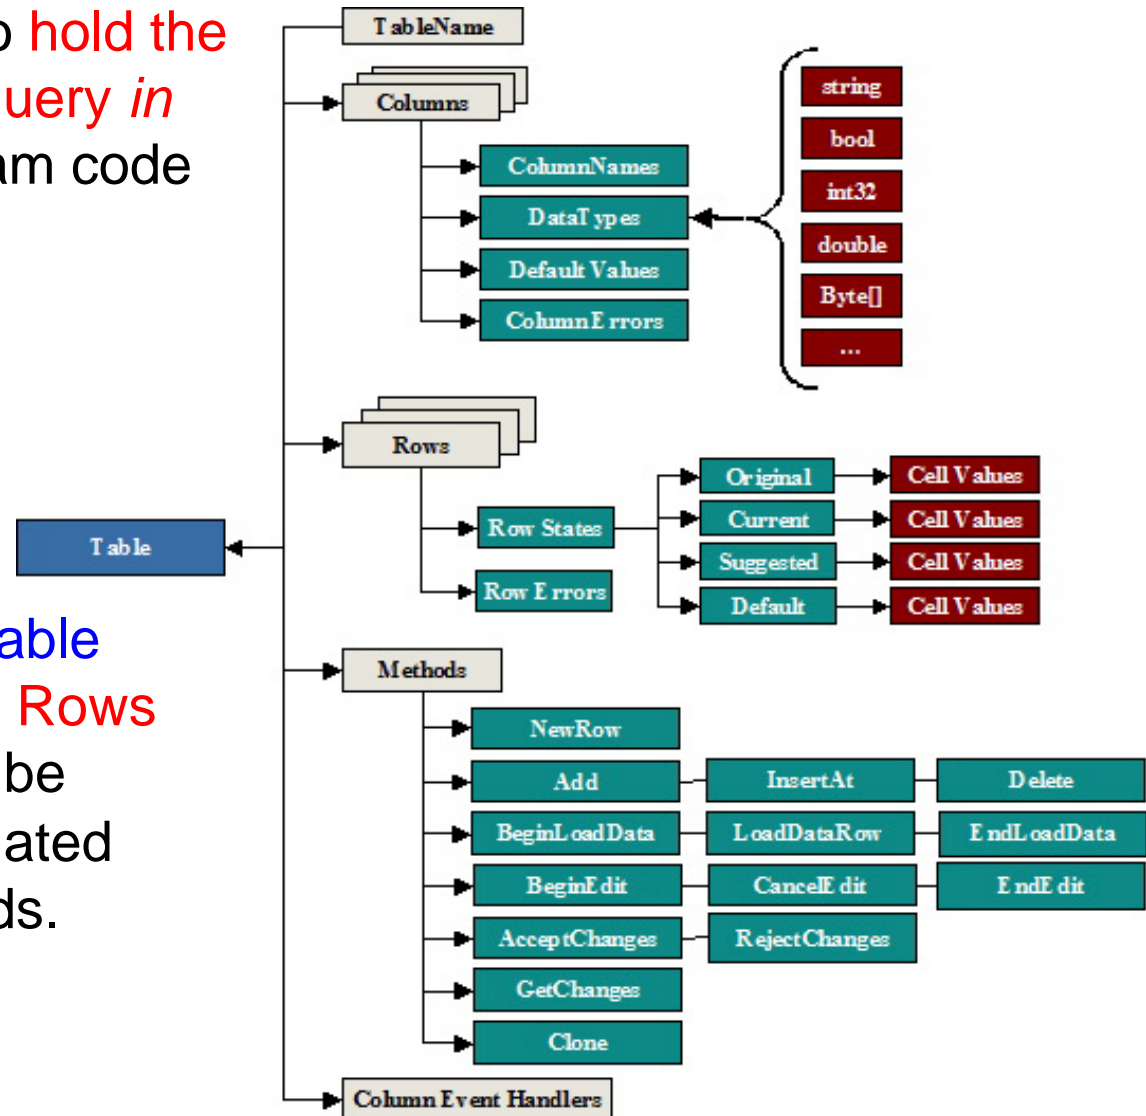
- ADO.NET offers a collection of **data providers**, which consist of objects that handle creation of database connections, sending queries and reading results.
- A **DataSet/DataTable** data structure provides a **disconnected** way to hold the data retrieved from a database.
- A **DataSet** can hold **several tables**, while a **DataTable** can hold **only one table**.





# ADO.NET: DATATABLE

- A **DataTable** is used to **hold the data resulting from a query in memory** so that program code can manipulate it.
- A **DataTable** can hold at most **one table** (i.e., a query result).
- A table within a **DataTable** contains **Columns** and **Rows** collections, which can be accessed and manipulated using standard methods.



# ADO.NET: EXAMPLE DATABASE ACCESS

```
String connectionString = "Data Source= ...";
OracleConnection conn = new OracleConnection(connectionString);
conn.Open();
string title;
string author;
DataTable dataTable = new DataTable();
OracleCommand cmd1 = conn.CreateCommand();
OracleCommand cmd2 = conn.CreateCommand();
cmd1.CommandText = "select trunc(avg(price),2) from Book";
cmd2.CommandText = "select title, firstName, lastName from Author natural join Book where subject='Art'";
decimal averagePrice = Convert.ToDecimal(cmd1.ExecuteScalar());
OracleDataAdapter da = new OracleDataAdapter(cmd2.CommandText, conn);
da.Fill(dataTable);
foreach (DataRow row in dataTable.Rows) {
    title = row["TITLE"].ToString();
    author = row["LASTNAME"].ToString() + ", " + row["FIRSTNAME"].ToString();
    Console.WriteLine(title + " by " + author + "\n");
}
Console.WriteLine("The average book price is " + averagePrice + ".");
conn.Close();
```

**OracleConnection**  
Set up a connection to the database.

**OracleCommand**  
Create the queries in the context of the connection.

Set the two queries.

Execute the two queries.

Loop over the result set.

This C# code fragment shows the **Connection**, **Command** and **DataTable** objects using the .NET Framework Data Provider for Oracle Database.

# ORACLE PL/SQL

- PL/SQL (Procedural Language/SQL) allows SQL statements to be embedded into a procedural programming language.
- It combines the data manipulation power of SQL with the data processing power of procedural languages.
- A **block**, which is delimited by **begin...end** and which can be nested, is the basic processing unit in which statements:
  - are **case insensitive**.
  - use C style comments **/\*...\*/**.
  - use **:=** operator to assign values to a variable.
  - use **=** operator for comparison.

Allowed SQL statements: **select**, **insert**, **update**, **delete** (i.e., DML)

Not allowed SQL statements: **create**, **drop**, **alter**, **rename** (i.e., DDL)

## ORACLE PL/SQL (CONT'D)

- A PL/SQL program is stored as a **database object** and can be
  - a **procedure**, which does not return a value.
  - a **function**, which returns a value using the **return** keyword.
- Both types of programs can accept parameters which can be one of
  - in** a read-only variable for giving input (the default)
  - out** a read-write variable for getting output
  - in out** a read-write variable for giving input and getting output
- A procedure is invoked using the **exec** keyword.
- A function is invoked by assigning its result to a variable or using it in a **select** statement.

# BASIC STRUCTURE & DATA TYPES

create or replace procedure *procedure\_name* [ as | is ]

**Declaration section:** contains declaration of variables, types, and local subprograms.

**begin** **Executable section:** contains procedural and SQL statements. This is the only section of a block that is required.

**exception** **Exception handling section:** contains error handling statements.

**end;**

## Variable Data Types

- A data type used to define the attributes of a table (i.e., **number**, **int**, **char**, **varchar2**, **date**, etc.).
- The same as an attribute (*table\_name.attribute\_name%type*) or a row (*table\_name%rowtype*).

# FLOW OF CONTROL STATEMENTS

- **Sequential control**

`goto` – branch to a label unconditionally

`null` – pass control to the next statement

`return` – returns control to the calling block and may return a value.

- **Conditional control**

`if-then`, `if-then-else`, `if-then-elsif` – conditional processing

`case` – selects one sequence of statements to execute

- **Iterative control**

`loop statements end loop`;

`while condition loop statements end loop`;

`for loop_variable in [reverse] lower_bound..upper_bound`  
`loop statements end loop`;

`exit` / `exit when condition` – exit the current loop possibly conditionally

`continue` / `continue when condition` – exit current loop iteration

# PL/SQL PROCEDURE EXAMPLE

**Increment the rating of a sailor if the rating is less than 5.**

```
create or replace procedure L10Example1 (sid in int) as
-- sailorName is the same type as sName in the Sailor table
sailorName Sailor.sName%type;
-- sailorRating is the same type as rating in the Sailor table
sailorRating Sailor.rating%type;
begin
-- Fetch the sailor's name and rating into the variables sailorName and sailorRating
select sName, rating into sailorName, sailorRating from Sailor where sailorId=sid;
if sailorRating<5 then
    update Sailor set rating=sailorRating+1 where sailorId=sid;
    -- Write record updated message to the Script Output pane
    dbms_output.put_line('Sailor ' || sailorName || '(' || sid || ') rating updated from ' ||
        sailorRating || ' to ' || (sailorRating+1) || '.');
else
    -- Write record NOT updated message to the Script Output pane
    dbms_output.put_line('Sailor ' || sailorName || '(' || sid || ') rating ' || sailorRating || ' NOT updated.');
```

```
end if;
end L10Example1;
```

# SELECT INTO STATEMENT

`select attribute_name into variable_name from table_name [where condition];`

- Retrieves a value from a table in the database and assigns it to *variable\_name*.
- The `select ... into` statement should **retrieve only one record** as a variable can hold only one value.
- If the `select ... into` statement returns more than one or no value, an exception will be raised => handle in the `exception` section.
- The number of columns and their data type in the `select` clause must match with the number of variables and their data types in the `into` clause.
- The values are retrieved and populated in the same order as specified in the `select` clause.



# CURSORS

- If a **select** statement returns more than one record, a cursor is normally used to **process the records one-at-a-time**.
- A cursor is like a pointer that **points to a single record** and allows access to the attribute values of that record.
- A cursor is defined in the **declare** section using the syntax:

**cursor** *cursor\_name* **is** *select\_statement*;

- A cursor can be used and managed **explicitly** using the **open**, **fetch** and **close** commands and by checking cursor status.
- It can also be used and managed **implicitly** using the **for...loop** statement where the *cursor\_name* replaces the range limit so the loop ranges from the first record of the cursor to the last record of the cursor.

# CURSOR STATUS

- The possible values of a cursor status are:

*cursor\_name%***found** Returns TRUE if the fetch operation succeeded; else returns FALSE.

*cursor\_name%***notfound** Returns TRUE if the fetch operation failed; else returns FALSE.

*cursor\_name%***isopen** Returns TRUE if the cursor is still open; else returns FALSE.

*cursor\_name%***rowcount** Returns the number of records fetched.

# PL/SQL CURSOR EXAMPLE

```
create or replace procedure L10Example2 as
currentSailorId Sailor.sailorId%type;
-- Declare the cursors for the sailor and reserves tables
cursor sailorCursor is select * from Sailor order by sName;
cursor reservesCursor is select count(boatid) reservations from reserves where sailorId=currentSailorId;
begin
-- Fetch the sailorCursor records one-by-one
for sailorRecord in sailorCursor loop
-- Assign the sailor id for the current sailor record
currentSailorId:=sailorRecord.sailorId;
-- Fetch the reservesCursor records one-by-one
for reservesRecord in reservesCursor loop
-- Insert into appropriate table
if reservesRecord.reservations=0 then
insert into NoReservations values (sailorRecord.sailorId, sailorRecord.sName);
else
insert into YesReservations values (sailorRecord.sailorId, sailorRecord.sName);
end if;
end loop;
end loop;
end L10Example2;
```

# PL/SQL EXCEPTIONS

- Predefined exceptions are raised implicitly by PL/SQL if the exception occurs.
- User-defined exceptions are declared in the declaration section,

*exception\_name* **exception**;

raised explicitly within a **begin...end** block

```
if condition then
  raise exception_name;
end if;
```

and handled in the **exception** section  
within the **begin...end** block.

```
exception
  when exception_name then
    ⋮
```

## Predefined Exceptions

ACCESS_INTO_NULL	ORA-06530
CASE_NOT_FOUND	ORA-06592
COLLECTION_IS_NULL	ORA-06531
CURSOR_ALREADY_OPEN	ORA-06511
DUP_VAL_ON_INDEX	ORA-00001
INVALID_CURSOR	ORA-01001
INVALID_NUMBER	ORA-01722
LOGIN_DENIED	ORA-01017
NO_DATA_FOUND	ORA-01403
NOT_LOGGED_ON	ORA-01012
PROGRAM_ERROR	ORA-06501
ROWTYPE_MISMATCH	ORA-06504
SELF_IS_NULL	ORA-30625
STORAGE_ERROR	ORA-06500
SUBSCRIPT_BEYOND_COUNT	ORA-06533
SUBSCRIPT_OUTSIDE_LIMIT	ORA-06532
SYS_INVALID_ROWID	ORA-01410
TIMEOUT_ON_RESOURCE	ORA-00051
TOO_MANY_ROWS	ORA-01422
VALUE_ERROR	ORA-06502
ZERO_DIVIDE	ORA-01476

# PL/SQL EXCEPTIONS EXAMPLE

```
create or replace procedure L10Example3 (sid in int) as
  -- sailorName is the same type as sName in the Sailor table
  sailorName Sailor.sName%type;
  -- sailorRating is the same type as rating in the Sailor table
  sailorRating Sailor.rating%type;
begin
  -- Fetch the sailor's name and rating into the variables sailorName and sailorRating
  select sName, rating into sailorName, sailorRating from Sailor where sailorId=sid;
  if sailorRating<5 then
    update Sailor set rating=sailorRating+1 where sailorId=sid;
    -- Write record updated message to the Script Output pane
    dbms_output.put_line('Sailor ' || sailorName || '(' || sid || ') rating updated from ' ||
      sailorRating || ' to ' || (sailorRating+1) || '.');
  else
    -- Write record NOT updated message to the Script Output pane
    dbms_output.put_line('Sailor ' || sailorName || '(' || sid || ') rating ' || sailorRating || ' NOT updated.');
```

end if;

```
exception
  when no_data_found then
    -- Write exception message to the Script Output pane
    dbms_output.put_line('There is no sailor with id ' || sid || '.');
end L10Example3;
```

# STRUCTURED QUERY LANGUAGE (SQL): SUMMARY

- Structured Query Language (SQL) is a relational query language that provides facilities to

## Query Relations

- Select-From-Where Statement
- Set Operations (Union, Intersect, Except)
- Nested Subqueries (to test for set membership, comparison, cardinality)
- Aggregate Functions (avg, min, max, sum, count)
- Group By with Having clause

## Create and Modify Relations

- Create, Alter, Drop Tables
- Specify integrity constraints: domain, key, foreign key, general
- Specify views
- Insert, Delete, Update Tuples

## Access a Database from a Programming Language