

# COMP 3311

# DATABASE MANAGEMENT

# SYSTEMS

## LECTURE 13

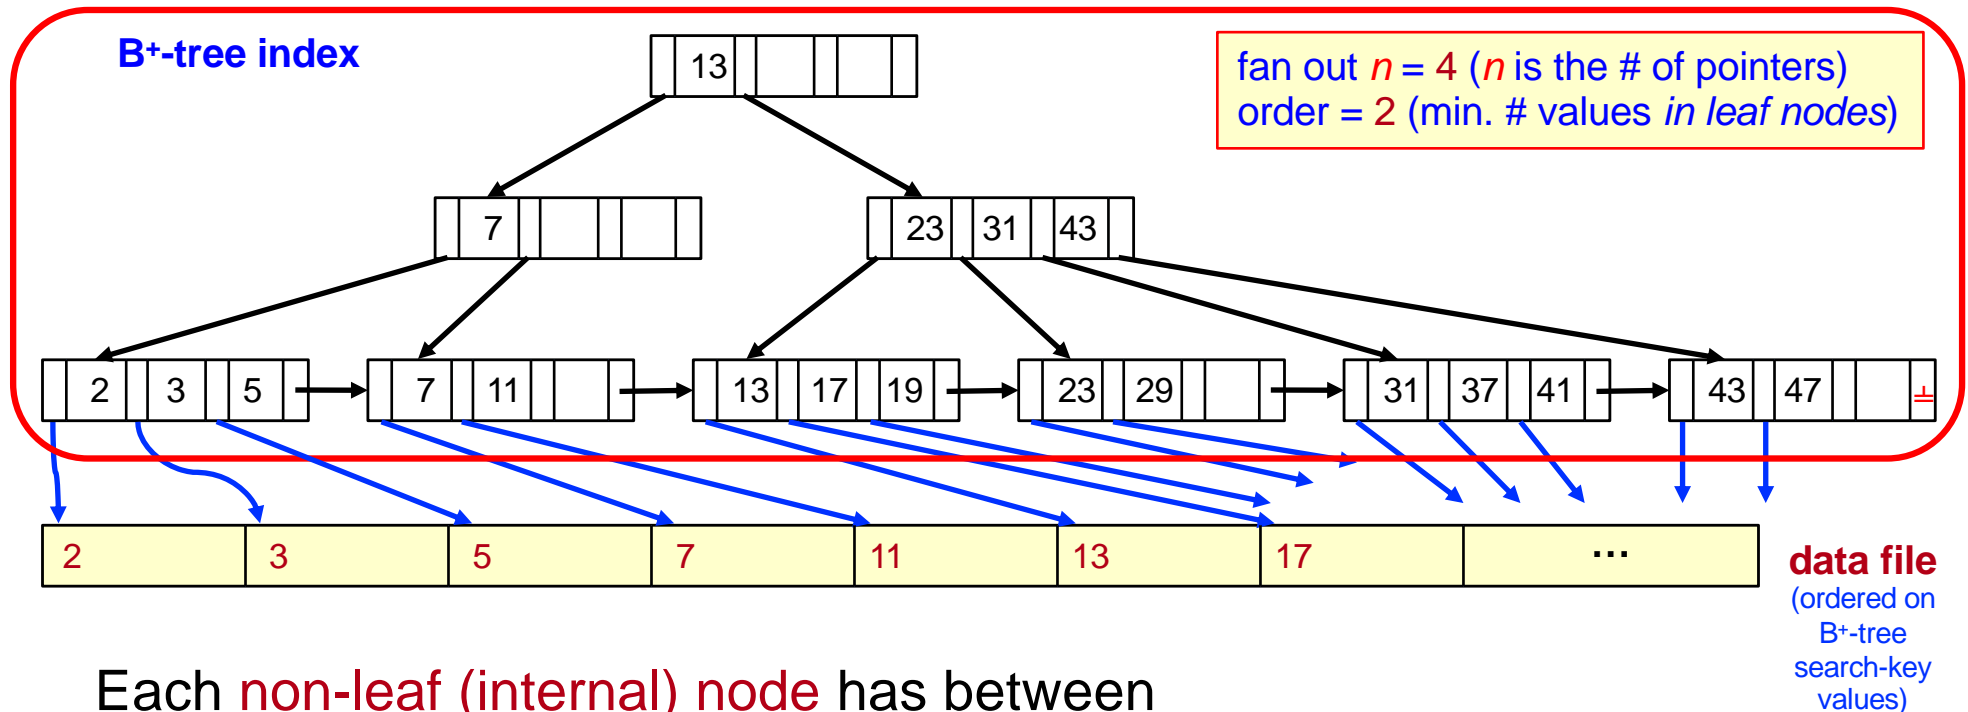
## INDEXING: B+-TREE INDEX

# B<sup>+</sup>-TREE INDEX

- **Disadvantage of index-sequential files**
  - Performance degrades as the file grows due to overflow pages.
  - Periodic reorganization of the entire file is required.
- **Advantage of B<sup>+</sup>-tree index files**
  - Balanced tree  $\Rightarrow$  every path from root to leaves is the same length.
  - Automatically reorganizes itself with small, local changes in the presence of insertions and deletions.
  - Reorganization of entire file is not required to maintain performance.
- **Disadvantage of B<sup>+</sup>-trees**
  - Extra insertion, deletion and space overhead.

**The advantages of B<sup>+</sup>-trees far outweigh their disadvantages.  
They are used extensively in all commercial products.**

# B<sup>+</sup>-TREE: STRUCTURE



Each **non-leaf (internal) node** has between

- $\lceil n/2 \rceil$  and  $n$  pointers ( $\lceil n/2 \rceil - 1$  (min) and  $n - 1$  (max) values).
- **In the example:** between 2 and 4 pointers (1 and 3 values).

Special cases for root:  
if **non-leaf** - min 1 value.  
if **leaf** - min 0 values.

Each **leaf node** is at least half full, i.e., has between

- $\lceil (n-1)/2 \rceil + 1$  and  $n$  pointers, e.g., ( $\lceil (n-1)/2 \rceil$  (min) and  $n - 1$  (max) values).
- **In the example:** between 3 and 4 pointers (2 and 3 values).

## B<sup>+</sup>-TREE: STRUCTURE (cont'd)

- **Balanced tree**
  - All **paths** from the root node to the leaf nodes are the **same length**.
- **Fan-out**
  - The **maximum number of pointers/children** in each node, denoted ***n***.
- **B<sup>+</sup>-tree order**
  - The value  $\lceil (n-1)/2 \rceil$  is called the **order** and corresponds to the **minimum number of values in a leaf node**.
- **Special cases**
  - If the **root is not a leaf**, it has **at least 2 children** (i.e., it has **at least 1 value**).
  - If the **root is a leaf** (i.e., there are no other nodes in the tree), it can have **between 0 and (*n*-1) values**.

# B<sup>+</sup>-TREE: NON-LEAF (INTERNAL) NODES


- Non-leaf (internal) nodes form a **multi-level, sparse index** on the leaf nodes (i.e., only some search-key values are present).

- For a non-leaf node with  **$n$  pointers**:

**First pointer  $P_1$ :** All the search-key values in the subtree to which  $P_1$  points are **less than  $K_1$** .

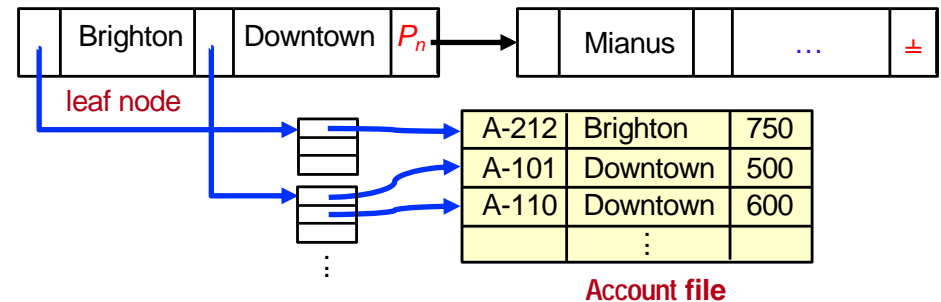
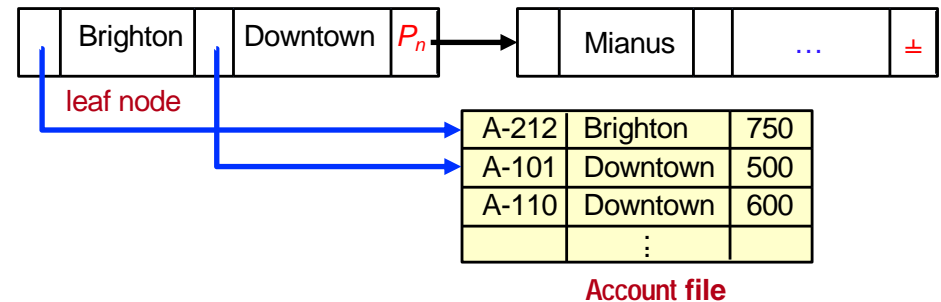
**Internal pointer:** For  $2 \leq i \leq n-1$ , all the search-key values in the subtree to which  $P_i$  points have values **greater than or equal to  $K_{i-1}$  and less than  $K_i$** .

**Last pointer  $P_n$ :** All the search-key values in the subtree to which  $P_n$  points are **greater than  $K_{n-1}$** .



# B<sup>+</sup>-TREE: LEAF NODES

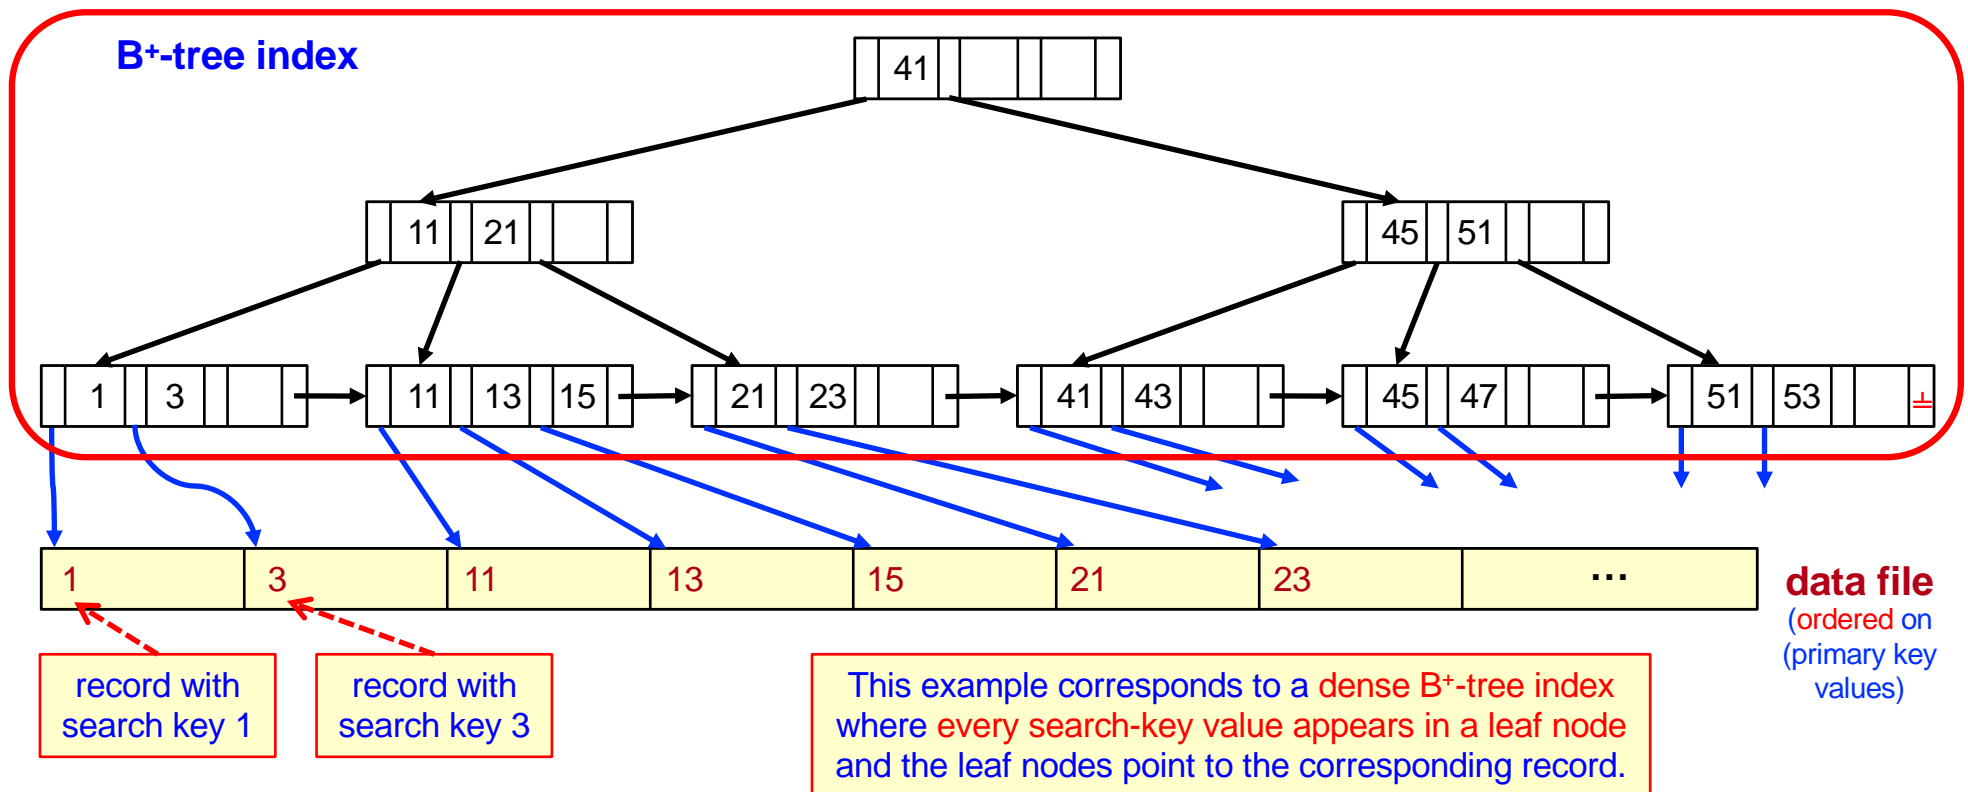
- For  $i = 1, 2, \dots, n-1$ ,  
pointer  $P_i$  either points
  - to a file record  
with search-key value  $K_i$ ,
  - or
  - to a “bucket” of pointers each  
of which points to a file record  
that has search-key value  $K_i$ .
- Search-key values in a node  
are kept in sorted order.
  - If  $L_i$  and  $L_j$  are leaf nodes and  $i < j$ ,  $L_i$ 's search-key values are less  
than  $L_j$ 's search-key value.
- The last pointer in a node,  $P_n$ , points to the next leaf node in the  
search-key order (right sibling node) or contains an end symbol.



## B<sup>+</sup>-TREE: OBSERVATIONS

- Since the inter-node connections are given by pointers, the “close” pages need not be “physically” close (i.e., no need for sequential storage).
- The non-leaf levels of the B<sup>+</sup>-tree form a sparse index (i.e., not every search-key value is present in the index).
- Since a B<sup>+</sup>-tree contains a relatively small number of levels (logarithmic in the size of the data file), search can be conducted efficiently.
- Insertions and deletions to the data file can be handled efficiently, as the index can be restructured in logarithmic time.

# EXAMPLE CLUSTERING B<sup>+</sup>-TREE ON PRIMARY KEY

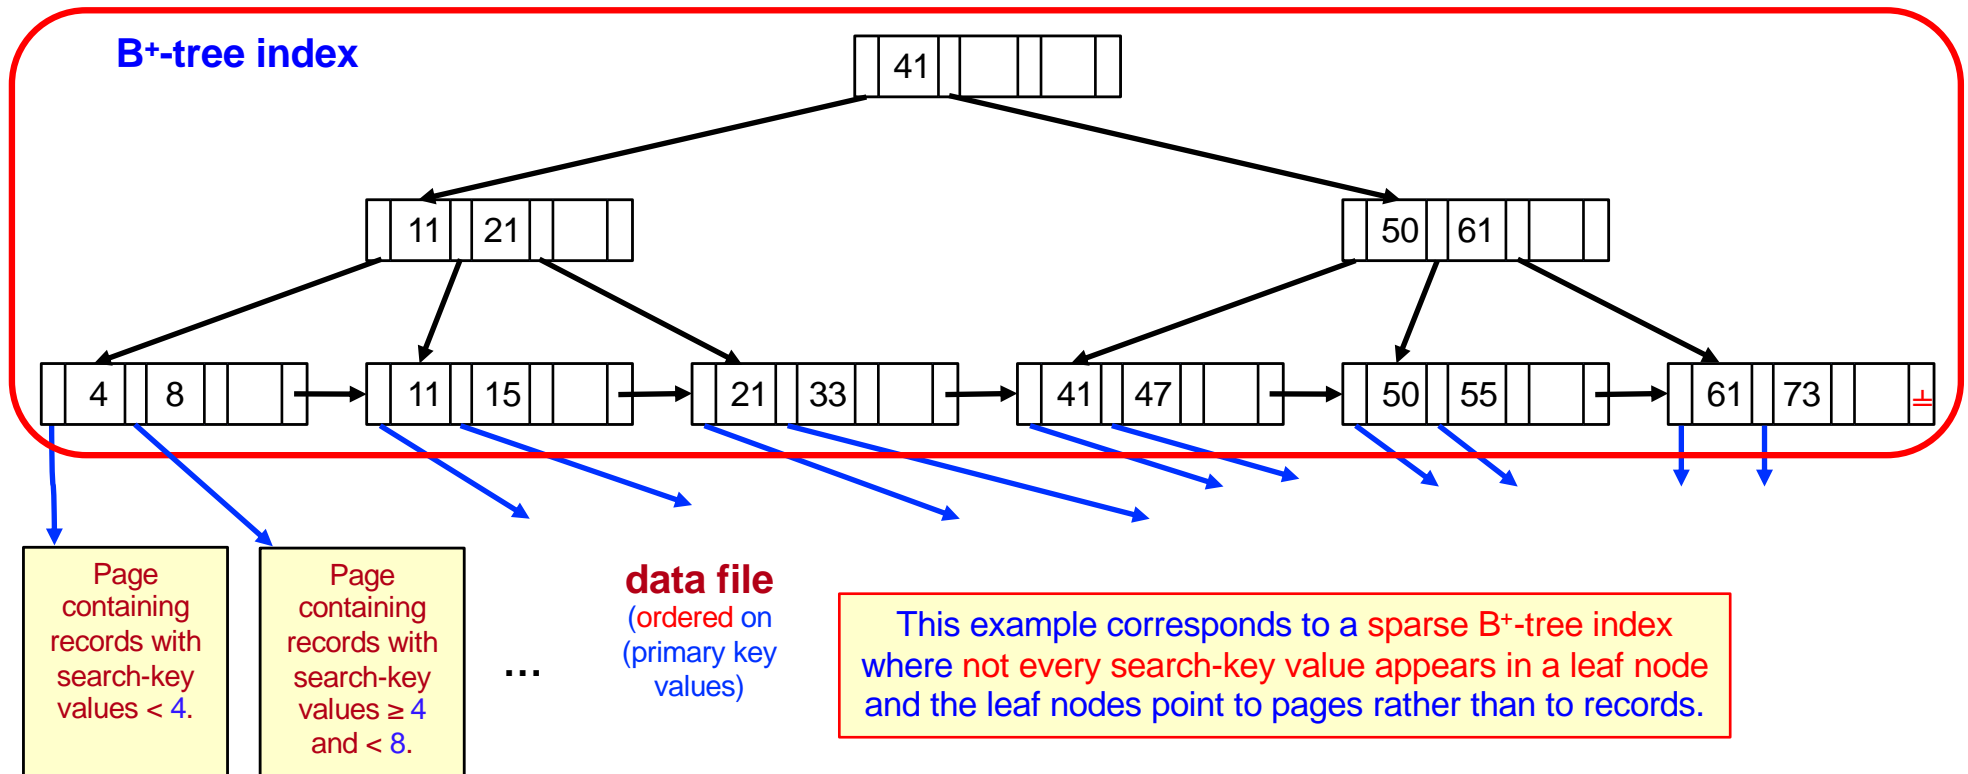


**Leaf nodes:** contain all the **primary key values** (dense) and **point to the record** in the file with that value.

**Data file:** records are ordered by the primary key.



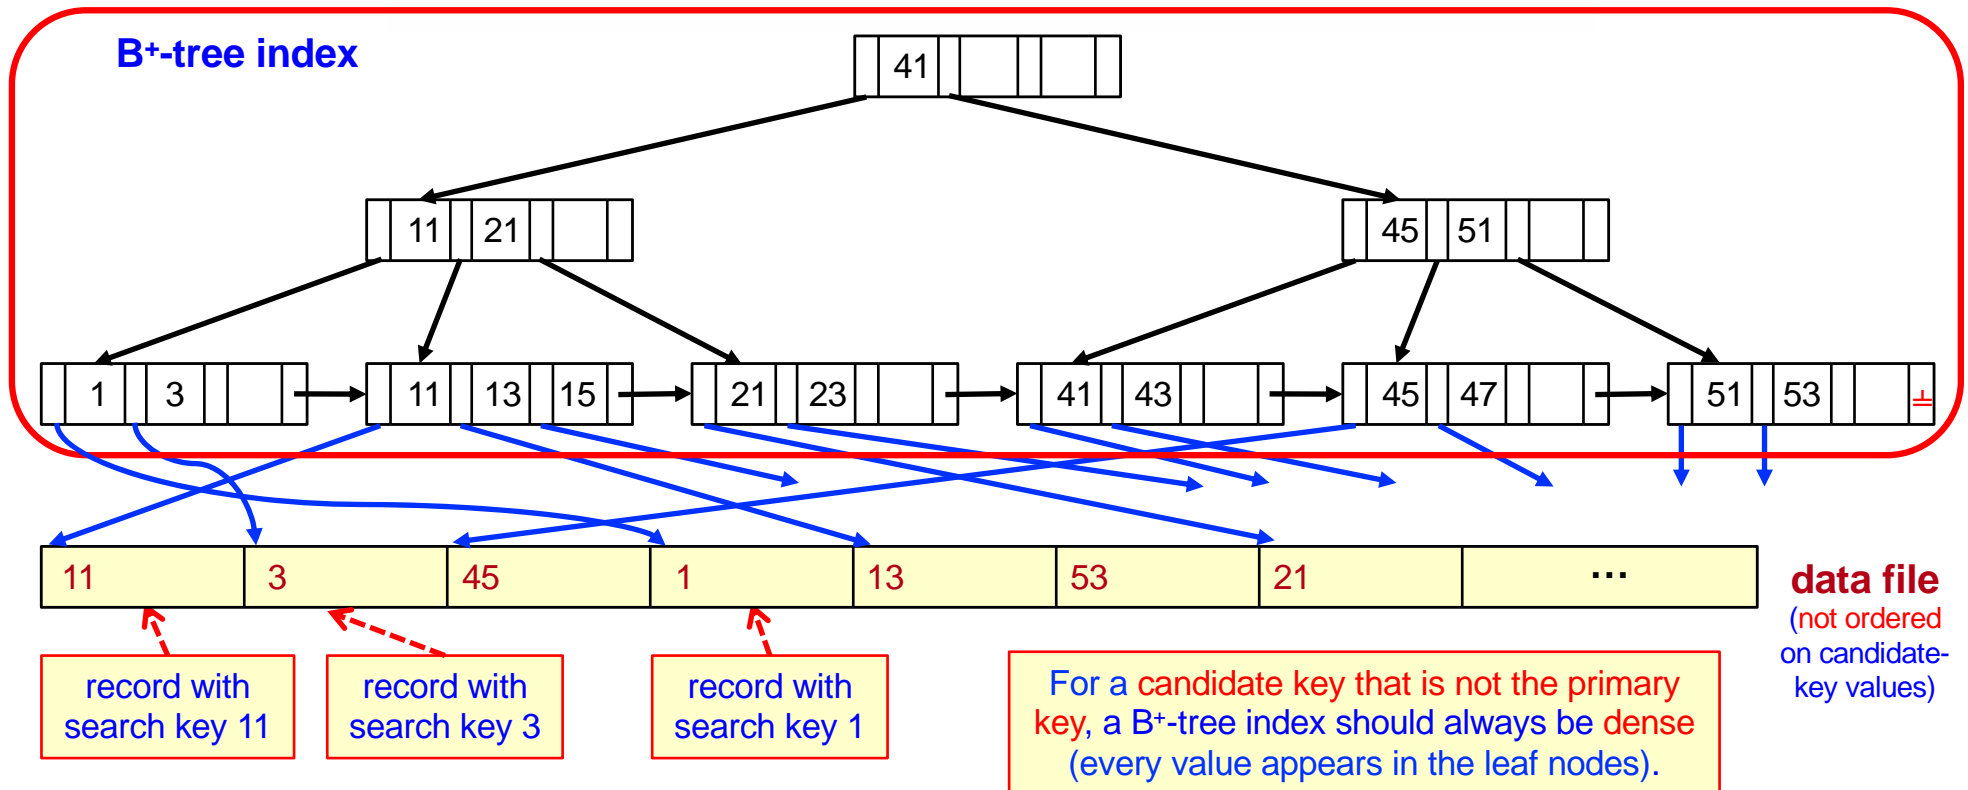
# EXAMPLE CLUSTERING B<sup>+</sup>-TREE ON PRIMARY KEY



**Leaf nodes:** contain only some of the **primary key values** (**sparse**) and **point to a page** with several records.

**Data file:** records are ordered by the **primary key**.

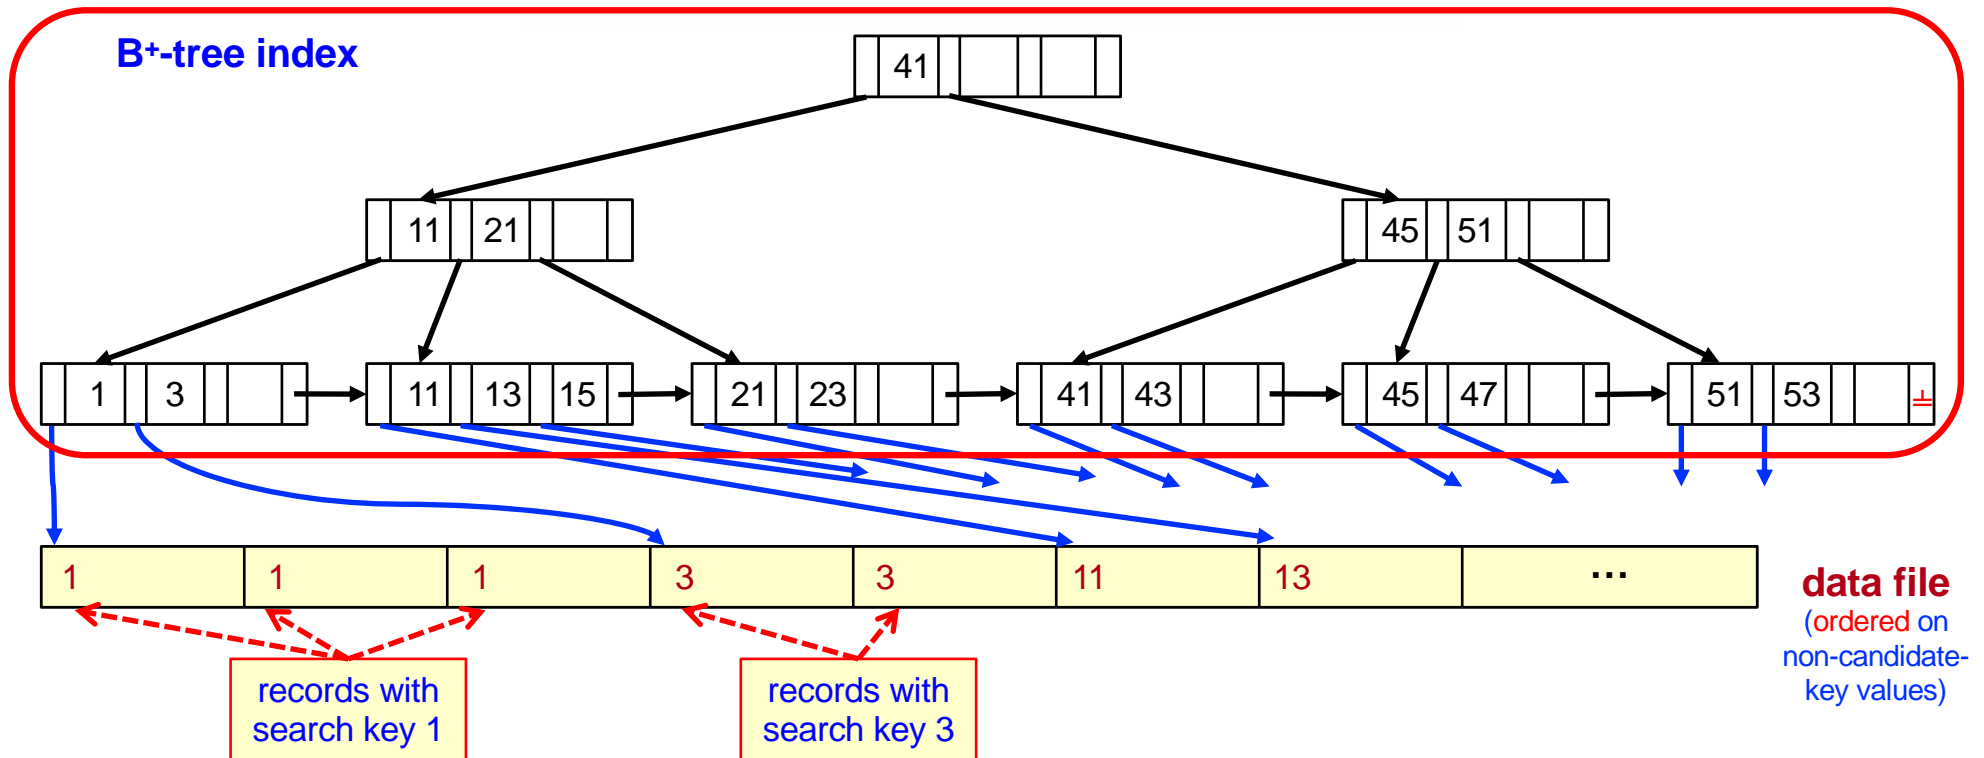
# EXAMPLE NON-CLUSTERING/SECONDARY B<sup>+</sup>-TREE ON CANDIDATE KEY



**Leaf nodes:** contain all the candidate-key values (dense) and point to the record in the file with that value.

**Data file:** records are not ordered by the candidate key.

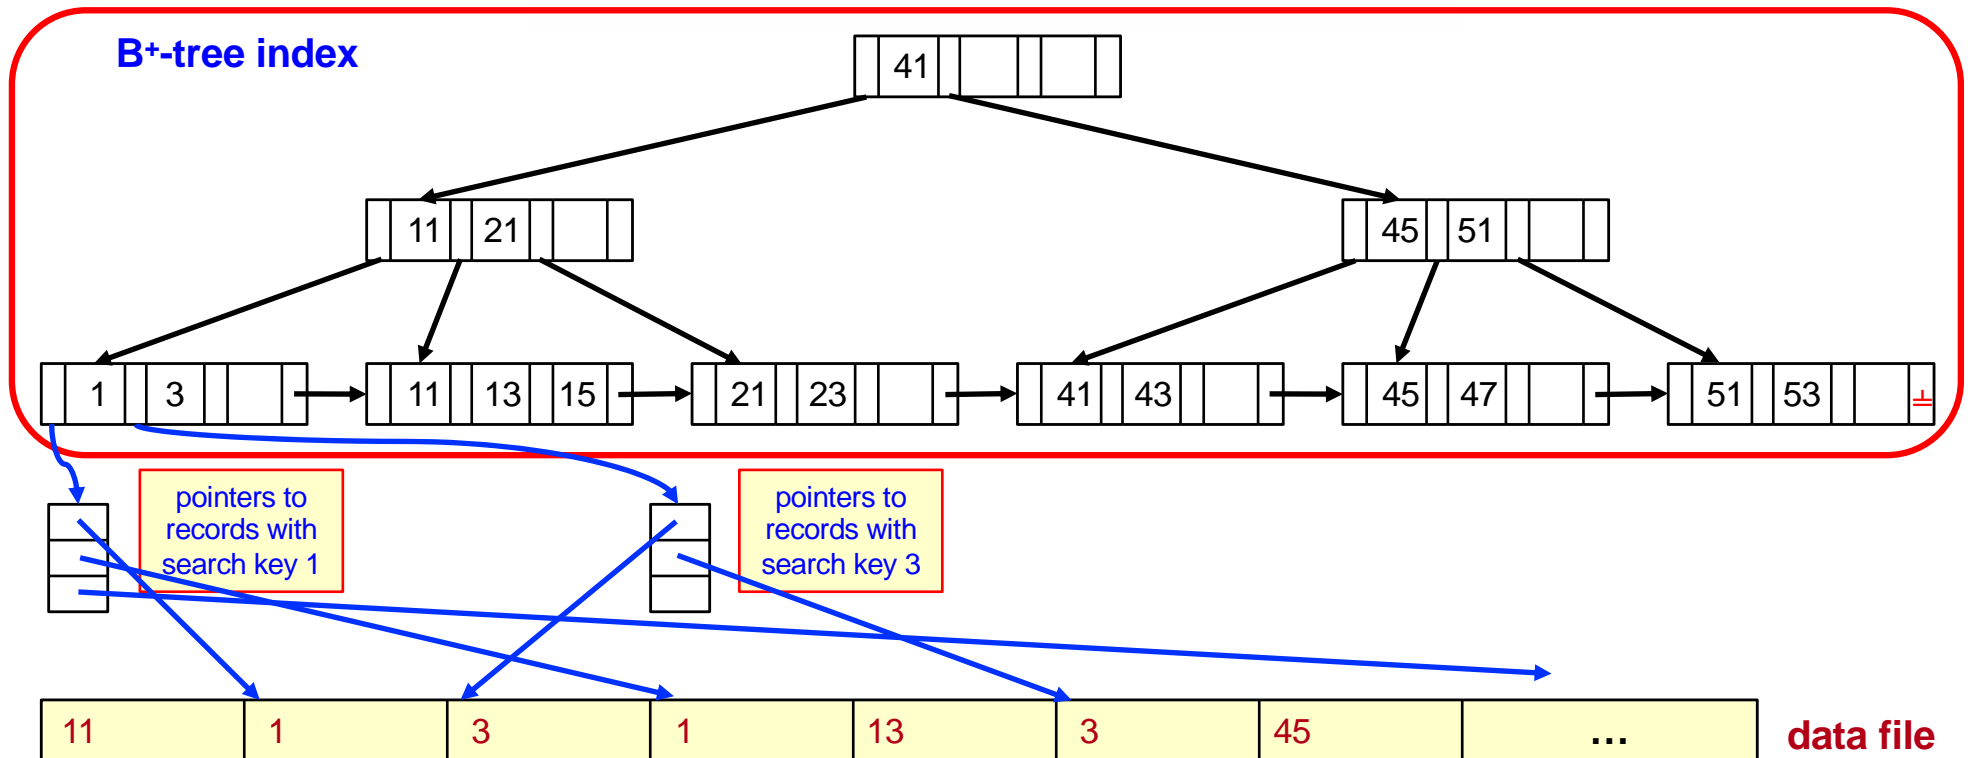
# EXAMPLE CLUSTERING B<sup>+</sup>-TREE ON NON-CANDIDATE KEY



**Leaf nodes:** contain all the unique non-candidate-key values (dense) and point to the first record in the file with that value.

**Data file:** records are ordered by the non-candidate key, which may be duplicated in different records.

# EXAMPLE NON-CLUSTERING B<sup>+</sup>-TREE ON NON-CANDIDATE KEY



**Leaf nodes:** contain all the unique non-candidate-key values (dense) which point to a list of pointers that point to the records in the file with that value.

**Data file:** records are not ordered by the non-candidate key.

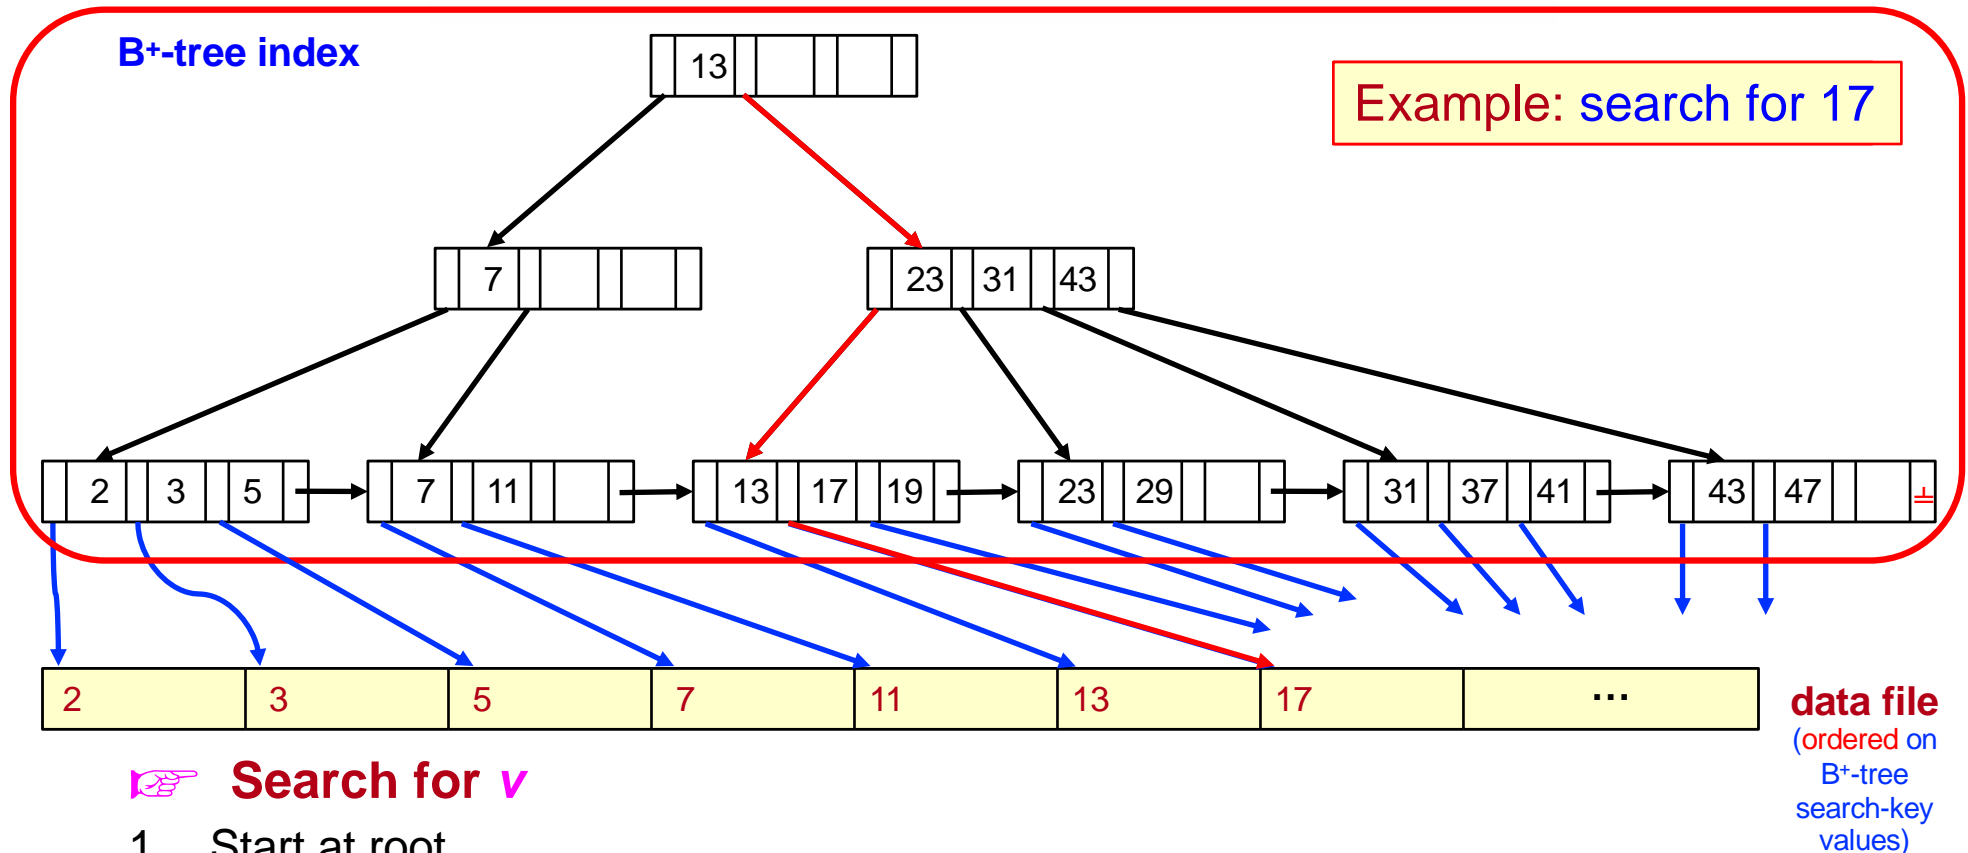
## B<sup>+</sup>-TREE: QUERY

- Find all records with a search-key value of  $v$ .
  1. Start at the root node.
    - If there is an entry with search-key value  $K_j = v$ , follow pointer  $P_{j+1}$ .
    - Otherwise, if  $v < K_{n-1}$  (there are  $n$  pointers in the node, i.e.,  $v$  is not larger than all values in the node) follow pointer  $P_j$ , where  $K_j$  is the smallest search-key value  $> v$ .
    - Otherwise, if  $v \geq K_{n-1}$ , follow  $P_n$  to the child node.
  2. If the node reached by following the pointer in step 1 is not a leaf node, repeat step 1 on the node, and follow the corresponding pointer.
  3. Reached a leaf node.
    - If for some  $i$ , key  $K_i = v$  follow pointer  $P_i$  to the desired record or list of pointers to records.
    - Else no record with search-key value  $v$  exists.

## B<sup>+</sup>-TREE: QUERY (cont'd)

- In processing a query, a **path** is traversed in the tree from the **root to some leaf node**.
- If there are  $K$  search-key values in the file, the path is no longer than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$ .
  - Assuming each node holds the minimum number of values.
- A node is generally the same size as a disk page, typically 4KB, and  $n$  is typically around 100 (assuming 40 bytes per index entry).
- With 1 million search-key values and  $n = 100$ , at most  $\log_{50}(1,000,000) = 4$  nodes are accessed in a search.

# B<sup>+</sup>-TREE: QUERY EXAMPLE



## 👉 Search for $v$

1. Start at root.
2. Find  $K_i$  such that  $K_i > v$ ; *follow left pointer*; if no  $K_i > v$ , *follow last pointer* in node.
3. At each node, repeat step 2 until a leaf node is reached.
4. At leaf node:
  - If  $v$  is found, *follow left pointer* to the record or list of pointers to records.
  - Else, if  $v$  is not found, then no record with the search-key value  $v$  exists.

## B<sup>+</sup>-TREE UPDATE

- When a record is inserted or deleted from a table, indexes on the table must be updated.
- Updates to a record can be modeled as a deletion of the old record followed by an insertion of the updated record.

 **Only need to consider B<sup>+</sup>-tree insertion and deletion operations.**

- To maintain the B<sup>+</sup>-tree properties, it may be necessary to
  - **split a node** when doing an insertion (**node overfull**).
  - **coalesce (combine) nodes** when doing a deletion (**node underfull**).

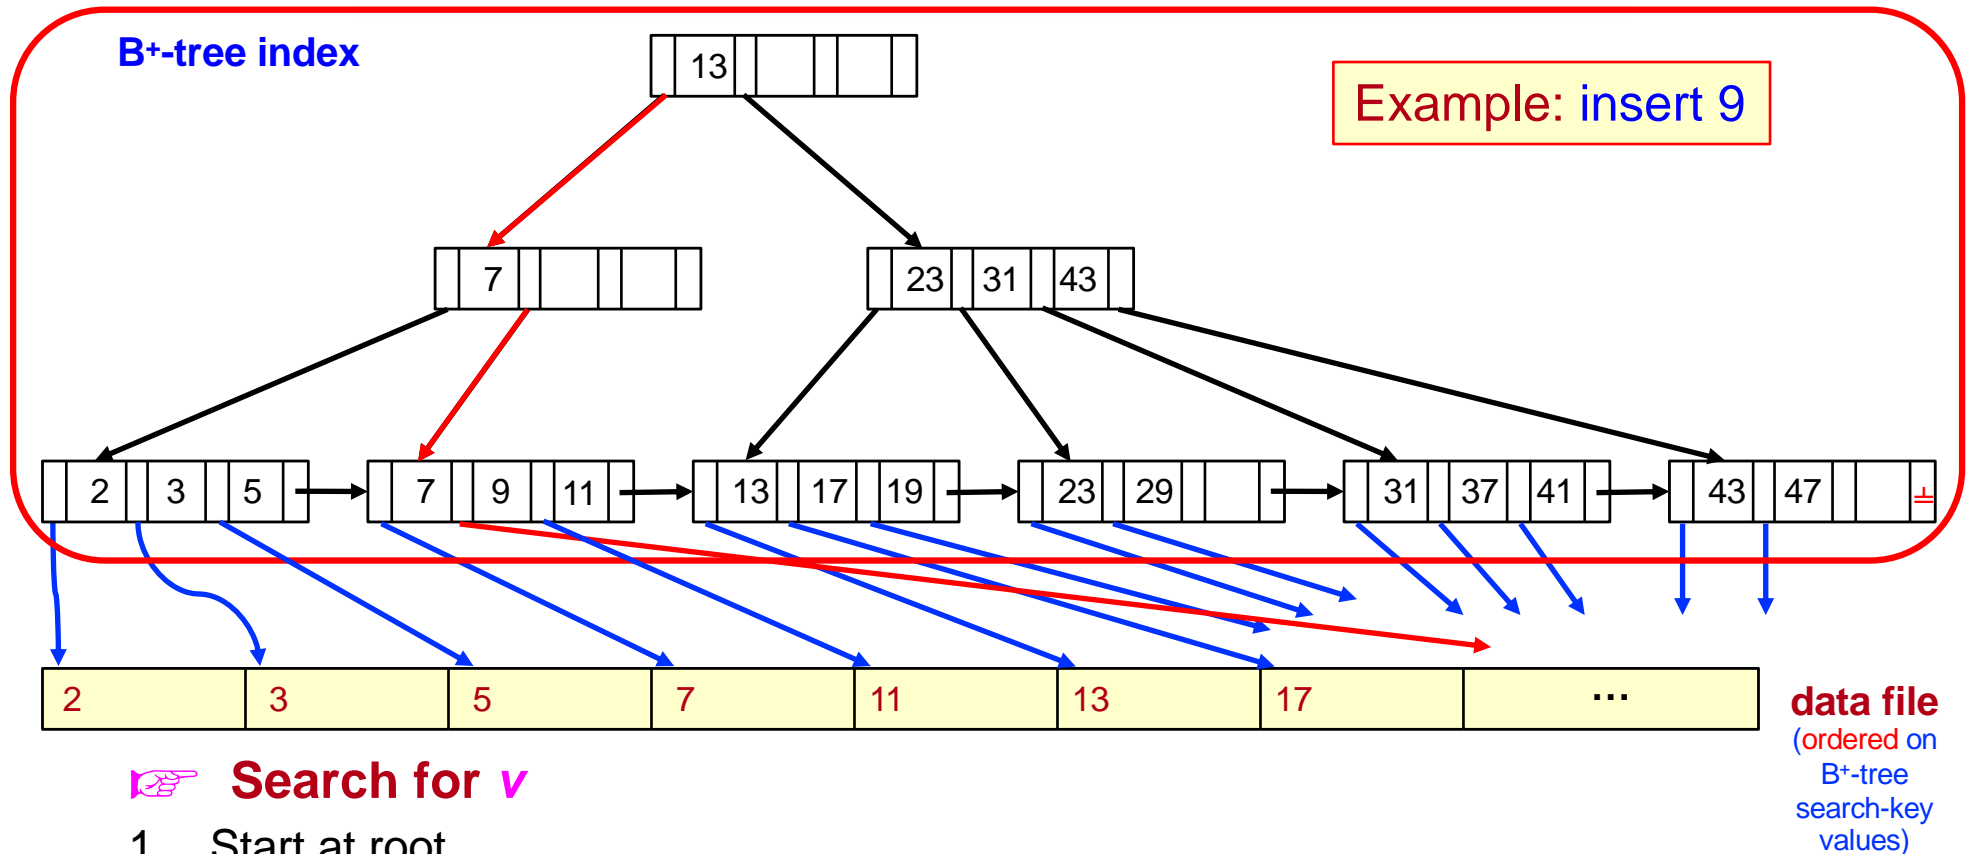
 **Must ensure that balance is preserved when doing an update.**



## B<sup>+</sup>-TREE INSERTION

- Use query strategy to find the leaf node  $L$  where the value belongs.
- Put the value into  $L$  (a leaf node).
  - If  $L$  has enough space, *done!*
  - Else (*leaf node overfull*), must split  $L$  (into  $L$  and a new node leaf node  $L'$ ).
    - Place first  $\lceil n/2 \rceil$  values into  $L$ , copy up value at  $\lceil n/2 \rceil + 1$  and place values starting at  $\lceil n/2 \rceil + 1$  up to and including the last value into  $L'$ .
    - Copy up: Insert an index entry (value at  $\lceil n/2 \rceil + 1$ , pointer to  $L'$ ) into the correct position in the (non-leaf) parent of  $L$ .
- Splits can happen recursively.
  - To split an index (*internal*) node  $N$  (into  $N$  and a new internal node  $N'$ ).
    - Place first  $\lceil n/2 \rceil - 1$  values into  $N$ , push up value at  $\lceil n/2 \rceil$  and place values starting at  $\lceil n/2 \rceil + 1$  up to and including the last value into  $N'$ .
    - Push up: Insert an index entry (value at  $\lceil n/2 \rceil$ , pointer to  $N'$ ) into the correct position in the parent of  $N$ .
- Splits “grow” tree; a root split increases the height.
  - B<sup>+</sup>-tree growth: The tree gets wider or one level taller at the top.

# B<sup>+</sup>-TREE: INSERTION EXAMPLE



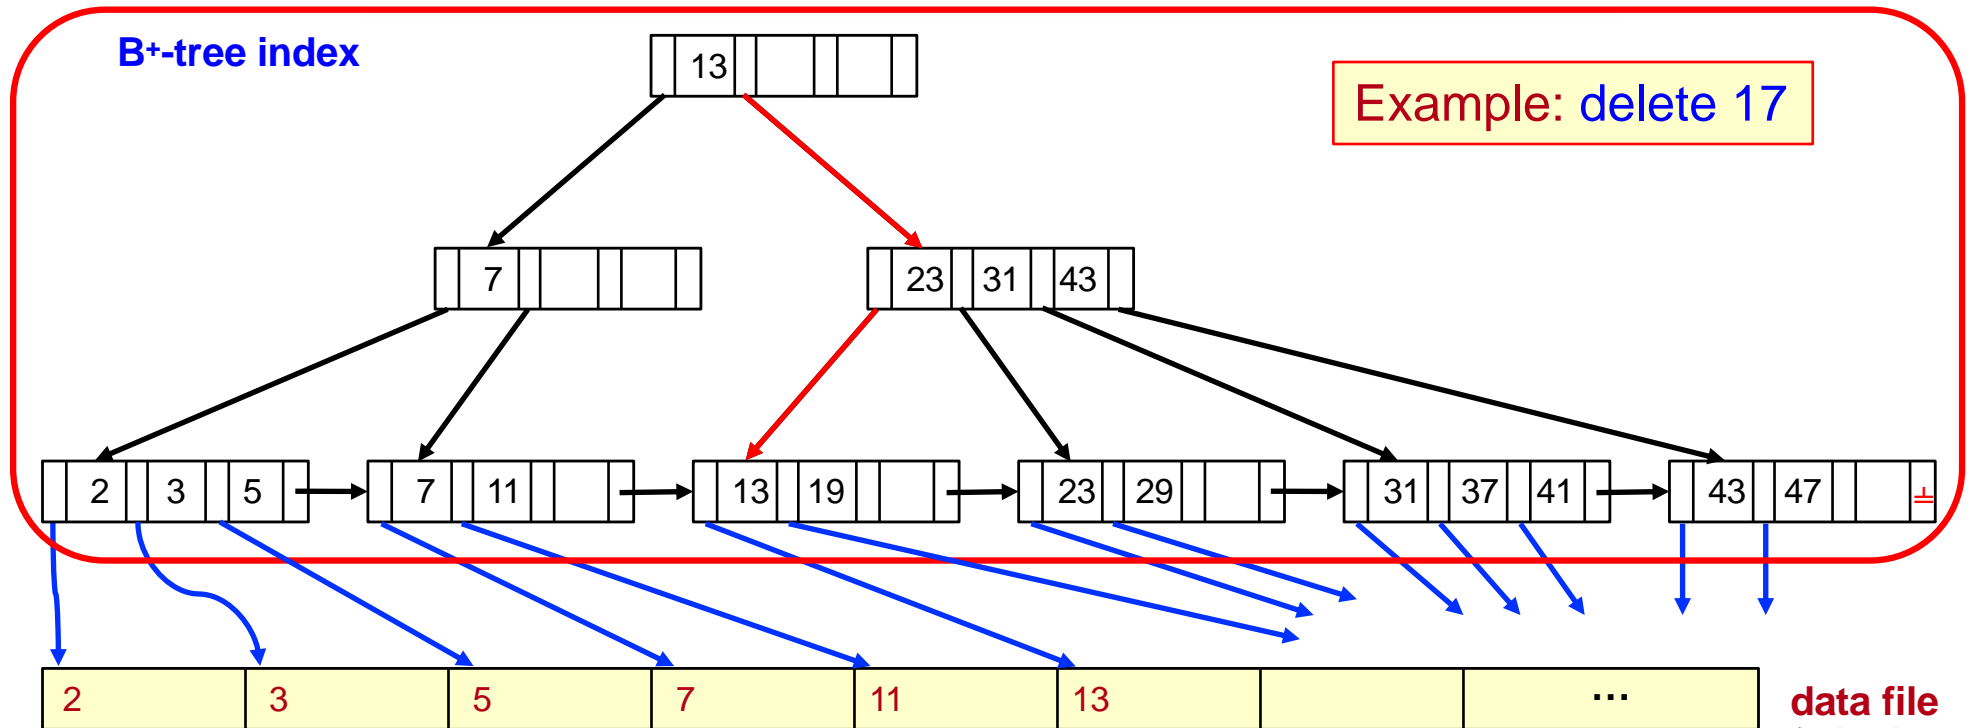
## 👉 Search for $v$

1. Start at root.
2. Find  $K_i$  such that  $K_i > v$ ; *follow left pointer*; if no  $K_i > v$ , *follow last pointer* in node.
3. At each node, repeat step 2 until a leaf node is reached.
4. At leaf node:
  - If enough space, *insert in node in order*.
  - Else, (*leaf node overfull*), must *split*.

## B<sup>+</sup>-TREE DELETION

- Use query strategy to find the leaf node  $L$  where the value belongs.
- Remove the leaf node entry.
  - If  $L$  has at least  $\lceil (n-1)/2 \rceil$  values, *done!*
  - Else  $L$  has less than  $\lceil (n-1)/2 \rceil$  values (*leaf node underfull*),
    - Try to re-distribute by borrowing values from a sibling node (an adjacent node to the right or the left *under the same parent*).
    - If re-distribution fails, merge  $L$  and a sibling.
- If a merge occurred, delete the entry (pointing to  $L$  or the sibling) from parent (non-leaf node) of  $L$ .
- A merge could propagate to the root, decreasing the height of the tree.
  - For non-leaf (internal) nodes, redistribution requires that a node have at least  $\lceil n/2 \rceil$  pointers.
  - If this criterion cannot be met by re-distribution, then nodes must be merged.

# B+-TREE: DELETION EXAMPLE



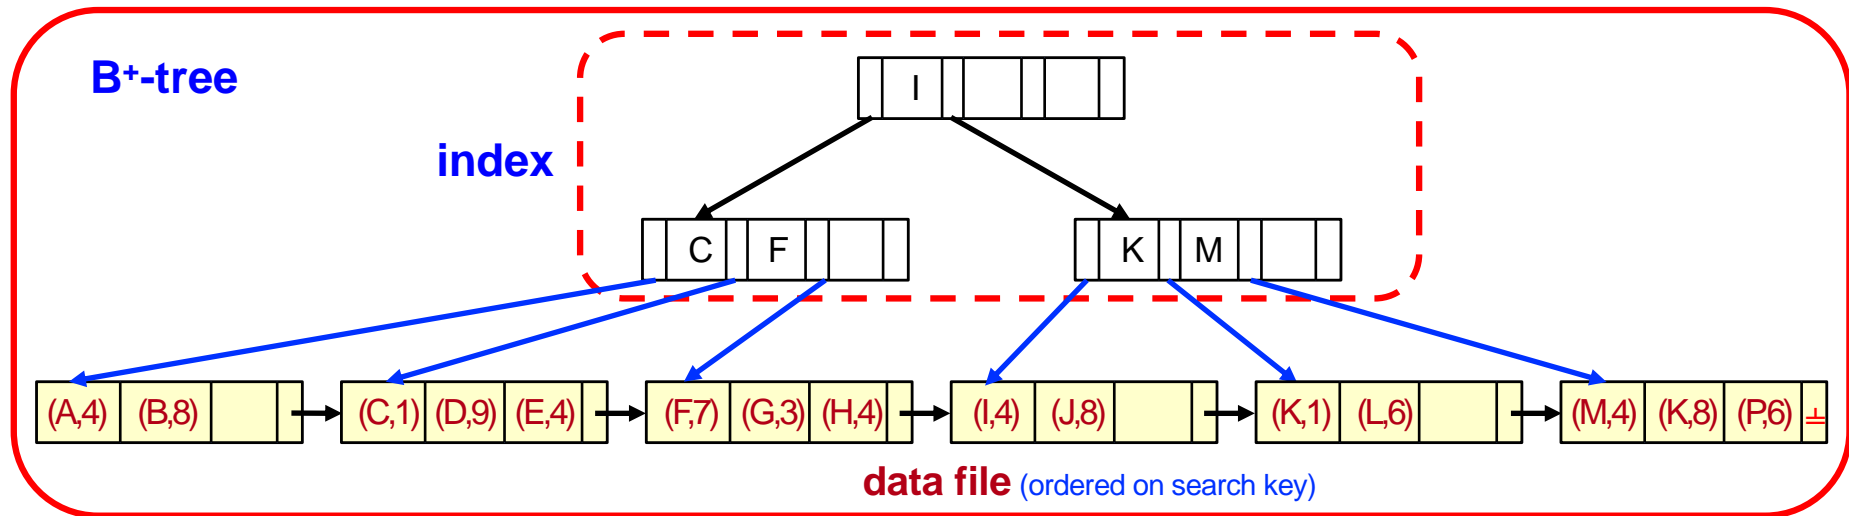
## 👉 Search for $v$

1. Start at root.
2. Find  $K_i$  such that  $K_i > v$ ; *follow left pointer*; if no  $K_i > v$ , *follow last pointer* in node.
3. At each node, repeat step 2 until a leaf node is reached.
4. At leaf node:
  - If  $L$  has at least  $\lceil (n-1)/2 \rceil$  values, *done!*.
  - Else,  $L$  has less than  $\lceil (n-1)/2 \rceil$  values (*leaf node underfull*).

## B<sup>+</sup>-TREE FILE ORGANIZATION

- The index file degradation problem is solved by using B<sup>+</sup>-tree indexes.
- The data file degradation problem can be solved by using B<sup>+</sup>-tree file organization.
- The leaf nodes in a **B<sup>+</sup>-tree file organization** store records, instead of pointers.
- Since records are larger than pointers, the maximum number of records that can be stored in a leaf node is less than the number of pointers in a non-leaf node.
- Leaf nodes are still required to be half full.
- Insertion and deletion are handled in the same way as insertion and deletion of entries in a B<sup>+</sup>-tree index.

# B<sup>+</sup>-TREE FILE ORGANIZATION (cont'd)

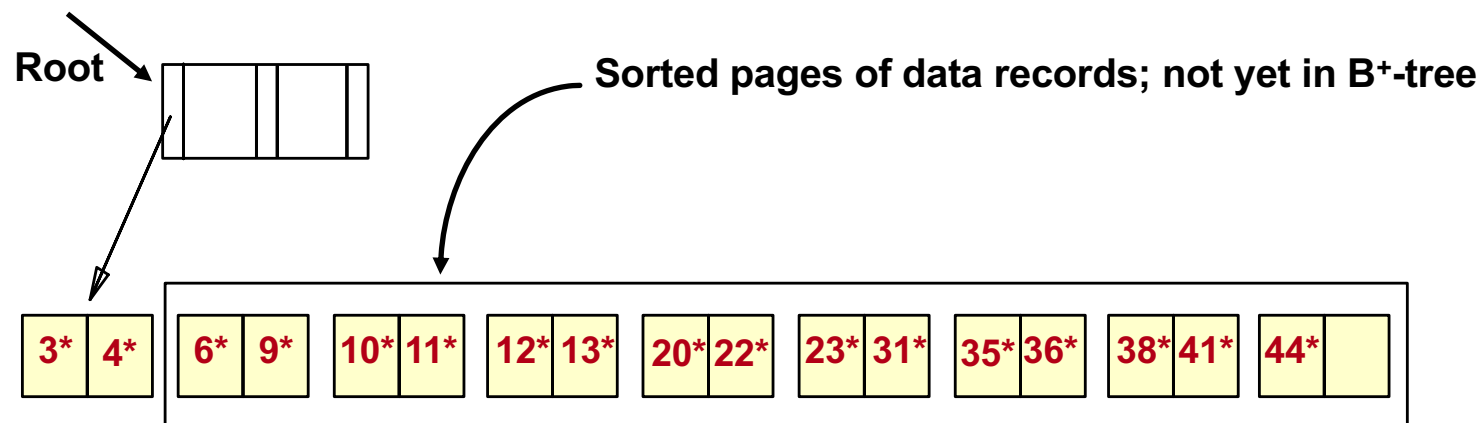


## Example of B<sup>+</sup>-tree File Organization

- In a B<sup>+</sup>-tree file organization, good space utilization is important since records use more space than pointers.
- To improve space utilization, we can involve more sibling nodes in redistribution during splits and merges.
  - Involving 2 siblings in redistribution (to avoid split / merge where possible) results in each node having at least  $\lfloor 2n/3 \rfloor$  entries.


## B<sup>+</sup>-TREE BULK LOADING

- Creating a B<sup>+</sup>-tree by repeatedly inserting records is very slow and can be costly (i.e., need to read and write leaf nodes).
- **Bulk Loading** can be done much more efficiently.
- **Initialization:** Sort all data entries (using external sorting), insert pointer to first (leaf) page in a new (root) page.
- Only need to **write each leaf node once**; **never have to read it**.



# B<sup>+</sup>-TREE BULK LOADING: EXAMPLE

- Index entries for leaf pages are always entered into the **right-most index page** just above the leaf level. When this fills up, it splits. (A split may go up right-most path to the root.)



- Much faster than repeated inserts!

