

# COMP 3311

# DATABASE MANAGEMENT

# SYSTEMS

## LECTURE 17

## QUERY PROCESSING:

## EXPRESSION EVALUATION

# PROJECTION OPERATION: USING SORTING

```
select distinct boatId
from Reserves;
```

Reserves(sailorId, boatId, rDate)

- For duplicate elimination we can use a modified **external sorting** approach.

**Modify Pass 0** of external sorting to **eliminate unwanted attributes**.

- Thus, sorted runs contain **smaller tuples**. (The size reduction depends on the number and the size of the attributes that are eliminated.)

**Modify merge passes** to **eliminate duplicates**.

- Thus, the number of **result** tuples is **smaller than** the **input**. (The difference depends on the number of duplicates.)


**Cost:** **Pass 0:** the original relation is read and the same number of **smaller tuples are written out**  $\Rightarrow$  **fewer write page I/Os than read I/Os**.

**Merge passes:** **smaller tuples are read** and **fewer tuples are written out** in each pass  $\Rightarrow$  **fewer write page I/Os than read I/Os**.

# PROJECTION OPERATION: USING HASHING

## Partitioning Phase

- Read  $r$  using one input page.
- For each tuple, discard unwanted attributes and apply hash function  $h_1$  to the remaining attributes to hash to one of  $M-1$  output pages where  $M$  is the available main memory pages.
- The result is  $M-1$  partitions (of tuples with no unwanted attributes) in which two tuples from different partitions are guaranteed to be distinct.



## Duplicate Elimination Phase

- Read each partition and build an in-memory hash table using hash function  $h_2$  ( $\neq h_1$ ) on all attributes; discard duplicates in the pages.

## Cost:

For partitioning: Read  $r$  and write out each tuple, but with fewer attributes  
 $\Rightarrow$  fewer write page I/Os than read I/Os.

For duplicate elimination: Read the result of the partitioning phase.

# PROJECTION OPERATION: DISCUSSION

- The **sort-based approach** is the **standard** since it handles skewed data better *and* the result is sorted.
- If an **index** on the relation **contains all wanted attributes in its search key**, we can do an *index-only scan*.
  - Apply projection on the index data entries, which are much smaller.
- If an **ordered (i.e., tree) index contains all wanted attributes as the *prefix* of the search key**, we can do even better.
  - Retrieve data entries in order using an index-only scan, discard unwanted attributes and compare the adjacent entries to check for duplicates.

 **For an index-only scan, the index must be dense. Why?**

# SET OPERATIONS: USING SORTING

1. Sort both relation  $r$  and  $s$  on the same attribute eliminating duplicates.
2. The merge phase depends on the operation.

**Intersection** ( $r \cap s$ ): Output a tuple *only if it is in both relations*.

**Union** ( $r \cup s$ ): Output a tuple *only once* if it belongs to both relations.

**Set difference** ( $r - s$ ): Output a tuple *only if it is in the first relation but not in the second one*.

# SET OPERATIONS: USING HASHING

1. Partition relation  $r$  and  $s$  using hash function  $h_1$  on all attributes.

➤  $s$  is the **build input**.

2. For each  $s$ -partition,

- build an in-memory hash table (using  $h_2$ ) on all attributes while discarding duplicates.
- scan the corresponding  $r$ -partition (page-by-page) and, for each tuple  $t_r$  of  $r$ , probe the  $s$ -partition and do the following.

**Intersection** ( $r \cap s$ ): Output  $t_r$  only if it is also in  $s$ .

**Union** ( $r \cup s$ ): Output  $t_r$  if it is not in  $s$ ; also output all tuples of  $s$ .

**Set difference** ( $r - s$ ): Output  $t_r$  only if it is not in  $s$ .  
This computes  $r - s$ . How to compute  $s - r$ ?  
If  $t_r$  is in  $s$ , then remove all matching  $t_s$  tuples from the  $s$ -partition. Output all remaining  $s$  tuples.

# AGGREGATE OPERATIONS

## Without Grouping

- In general, requires scanning the relation.
- If we have a tree index whose search key includes *all attributes* in the *select* or *where* clauses, we can do an index-only scan (e.g., “Find the average age of all sailors” given a dense index on age).

## With Grouping

- Use the duplicate elimination techniques (i.e., sorting or hashing), but instead of eliminating duplicates, *gather them into groups and apply aggregate operations on each group*.

 **Cost is the same as for duplicate elimination.**

 **Do not need to actually form groups—can form groups and compute aggregate for each group on-the-fly.**

- If we have a tree index whose search key includes *all attributes* in *select*, *where* and *group by* clauses, we can do an index-only scan.
- If the group-by attributes form a prefix of the search key, we can retrieve index data entries/tuples in group-by order.

# RELATIONAL ALGEBRA TREE

- Recall that a **relational algebra (operator) tree** represents a relational algebra expression (an SQL query) as a tree.
  - It is a graphical way of representing the **evaluation order** of a relational algebra expression. (**No need for parenthesis!**)
  - Evaluation is **bottom up** using materialization or pipelining.
- 👉 Can annotate the tree to indicate access methods to use for each relation and implementation method to use for each relational operator.

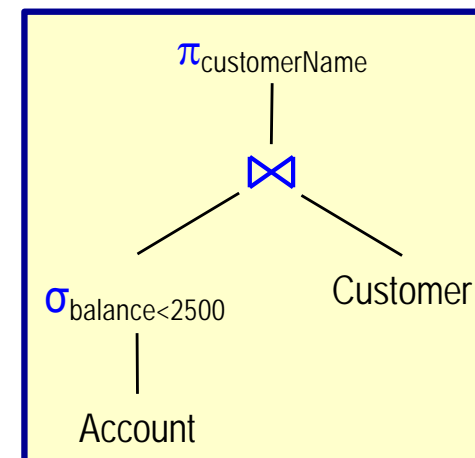
## SQL query

```
select customerName
from Customer natural join Account
where balance < 2500;
```

## Relational algebra expression

$$\pi_{\text{customerName}}(\sigma_{\text{balance} < 2500}(\text{Customer JOIN Account}))$$

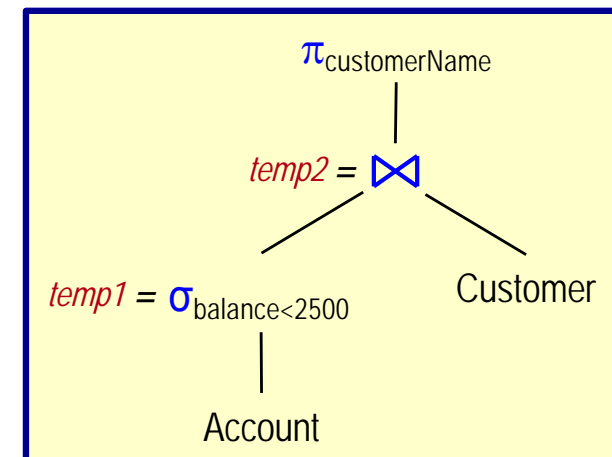
## Relational algebra tree





# EVALUATION USING MATERIALIZATION

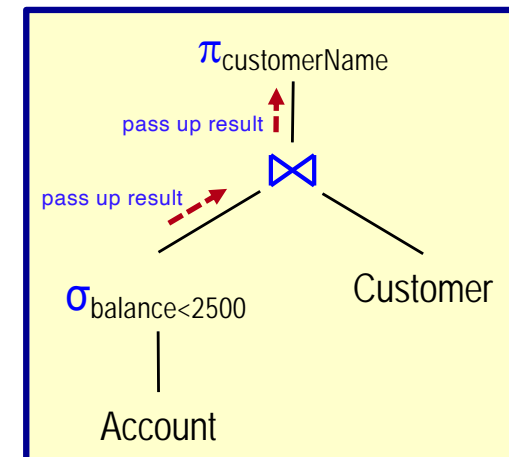
- Evaluate **one operation at a time**, store intermediate result into a **temporary relation** and use it to evaluate the next operation.
- For the query in the figure:
  1. compute and **store** the selection  $\sigma_{\text{balance} < 2500}(\text{Account})$  into *temp1*.
  2. compute and **store** the join of *temp1* with *Customer* in *temp2*.
  3. compute the projection of *temp2* on *customerName*.
- Materialized evaluation is **always applicable**, but the cost of writing and reading intermediate results can be quite high.



**Overall cost:** sum of costs of individual operations +  
cost of writing intermediate results to disk

# EVALUATION USING PIPELINING

- Evaluate **several operations simultaneously**, passing the results to the next operation (**lazy** or **eager**).
- For the query in the figure:
  - compute the selection  $\sigma_{\text{balance} < 2500}(\text{Account})$ , but **do not store** the result; instead, **pass the tuples directly** to the join operation.
  - Similarly, compute the join, but **do not store** the result; instead, **pass the tuples directly** to the projection operation.
- Pipelining is much **cheaper than materialization** since there is **no need to store temporary relations to disk**.
- Need evaluation algorithms that generate output tuples even as tuples are received for input to the operation.



# EVALUATION USING PIPELINING (CONTD)

## Demand-driven or Lazy

- The system repeatedly requests the next tuple from the top level operation.
- Each operation requests the next tuple from its pipelined inputs as required in order to output its next tuple.
- In between calls, the operation has to maintain “state” so it knows what to return next.

## Producer-driven or Eager

- Operations produce tuples continuously and pass them into their pipeline output.
  - A buffer is needed between operations – one operation puts tuples into the buffer, while another operation removes tuples from the buffer.
  - If the buffer is full, the operation waits until there is space in the buffer before generating more tuples.

## EVALUATION USING PIPELINING (CONTD)

- Algorithms that are not able to output tuples as they get input tuples are called **blocking**.

✎ Sort-merge join or hash join are blocking since they require intermediate results to be written to disk and then read back.

- There are algorithm variants that generate (at least some) results on-the-fly as input tuples are read.

✎ Hybrid hash join generates output tuples even as probe relation tuples in the in-memory partition (partition 0) are read.

**Pipelined Join:** Hybrid hash join can be modified to buffer partition 0 tuples of both relations in-memory, reading them as they become available, and outputting matching result tuples between partition 0 tuples.

- When a new  $r_0$  tuple is found, match it with existing  $s_0$  tuples, output matches, and save it in  $r_0$ .
- This can also be done symmetrically for  $s_0$  tuples.

# QUERY PROCESSING: SUMMARY

## Query Processing

1. Transform the SQL query into a relational algebra expression.
2. Generate different query execution plans and evaluate their cost.
3. Select an optimal query execution plan.

## Operation Processing

**Selection:** linear scan; index lookup.

**Sorting:** external sort-merge if too large to fit into memory

**Join:** block nested-loop; indexed nested-loop; merge join; hash join

**Projection, duplicate elimination, set operations:** sorting; hashing

**Aggregate operations:** linear scan; index-only scan; sorting; hashing

**Expression evaluation:** materialization; pipelining.