

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 21

CONCURRENCY CONTROL:

LOCK-BASED PROTOCOLS

CONCURRENCY CONTROL: OUTLINE

Lock-based Protocols

- Two-phase Locking Protocols
- Deadlock Handling
- Graph-based Protocols

Timestamp-based Protocols

- Timestamp-ordering Protocols
- Validation-based Protocols

Multiversion Schemes

Snapshot Isolation

GOAL OF CONCURRENCY CONTROL SCHEMES

Atomicity Handled by database recovery.

- Either **all** transaction operations are properly **reflected** in the database **or none** are.

Consistency Handled by concurrency control.

- Execution of a transaction in isolation **preserves the consistency** of the database.

Isolation Handled by concurrency control.

- Although multiple transactions may execute concurrently, **each transaction must be unaware of other concurrently executing transactions.**

Durability Handled by database recovery.


- After a transaction completes successfully, the **changes** it made to the database **persist**, even if there are system failures.

 **Generate schedules that enforce the isolation property.**

CONCURRENCY CONTROL SCHEME CORRECTNESS

T_1	T_2	T_3
	read(Z) read(Y) write(Y)	
		read(Y) read(Z)
read(X) write(X)		
		write(Y) write(Z)
	read(X)	
read(Y) write(Y)		
	write(X)	

Based on serializability



The schedule is **not** serializable.
Therefore, it should **not** be generated
by any concurrency control scheme.

LOCKING

Lock: a mechanism to control concurrent access to a data item.



- A data item Q can be locked in one of two modes:
 1. **shared-mode (shared lock)**
 - Can only read Q . An **s-lock** is requested using a **lock-s(Q)** instruction.
 2. **exclusive-mode (exclusive lock)**
 - Can both read and write Q . An **x-lock** is requested using a **lock-x(Q)** instruction.
 - A data item Q is unlocked using an **unlock(Q)** instruction.
- A transaction must make a lock request to the **concurrency-control manager** before accessing a data item.
- A transaction can proceed only after a request is granted.

👉 **The concurrency control manager should allow *only* conflict-serializable schedules.**

LOCK-COMPATIBILITY MATRIX

Lock Mode	Shared (S)	Exclusive (X)
Shared (S)	true (grant)	false (deny)
Exclusive (X)	false (deny)	false (deny)

Shared lock - can *only* read

Exclusive lock - can read and write

- A transaction may be granted a lock on a data item *Q* if the requested lock is *compatible* with locks already held on *Q* by other transactions.
- Any number of transactions can hold **shared locks** on a data item *Q*.
- If any transaction holds an **exclusive lock** on the data item, then *no other transaction may hold any lock on the data item*.
- If a lock cannot be granted, *the requesting transaction is made to wait* until all incompatible locks held by other transactions have been **released**. *The lock is then granted*.

LOCKING EXAMPLE

Start with A=100, B=200		
	T_1	T_2
	read(B)	
	B := B - 50	
B=150	write(B)	
	read(A)	
	A := A - 50	
A=50	write(A)	
		read(A)
		read(B)
		display(A+B)
		200

Locking requires a set of rules, called a locking protocol, that all transactions follow when requesting and releasing locks. The rules restrict the set of possible schedules.

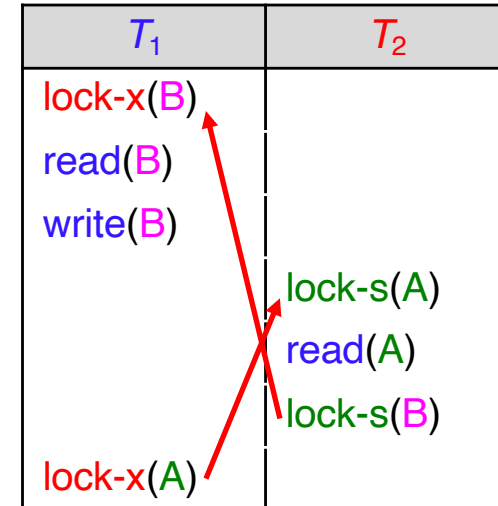
T_1	T_2	CC Manager
lock-x(B)		grant-X(B)
read(B)		
B := B - 50		
write(B)		
unlock(B)		
	lock-s(A)	grant-S(A)
A=100	read(A)	
	unlock(A)	
	lock-s(B)	grant-S(B)
B=150	read(B)	
	unlock(B)	
	display(A+B)	250 INCORRECT!
lock-x(A)		grant-X(A)
read(A)		
A := A - 50		
write(A)		
unlock(A)		

Serializability may not be ensured if a transaction unlocks a data item immediately after its final access of that data item.

DEADLOCKS

- A system is **deadlocked** if there is a set of transactions such that **every transaction is waiting for another transaction in the set**.
- Neither T_1 nor T_2 can make progress — executing **lock-s(B)** causes T_2 to wait for T_1 to release its lock on **B**, while executing **lock-x(A)** causes T_1 to wait for T_2 to release its lock on **A**.

T_1	T_2
lock-x(B) read(B) write(B)	
	lock-s(A) read(A) lock-s(B)
lock-x(A)	



☞ Such a situation is called a **deadlock**.

☞ To handle a deadlock, either T_1 or T_2 must be **rolled back and its locks released**.


☞ A rollback requires all transaction operations to be undone.

☞ A locking protocol needs to be able to handle deadlocks.

STARVATION

Starvation (waiting forever) is possible if the concurrency control manager is **badly designed**.

- A transaction may be waiting for an **x-lock** on an item, while a sequence of other transactions request, and are granted, an **s-lock** on the same item.



👉 **A locking protocol needs to be able to handle starvation.**

TWO-PHASE LOCKING (2PL) PROTOCOL

Phase 1: Growing Phase

A transaction may **obtain locks**, but may not release any locks.

Phase 2: Shrinking Phase

A transaction may **release locks**, but may not obtain any new locks.

Using 2PL, transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its **final lock**).

✎ If a schedule is **executed by 2PL**, then
it must be **conflict serializable**.

✎ If a schedule is **conflict serializable**, then
it may or may not be **executed by 2PL**.

✎ **Not all conflict serializable schedules are allowed by 2PL.**

LOCK CONVERSIONS

- Allow **shared locks** to be **upgraded** to **exclusive locks** and **exclusive locks** to be **downgraded** to **shared locks**.

Phase 1: Growing Phase (**request or upgrade locks**)

- A transaction can acquire a **lock-s** on a data item.
- A transaction can acquire a **lock-x** on a data item.
- A transaction can convert a **lock-s** to a **lock-x** (**upgrade**) on a data item.

Phase 2: Shrinking Phase (**release or downgrade locks**)

- A transaction can release a **lock-s** on a data item.
- A transaction can release a **lock-x** on a data item.
- A transaction can convert a **lock-x** to a **lock-s** (**downgrade**) on a data item.

 **Schedules are cascadeless if exclusive–locks are held until the end of the transaction.**

STRICT AND RIGOROUS TWO-PHASE LOCKING

Under 2PL, cascading roll-back is possible.

Strict Two-phase Locking

- Requires that **all exclusive-mode locks** be **held until a transaction commits** (shared-mode locks can be released anytime).
- Ensures that any data written by an uncommitted transaction are locked in exclusive mode until the transaction commits.

Rigorous Two-phase Locking

- Requires that **all locks** (both shared-mode and exclusive-mode) be **held until a transaction commits**.
- The transactions can be serialized in the order in which they commit.

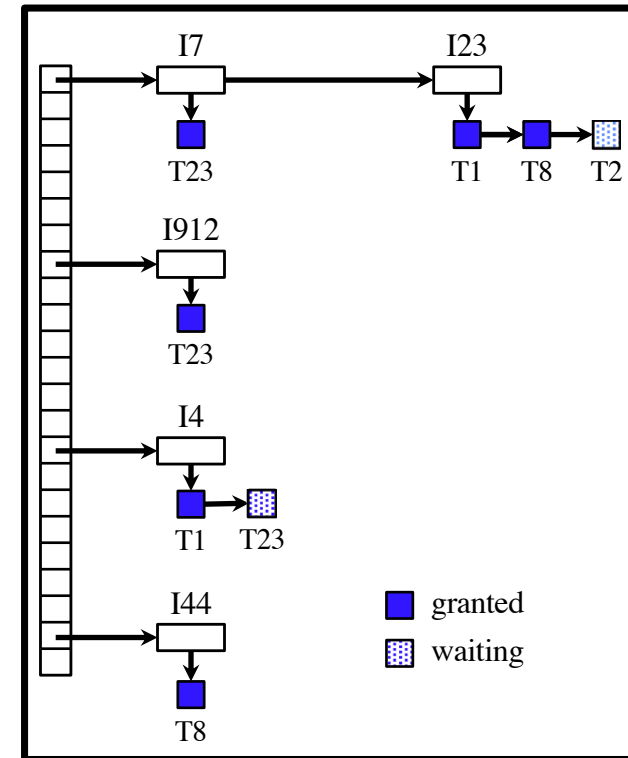
 **Strict and rigorous 2PL schedules are cascadeless.**

IMPLEMENTATION OF LOCKING

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests.
- The lock manager replies to a lock request by sending a **lock grant message** (or a message asking the transaction to roll back in case of a deadlock).
- The requesting transaction **waits** until its request is answered.
- The lock manager maintains a data structure called a **lock table** to record granted locks and pending requests.
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked.

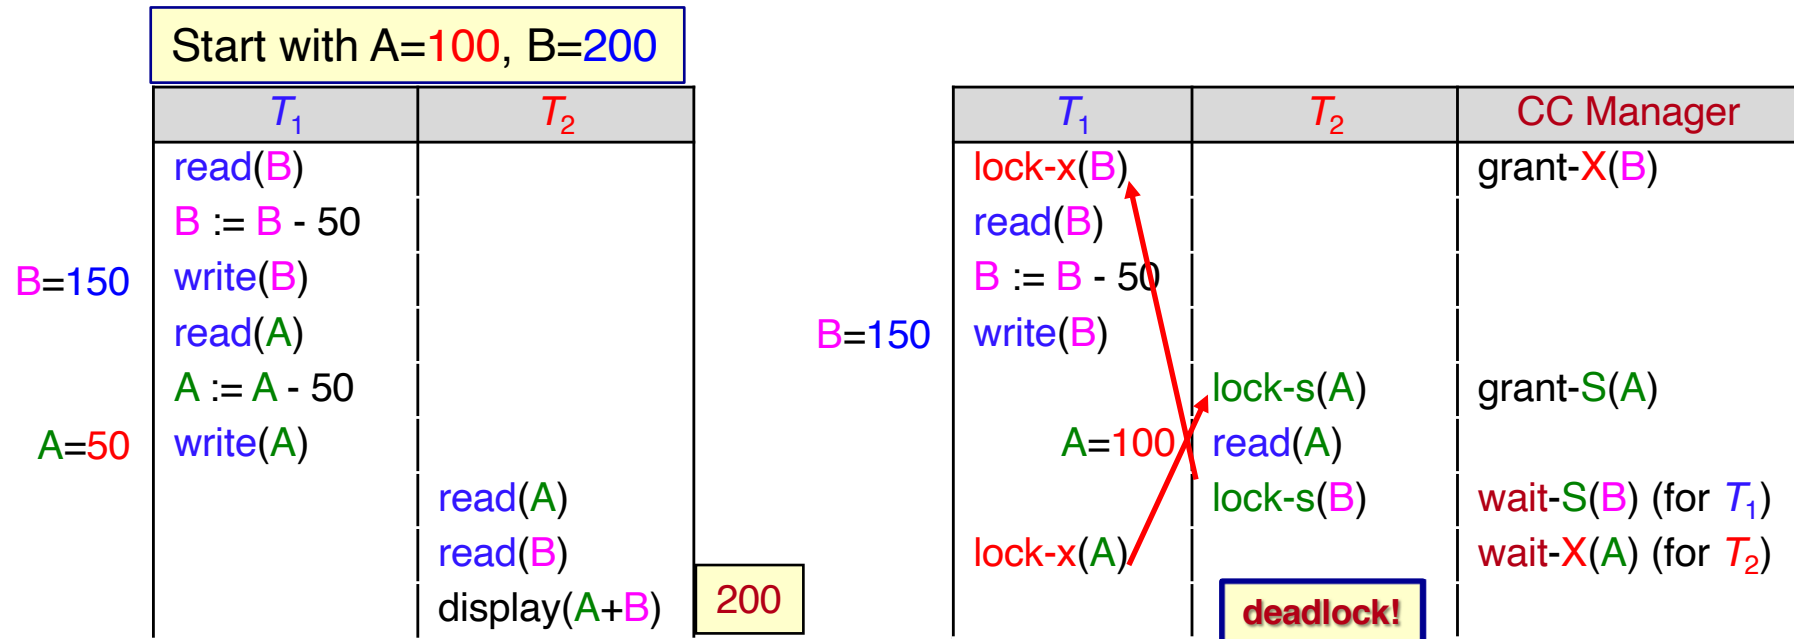
LOCK TABLE EXAMPLE

- Dark blue rectangles indicate **granted locks**, light blue ones indicate **waiting requests**.
 - The lock table also records the **type of lock** granted or requested.
 - A new request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks.
 - Unlock requests result in the request being deleted and later requests being checked to see if they can now be granted.
 - If a transaction aborts, all waiting or granted requests of the transaction are deleted.
- The lock manager also may keep a list of locks held by each transaction, to implement this efficiently.



➡ **This algorithm guarantees freedom from starvation.**

DEADLOCK HANDLING



👉 2PL permits deadlocks.

The *potential* for deadlock exists in most locking protocols.
Deadlocks are a necessary evil of locking protocols.

DEADLOCK PREVENTION

- A **deadlock prevention protocol** ensures that the system will *never* enter into a deadlock state.

Strategy 1: Order lock requests

- Require that each transaction locks all its data items *before* it begins execution (**pre-declaration**).
 - Often hard to predict, beforehand, what data items need to be locked.
 - Locked data items may be unused for a long time.
- Impose a partial/total ordering of all data items and require that a transaction can *lock data items only in the order specified by the partial/total order* (e.g., **tree protocol**).

Strategy 2: Preemption and/or rollback

- Preempt and/or rollback a transaction when needed.
- Use **transaction timestamps** to control preemption and rollback.

DEADLOCK PREVENTION (CONT'D)

Wait-die Scheme — non-preemptive

- An **older transaction** may **wait** for a **younger one** to release a data item.
- A **younger transaction** never waits for an **older one**; it is rolled back instead.
- A transaction **may die several times** before acquiring the needed data item.

Wound-wait Scheme — preemptive

- An **older transaction** **wounds** (forces the rollback of) a **younger transaction** instead of waiting for it.
- A **younger transaction** may wait for an **older one** to release a data item.
- There may be **fewer rollbacks** than in the wait-die scheme.
- In both schemes, a **rolled back transaction is restarted** with its original timestamp so that older transactions have precedence over younger ones, thus, avoiding starvation.

DEADLOCK PREVENTION (CONT'D)

Timeout-Based Schemes

- A transaction waits for a lock only for a specified amount of time.
- After a pre-defined waiting period, the transaction is rolled back.
- Simple to implement; but starvation is possible.
- Often difficult to determine a good value of the timeout interval.

DEADLOCK DETECTION


- Deadlocks can be detected using a **wait-for graph** $G = (V, E)$, where
 - V is a set of vertices (all the transactions in the system).
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
 - If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j implying that T_i is waiting for T_j to release a data item.
 - When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted into the wait-for graph.
 - This edge is removed only when T_j is no longer holding a data item needed by T_i .

 **The system is in a deadlock state if and only if the wait-for graph has a cycle.**


 **The system must invoke a deadlock-detection algorithm periodically to look for cycles in the wait-for graph.**

DEADLOCK DETECTION (CONT'D)

Example



Wait-for graph
without a cycle.



Wait-for graph
with a cycle.

✋ This is not the same as a precedence graph!

DEADLOCK RECOVERY

Victim Selection

- Select as a **victim** the transaction that will **incur minimum cost**.

Rollback

- Need to determine how far to roll back the transaction.
 - **Total rollback**: abort the transaction and then restart it.
 - **Partial rollback**: roll back the transaction only as far as necessary to break the deadlock. (Requires the system to maintain additional information.)

Starvation


- Can happen if the same transaction is always chosen as the victim.
- Can include the number of rollbacks in the cost factor to avoid starvation.

GRAPH-BASED PROTOCOLS

- For protocols that are not two phase, additional information is needed about how each transaction accesses the database.
- The simplest such protocols require knowledge about the order in which the data items will be accessed.
- This knowledge can be acquired by imposing a partial ordering on the set $\mathbf{D} = \{d_1, d_2, \dots, d_h\}$ of all data items.
 - If d_i precedes d_j in the ordering, then any transaction accessing both d_i and d_j must access d_i before accessing d_j .
 - The set \mathbf{D} may be viewed as a directed acyclic graph, called a database graph.
- For simplicity, we consider only those graphs that are rooted trees.

TREE PROTOCOL



1. Only **lock-x** instructions are allowed.
2. The first lock by T_i may be on any data item.
3. Subsequently, a data item Q can be **locked** by T_i only if the parent of Q is **currently locked** by T_i .
4. Data items may be **unlocked at any time**.
5. A data item that has been unlocked by T_i **cannot be locked again** by T_i .



👉 **All legal schedules under the tree protocol are conflict serializable.**

TREE PROTOCOL: SERIALIZABLE SCHEDULE

T_1	T_2	T_3	T_4
lock-X(B)	lock-X(D)		
	lock-X(H)		
	unlock(D)		
lock-X(E)			
lock-X(D)			
unlock(B)			
unlock(E)			
		lock-X(B)	
		lock-X(E)	
	unlock(H)		
lock-X(G)			
unlock(D)			
			lock-X(D)
			lock-X(H)
			unlock(D)
			unlock(H)
		unlock(E)	
		unlock(B)	
unlock(G)			



Precedence Graph

Equivalent serial
schedule:


T_2, T_1, T_4, T_3

or

T_2, T_1, T_3, T_4

TREE PROTOCOL: PROPERTIES

- The tree protocol
 - ensures conflict serializability.
 - is deadlock free.
- Unlocking may occur earlier than in the two-phase locking protocol.
 - Shorter waiting times and an increase in concurrency.
- However, a transaction may have to lock data items that it does not access.
 - Increased locking overhead and additional waiting time.
 - Potential decrease in concurrency.



👉 **Schedules not possible under two-phase locking are possible under the tree protocol, and vice versa.**