

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 22

CONCURRENCY CONTROL:

TIMESTAMP-BASED PROTOCOLS

TIMESTAMPS



Each transaction is issued a fixed **timestamp** $TS(T_i)$ by the system before it starts execution.

If an **old transaction** T_i has timestamp $TS(T_i)$, a **new transaction** T_j is assigned timestamp $TS(T_j)$ such that $TS(T_i) < TS(T_j)$.

👉 **The timestamps determine the serializability order.**

- To implement serializability, **two timestamp values** are associated with **each data item** Q

WTS(Q) (**write timestamp**) the largest timestamp of any transaction that executed **write**(Q) successfully.

RTS(Q) (**read timestamp**) the largest timestamp of any transaction that executed **read**(Q) successfully.

and updated whenever a new **read**(Q) or **write**(Q) instruction is executed.

TIMESTAMP-ORDERING PROTOCOL: READ

The timestamp-ordering protocol ensures that any **conflicting read** and **write** operations are **executed in timestamp order**.

Suppose a transaction T_i issues a **read**(Q).


1. If $TS(T_i) < WTS(Q)$, then T_i needs to read a value of Q that was already overwritten.



✎ **Reject** the **read** operation and **roll back** T_i .

✎ **Restart** T_i with a new (larger) timestamp $TS'(T_i)$.

2. If $TS(T_i) \geq WTS(Q)$, then execute the **read** operation and set $RTS(Q)$ to the maximum of $RTS(Q)$ and $TS(T_i)$.



TIMESTAMP-ORDERING PROTOCOL: **WRITE**

Suppose that transaction T_i issues a **write**(Q).

1. If $TS(T_i) < RTS(Q)$, then the value of Q that T_i is writing was **needed previously** and the system assumed that that value would never be written (since a different value was already read).



✋ **Reject** the **write** operation and **roll back** T_i .

2. If $TS(T_i) < WTS(Q)$, then T_i is attempting to write an **obsolete value** of Q .

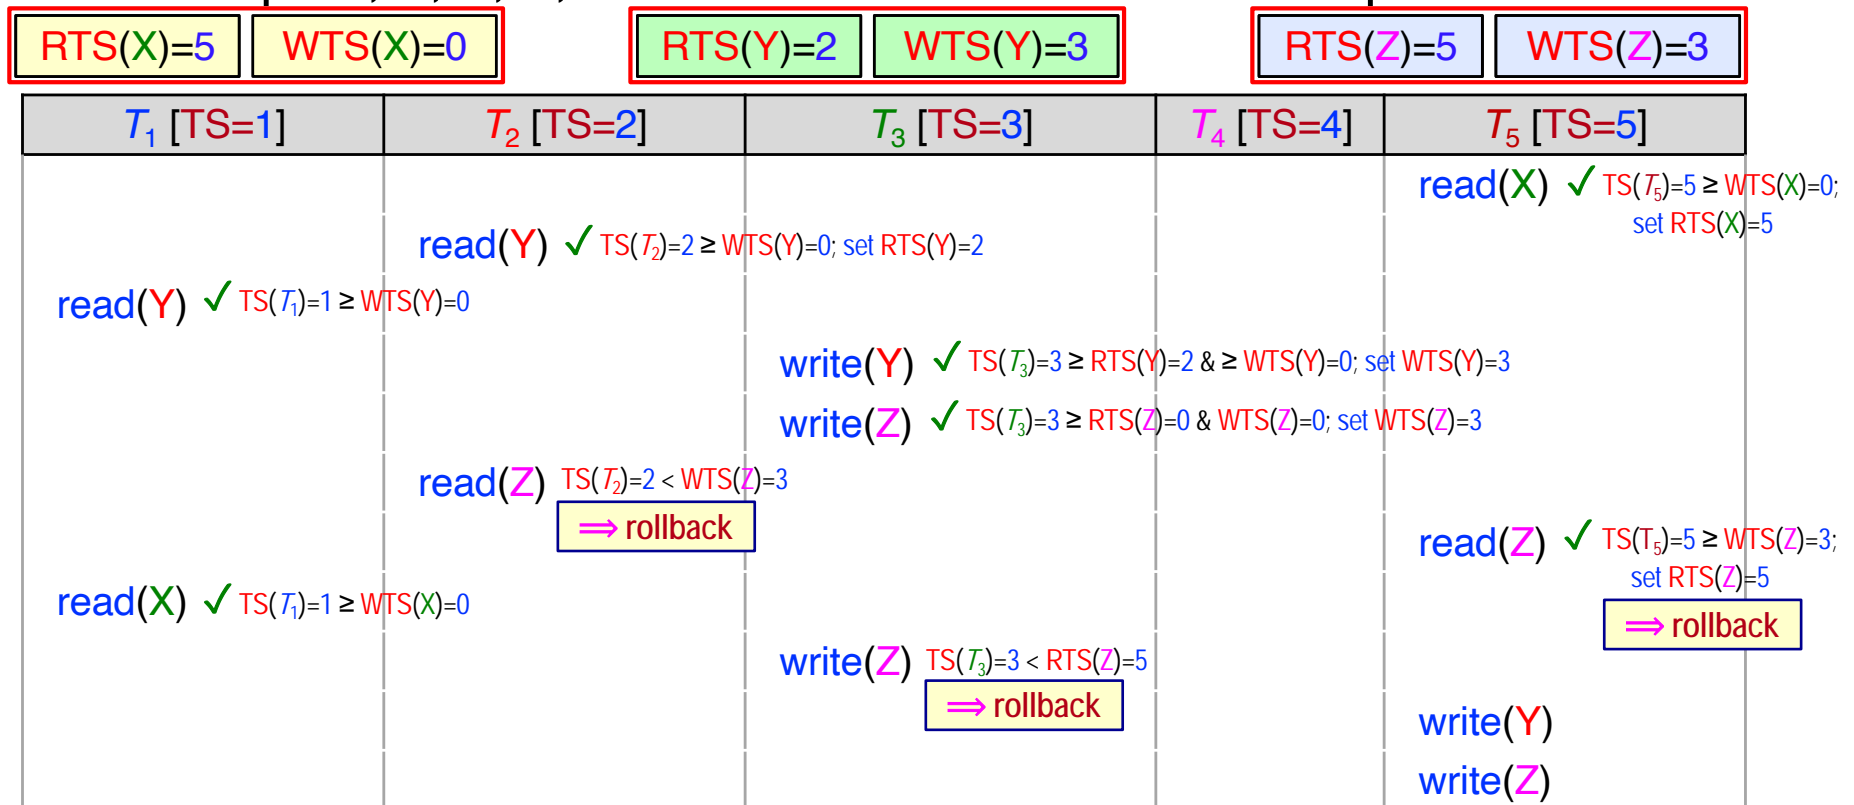


✋ **Reject** the **write** operation and **roll back** T_i .

3. Otherwise, execute the **write** operation and set $WTS(Q)$ to $TS(T_i)$.

TIMESTAMP-ORDERING PROTOCOL EXAMPLE

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5. Assume initial R/W timestamp of all items is 0.



Read

If $TS(T_i) < WTS(Q)$ **rollback**
 If $TS(T_i) \geq WTS(Q)$
 $RTS(Q) = \max(TS(T_i), RTS(Q))$


Write

If $TS(T_i) < RTS(Q)$ **rollback**
 If $TS(T_i) < WTS(Q)$ **rollback**
 Otherwise $WTS(Q) = TS(T_i)$



TIMESTAMP-ORDERING PROTOCOL CORRECTNESS

- The timestamp-ordering protocol ensures conflict serializability since all the arcs in the precedence graph are of the form:



✋ **Conflicting instructions are processed in timestamp order.**

- Since, there can be no cycles in the precedence graph, the timestamp protocol ensures freedom from deadlock and no transaction ever waits.
- However, starvation of long transactions is possible and a schedule may not be recoverable.

RECOVERABILITY AND CASCADING ROLLBACK

- The problem with the timestamp-ordering protocol:
 - Suppose T_i aborts, but T_j has read a data item written by T_i .
 - Then T_j must abort; if T_j had been allowed to commit before T_i aborted, the schedule is not recoverable.
 - Further, any transaction that has read a data item written by T_j must abort.

 **This can lead to cascading rollback.**

Solution

- A transaction is structured such that *all its writes* are performed at the end of its processing.
- **All writes of a transaction form an atomic action**; no transaction may execute while a transaction is being written.
- A transaction that aborts is restarted with a new timestamp.

THOMAS' WRITE RULE

Read

If $TS(T_i) < WTS(Q)$ **rollback**
 If $TS(T_i) \geq WTS(Q)$
 $RTS(Q) = \max(TS(T_i), RTS(Q))$

Write

If $TS(T_i) < RTS(Q)$ **rollback**
 If $TS(T_i) < WTS(Q)$ **rollback**
 Otherwise $WTS(Q) = TS(T_i)$

$RTS(Q)=1$ $WTS(Q)=2$

T_1 [TS=1]	T_2 [TS=2]
read(Q) ✓	$TS(T_1)=1 \geq WTS(Q)=0$; set $RTS(Q)=1$
	write(Q) ✓ $TS(T_2)=2 \geq RTS(Q)=1$ & $\geq WTS(Q)=0$; set $WTS(Q)=2$
write(Q)	$TS(T_1)=1 < WTS(Q)=2 \Rightarrow$ rollback

👉 **Rollback of T_1 is actually unnecessary!**

- Since T_2 has already written Q (with $TS=2$), the value T_1 is attempting to write (with $TS=1$) *will never be read*.
 - Any transaction with timestamp less than $TS(T_2)$ that attempts to read Q will be rolled back since its timestamp will be less than $WTS(Q)$.
 - Any transaction with a timestamp greater than $TS(T_2)$ must read the value of Q written by T_2 , rather than the one T_1 is attempting to write.

👉 **Ignore write of T_1 !**

Read (unchanged)

If $TS(T_i) < WTS(Q)$ **rollback**
 If $TS(T_i) \geq WTS(Q)$
 $RTS(Q) = \max(TS(T_i), RTS(Q))$

Write (Thomas' write rule)

If $TS(T_i) < RTS(Q)$ **rollback**
 If $TS(T_i) < WTS(Q)$ **ignore**
 Otherwise $WTS(Q) = TS(T_i)$

VALIDATION-BASED PROTOCOLS

Assumption: Most transactions are read-only.

- Transactions execute in two or three phases:
 1. **Read phase** write data items in temporary, local variables.
 2. **Validation phase** determine whether the transaction can proceed to the write phase; abort transaction if test fails.
 3. **Write phase** apply write operations to database. Read-only transactions omit this phase.
- The system maintains three timestamps for each transaction:
 1. **start**(T_i) – the time when T_i started execution.
 2. **validation**(T_i) – the time when T_i finished its read phase and started its validation phase.
 3. **finish**(T_i) – the time when T_i finished its write phase.

VALIDATION TEST

- For all transactions T_k with $TS(T_k) < TS(T_i)$, one of the following conditions must hold:
 1. $finish(T_k) < start(T_i)$.
Serializable since T_k completes before T_i started.
 2. data items written by (T_k) \cap data items read by (T_i) = \emptyset and $(start(T_i) < finish(T_k) < validation(T_i))$ (ensures that the writes of T_k and T_i do not overlap).
Serializable since the writes of T_k do not affect the read of T_i , and T_i cannot affect the read of T_k .
- Ensures no cascading rollbacks, but starvation of long transactions is possible.

 **Optimistic concurrency-control.**

VALIDATION-BASED PROTOCOL EXAMPLE

T_1	T_2
read(B)	read(B)
	$B := B - 50$
	read(A)
	$A := A - 50$
read(A)	
validate	
display(A+B)	validate
	write(B)
	write(A)


- Note that the writes to the actual variables are performed only after the validation phase of T_2 .
- Thus, T_1 reads the old values of A and B, and this schedule is serializable.

MULTIVERSION SCHEMES

- Multiversion schemes **keep old versions of data items**, **labelled with timestamps**, to increase concurrency.
 - Multiversion Timestamp Ordering
 - Multiversion Two-Phase Locking
- Each data item Q has a sequence of versions $\langle Q_1, Q_2, \dots, Q_m \rangle$.
- Each version Q_k contains three data fields:
 - Content** The value of version Q_k .
 - WTS**(Q_k) The timestamp of the transaction that created (wrote) version Q_k .
 - RTS**(Q_k) The **largest timestamp** of a transaction that **successfully read** version Q_k .
- Each successful write by a transaction T_i creates a **new version** Q_k of Q , and sets Q_k 's **RTS** and **WTS** to $TS(T_i)$.

MULTIVERSION TIMESTAMP ORDERING: **READ**

Let Q_k be the version of Q whose write timestamp is the largest write timestamp *less than or equal to* $TS(T_i)$.



- If transaction T_i issues a $read(Q)$, then the value returned is the content of version Q_k (i.e., T_i reads the most recent version that comes before it in time).

👉 If $TS(T_i) > RTS(Q_k)$, then set $RTS(Q_k) = TS(T_i)$.

👉 **Reads never fail and never wait as an appropriate version can always be found.**

Read

Reads always succeed

set $RTS(Q_k) = \max(TS(T_i), RTS(Q_k))$

MULTIVERSION TIMESTAMP ORDERING: **WRITE**

Let Q_k be the version of Q whose write timestamp is the largest write timestamp *less than or equal to* $TS(T_i)$.


If transaction T_i issues a **write**(Q), then

- a) if $TS(T_i) < RTS(Q_k)$, then **roll back** T_i .


Since some other transaction T_j that should read T_i 's write (in the serialization order defined by the timestamp values) has already read a version created by a transaction older than T_i .



- b) if $TS(T_i) = WTS(Q_k)$, **overwrite** the **contents** of Q_k
 $\Rightarrow Q_k$ was written before also by T_i .



- c) if $TS(T_i) > WTS(Q_k)$ **create** a **new version** of Q and set **R/WTS** of new version to $TS(T_i)$.



 **Conflicts are resolved through aborting transactions.**

Write

If $TS(T_i) < RTS(Q)$ **rollback**

If $TS(T_i) = WTS(Q)$ **overwrite contents**

If $TS(T_i) > WTS(Q)$ **create new version**

set $R/WTS(Q') = TS(T_i)$

MULTIVERSION SCHEMES: NOTES


- The schedule of **Exercise 2** will terminate successfully although it is **not conflict serializable**.
- The concept of **conflict serializability does not apply to multi-version protocols**, since **read operations do not conflict with previous writes** (there is always a version to read).

Versions are removed as follows:

- Suppose that there are two versions Z_1 and Z_2 with timestamps 3 and 1 (as could happen in **Exercise 2**) that both have a smaller timestamp than the oldest transactions (i.e., T_1 , T_2 and T_3 have committed and the new transactions have timestamps ≥ 4).
- Then the version Z_2 with timestamp 1 will not be used again and is deleted.
- If any subsequent transaction needs to read Z , it will read Z_1 with timestamp 3.
- In **Exercise 2** it is as if the last write(Z) was never executed, which is OK because T_3 executes a **blind** write(Z) (i.e., write with no previous read).
- If, on the other hand, T_3 had read Z before writing it, the schedule would fail at the last write of T_1 .

SNAPSHOT ISOLATION

- Each transaction **works on its own private copy (snapshot)** of the data items it reads and writes.
- If most transactions are read-only or if their updated data items do not overlap, then they cannot conflict with each other.
- 👉 **Results in higher concurrency levels.**
- Concurrency control is only needed if **updates conflict**.



SNAPSHOT ISOLATION VALIDATION STEPS

First Committer Wins

if a data item that T intends to write has already been updated by any concurrent transaction:
then T **aborts**;
else T **commits** and writes its updates to the database;

First Updater Wins


if T 's request for an **x-lock** on a data item it intends to update is granted:
 if the data item has already been updated by any concurrent transaction:
 then T **aborts**;
 else T **proceeds with its execution** including possibly committing;
else // data item is locked \Rightarrow wait until transaction holding the lock **aborts** or **commits**
 if the transaction holding the lock **aborts**:
 then the lock is released and T locks the data item;
 if the data item has already been updated by any concurrent transaction:
 then T **aborts**;
 else T **proceeds with its execution** including possibly committing;
 if the transaction holding the lock **commits**:
 then T **aborts**;

SNAPSHOT ISOLATION: NOTES

- Snapshot isolation is attractive because **overhead is low** and **no abort occurs unless two concurrent transactions update the same data item**.
- Snapshot isolation **does not ensure serializability!**

T_1	T_2
read(A)	
read(B)	
	read(B)
	read(A)
write(A)	
	write(B)

Precedence Graph



- The schedule is allowed by snapshot isolation since T_1 and T_2 update different data items.
- However, it is **not serializable!**

SNAPSHOT ISOLATION: NOTES (CONTD)

- Although **snapshot isolation** permits schedules that are not serializable, **it is widely used in commercial systems** (e.g., Oracle, SQL Server).
 - If database consistency is preserved, then non-serializable executions are OK.
 - Constraints can be used to help ensure database consistency when non-serializable schedules are allowed.

CONCURRENCY CONTROL: SUMMARY

While most protocols discussed **ensure correctness**,

✎ **A correct schedule may not be permitted by a protocol.**

✎ **The more correct schedules allowed by a protocol, the greater the degree of concurrency.**

- The protocols differ in the **way they handle conflicts**:
 - i. **Lock-based protocols** make transactions **wait** (thus, they can result in deadlocks).
 - ii. **Timestamp-based protocols** make transactions **abort** (thus, there are no deadlocks, but aborting a transaction may be more expensive).
- **Recoverability** (i.e., no rollback after a commit) is a necessary property of a schedule.

✎ **To ensure recoverability, transaction T_i can commit only after all transactions that wrote items that T_i read have committed.**

CONCURRENCY CONTROL: SUMMARY (CONT'D)

- A **cascading rollback** happens when an *uncommitted* transaction must be rolled back because it read an item written by a transaction that aborted.

✎ It is desirable to have **cascadeless schedules**.

✎ To achieve this property a transaction should **only be allowed to read items written by committed operations**.

- If a schedule is **cascadeless**, it is also **recoverable**.
 - **Strict 2PL** ensures cascadeless schedules by **releasing all exclusive locks of transaction T_i after T_i commits** (therefore other transactions cannot read the items locked by T_i at the same time).
 - **Timestamp-based protocols** can also achieve cascadeless schedules by **performing all the writes at the end of the transaction as an atomic operation**.